

PathFence: Reducing Cross-Path Dependencies in Microservices

Xuhang Gu and Qingyang Wang

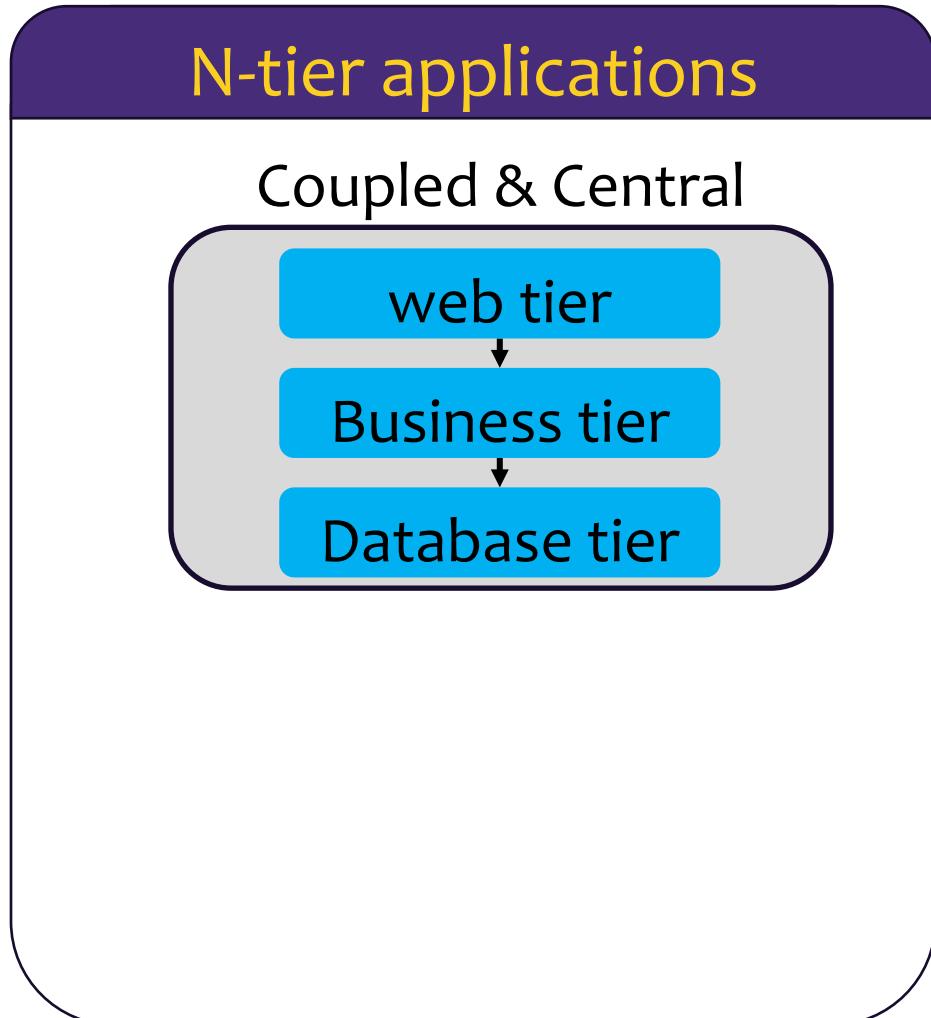
Louisiana State University

From Monolithic N-tier to Microservices

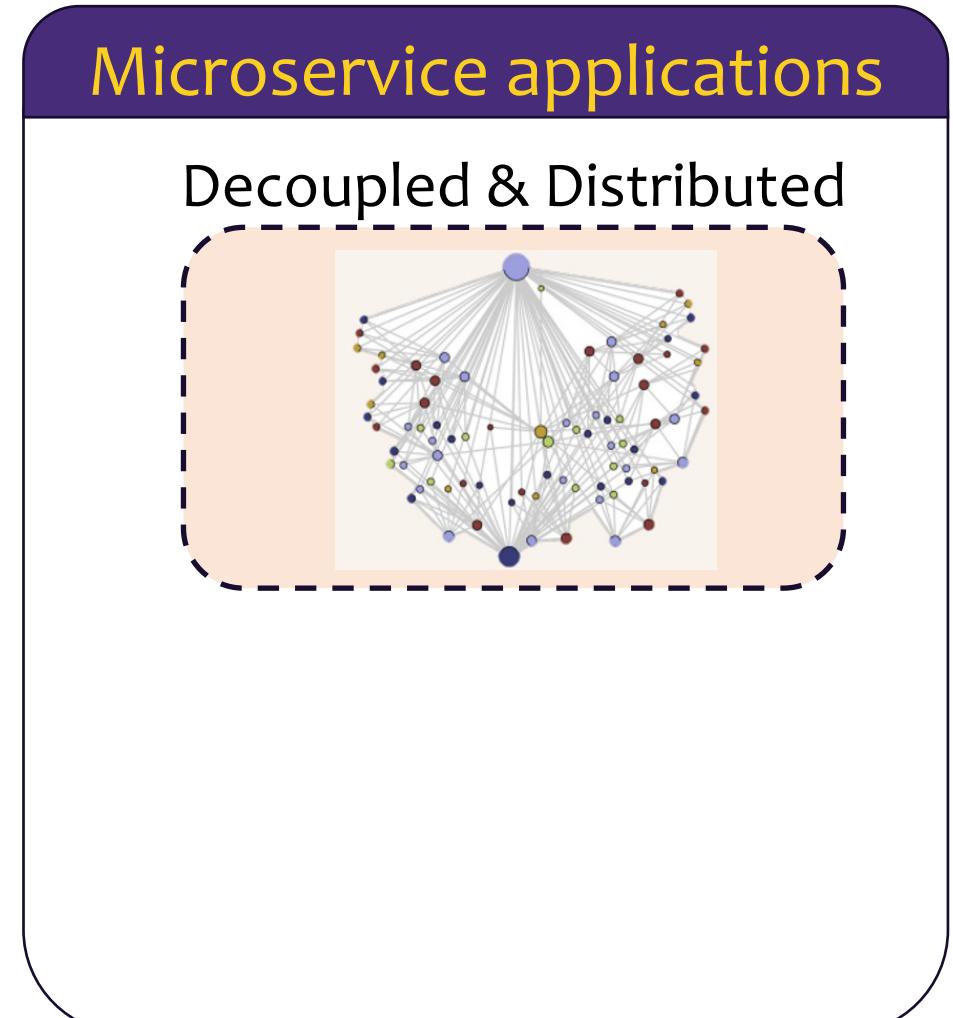
N-tier applications

Microservice applications

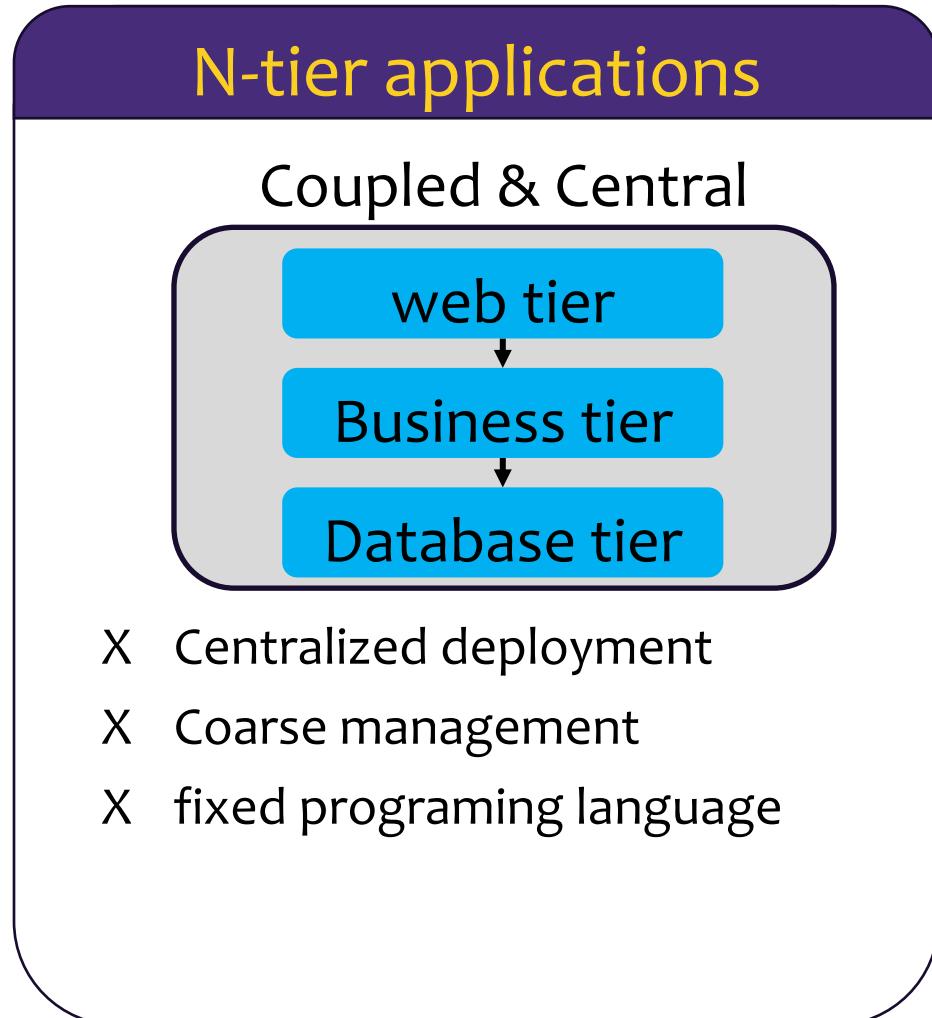
From Monolithic N-tier to Microservices



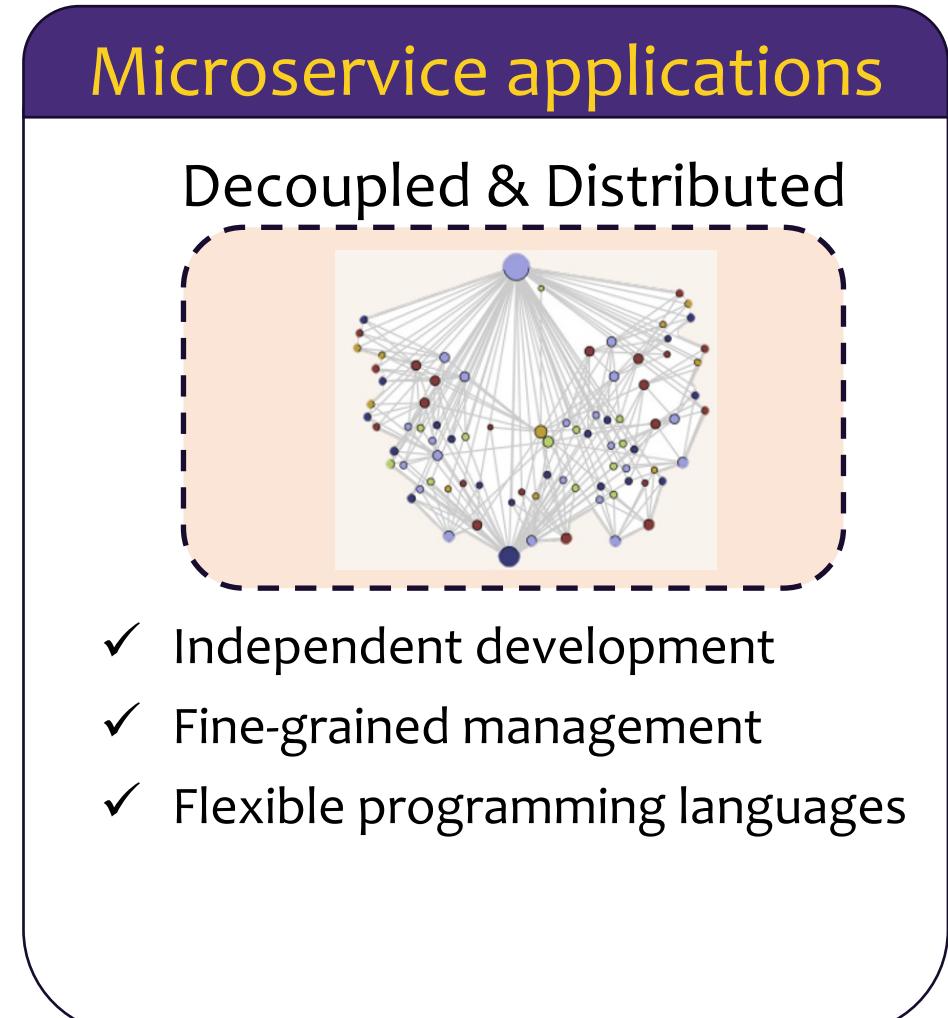
Shifting to



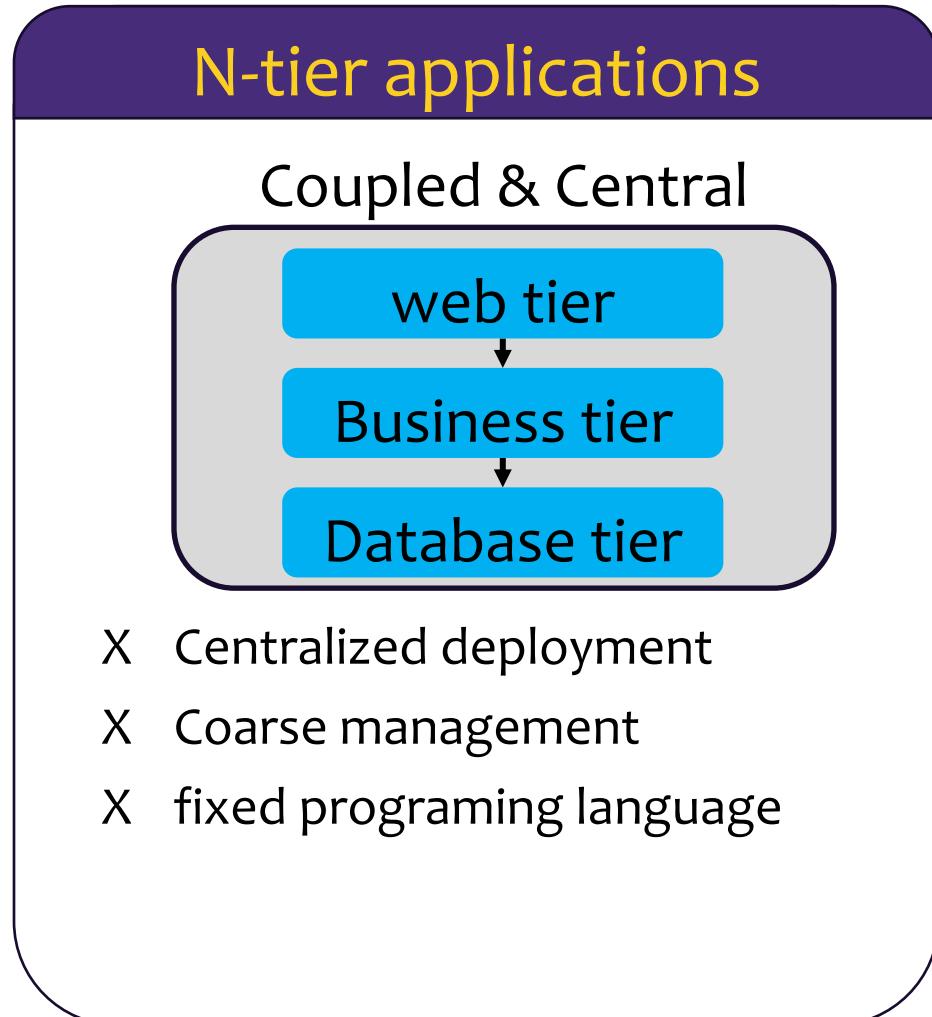
From Monolithic N-tier to Microservices



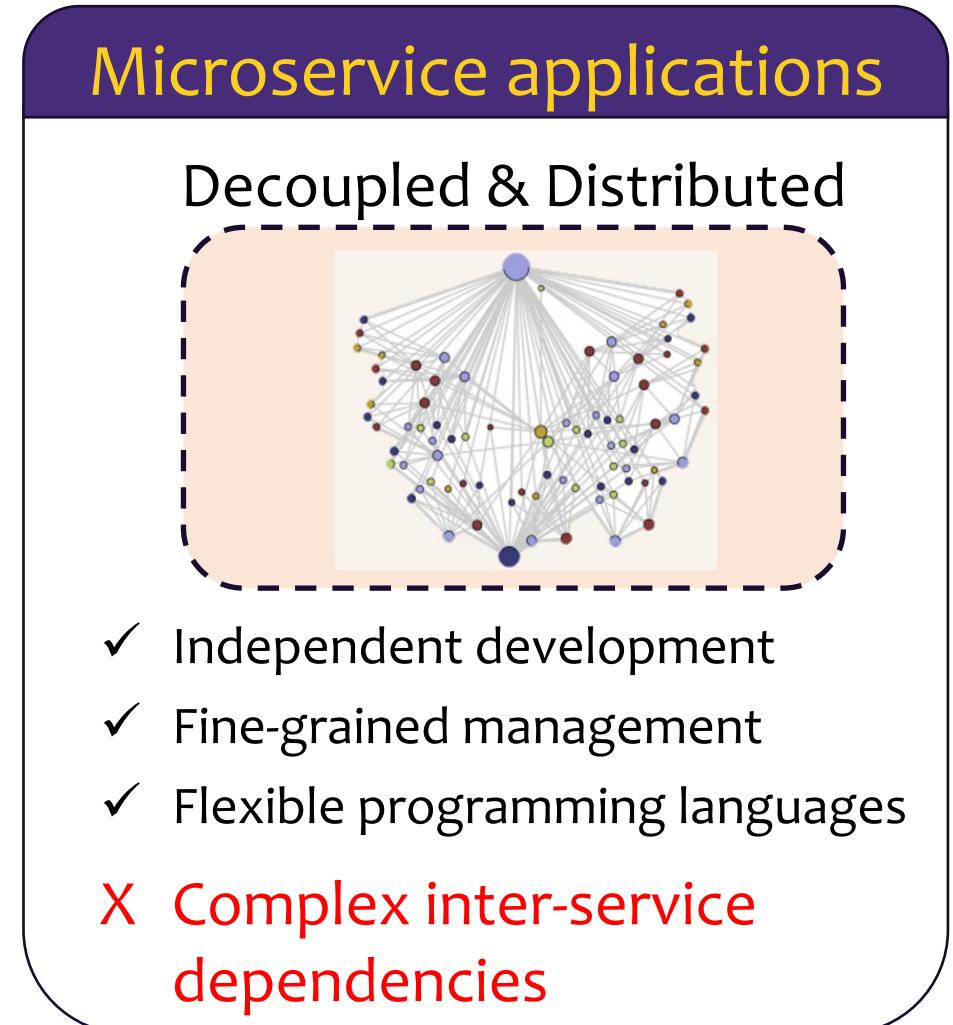
Shifting to



From Monolithic N-tier to Microservices



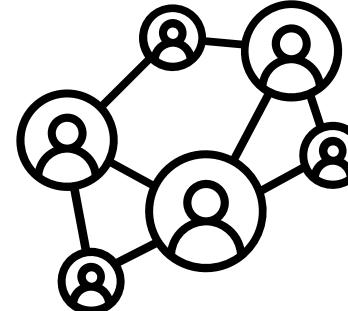
Shifting to



Stringent Latency Requirement of Online Services



E-commerce



Social Media



Online Navigation

Stringent Latency Requirement of Online Services

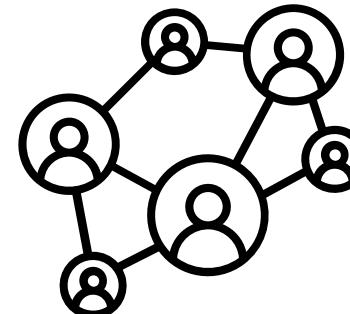
- Amazon: 100ms latency increase → \$0.7 billion loss



-[Kohavi et al. Computer' 07]



E-commerce



Social Media



Online Navigation

Stringent Latency Requirement of Online Services

- **Amazon:** 100ms latency increase → \$0.7 billion loss



-[Kohavi et al. Computer' 07]

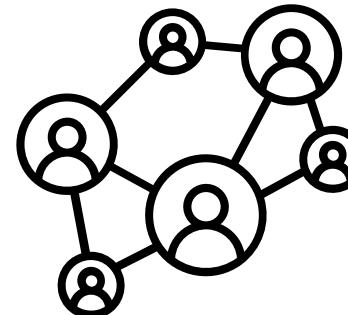
- **Alibaba:** adopting real-time data analytics for Singles' Day advertising → a 30% increase in purchase conversion rate



-[Jiang et al. AliReport'17]



E-commerce



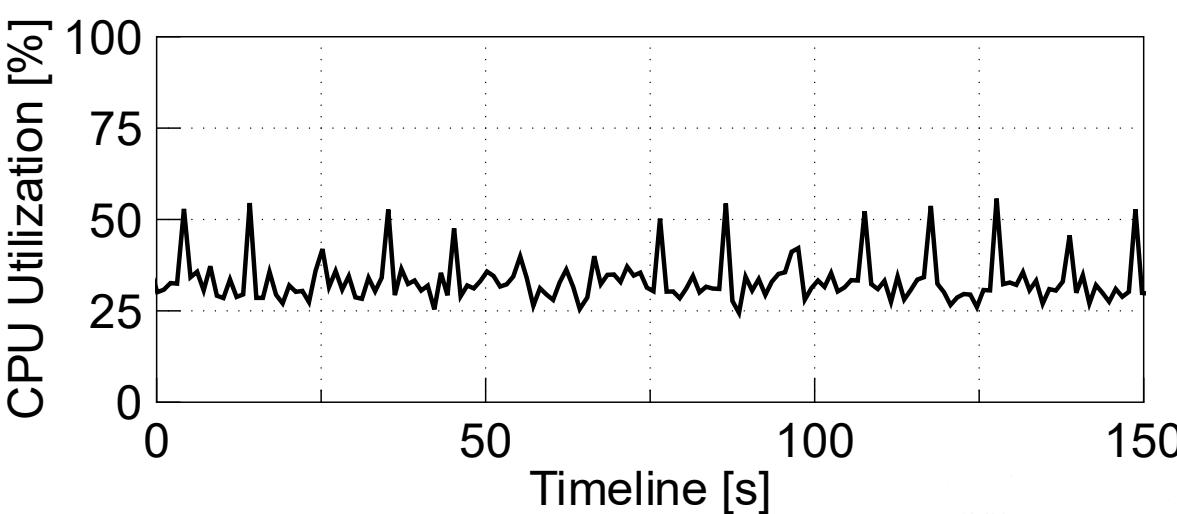
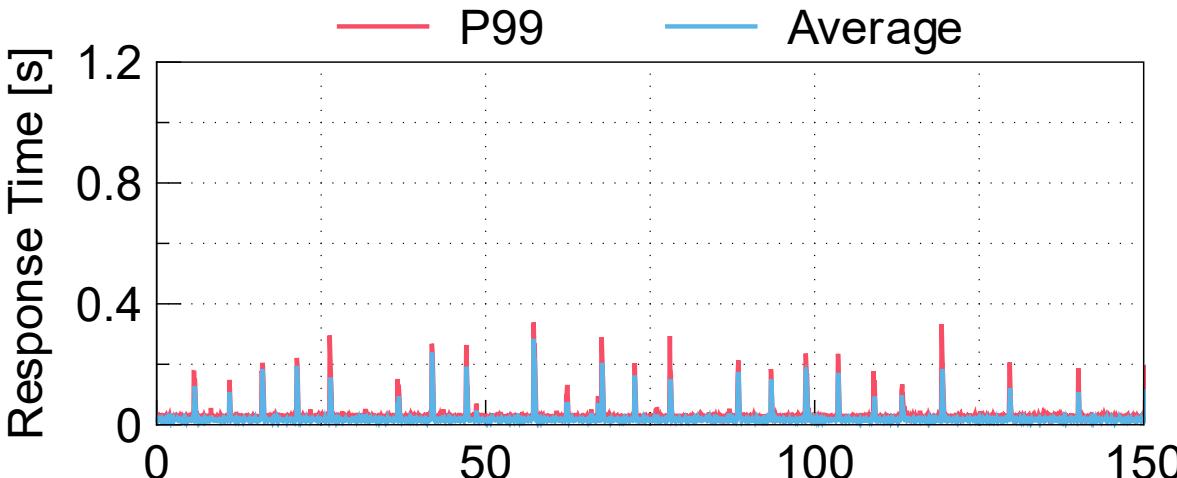
Social Media



Online Navigation

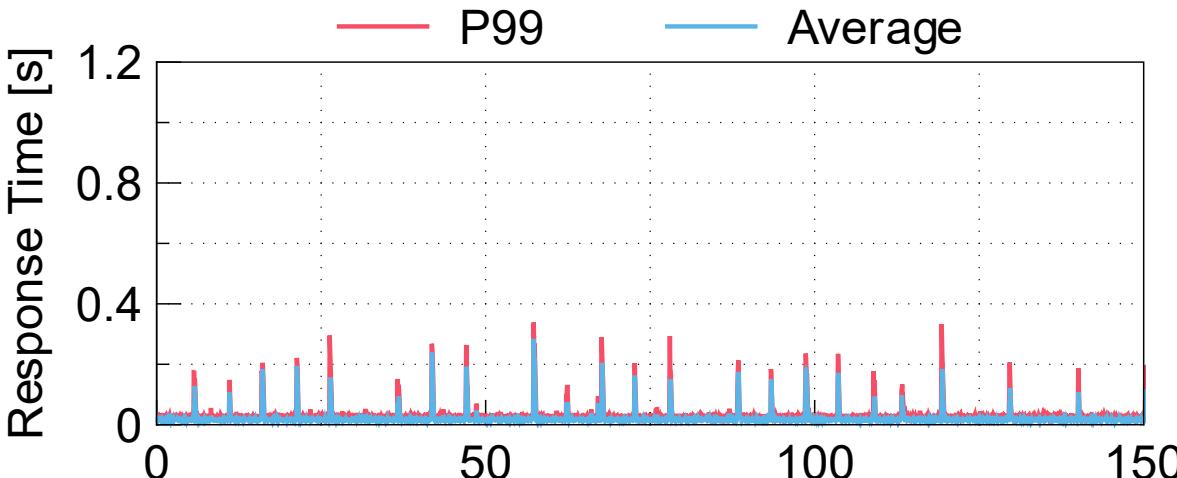
Research Problem

- The challenge: Fragile response time stability in cloud microservices

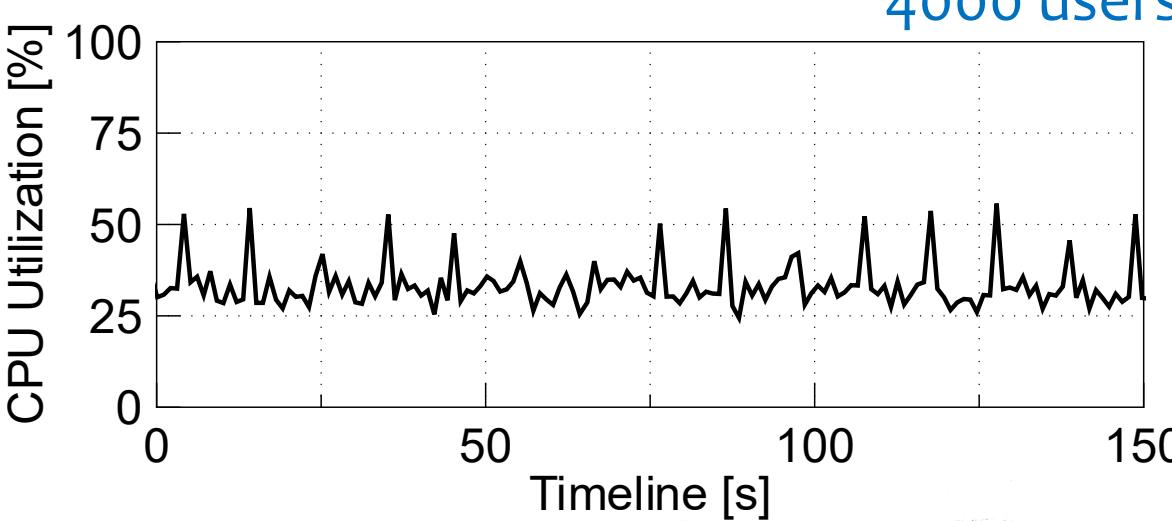


Research Problem

- The challenge: Fragile response time stability in cloud microservices

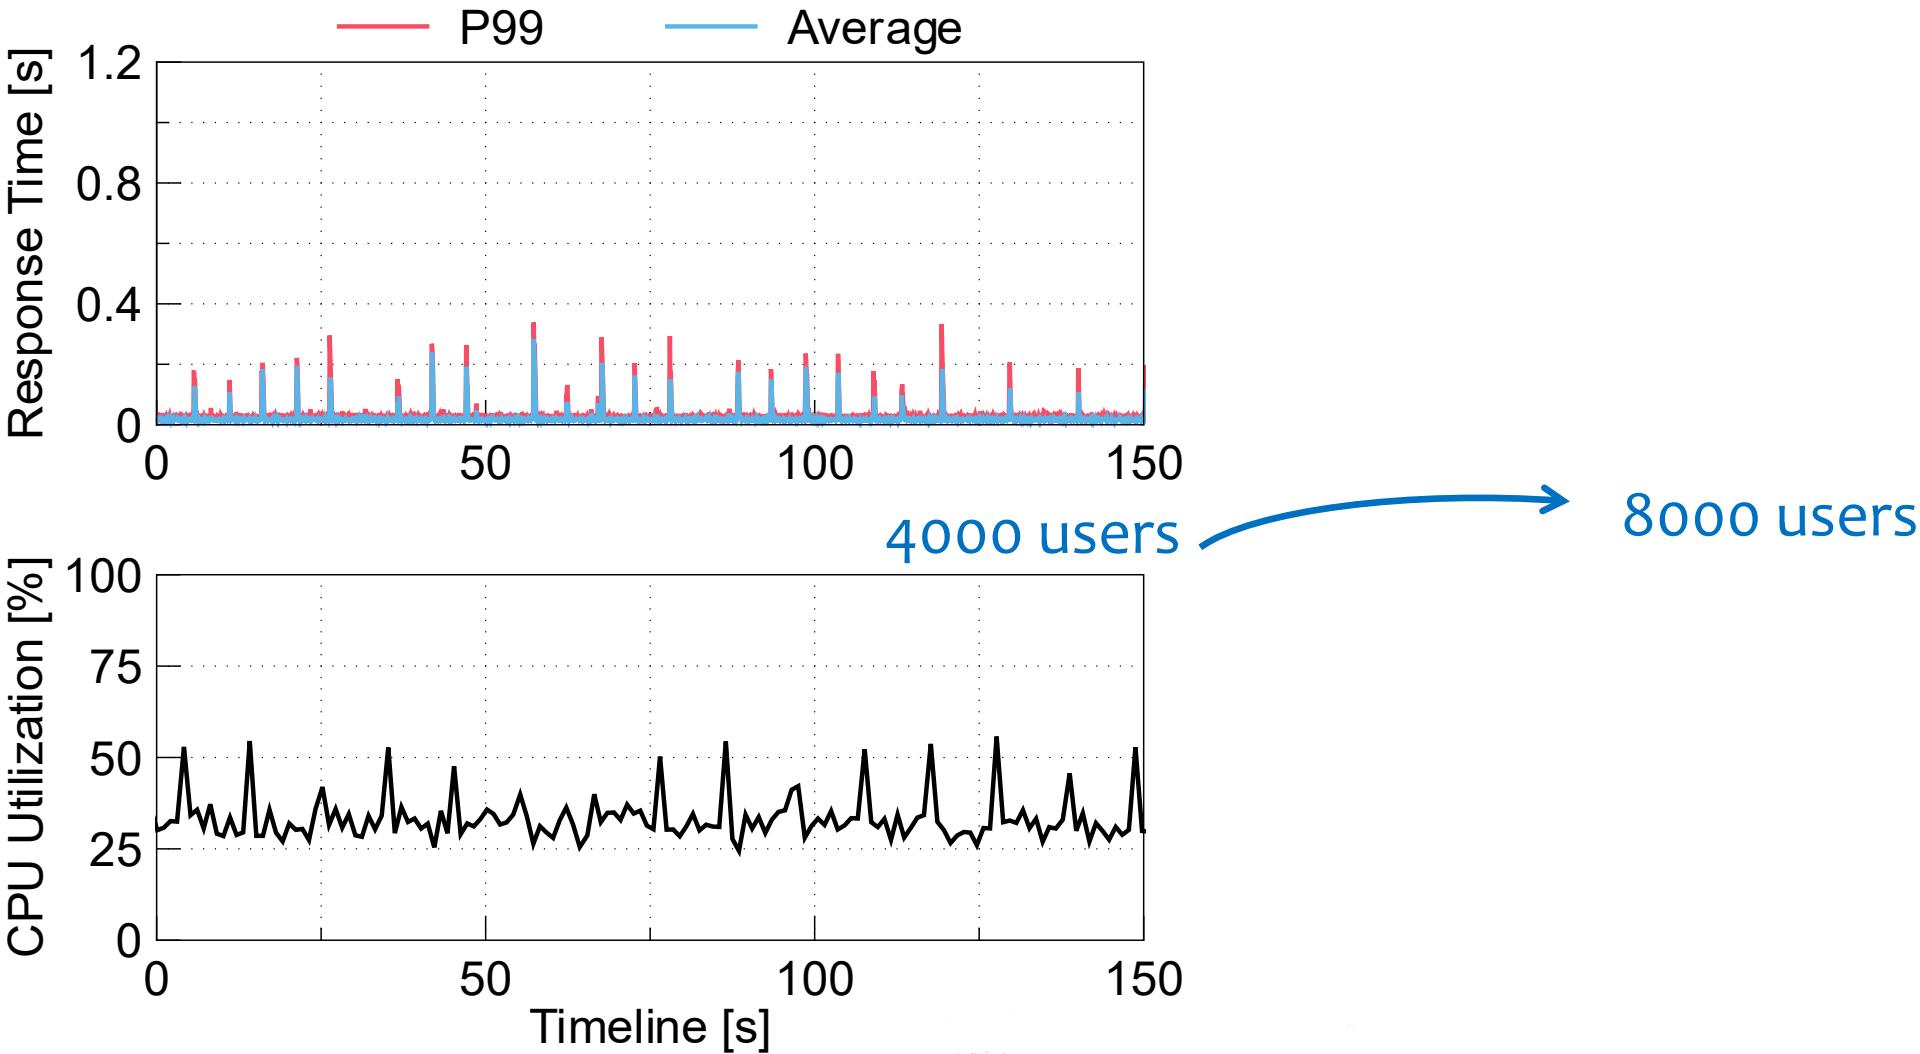


4000 users



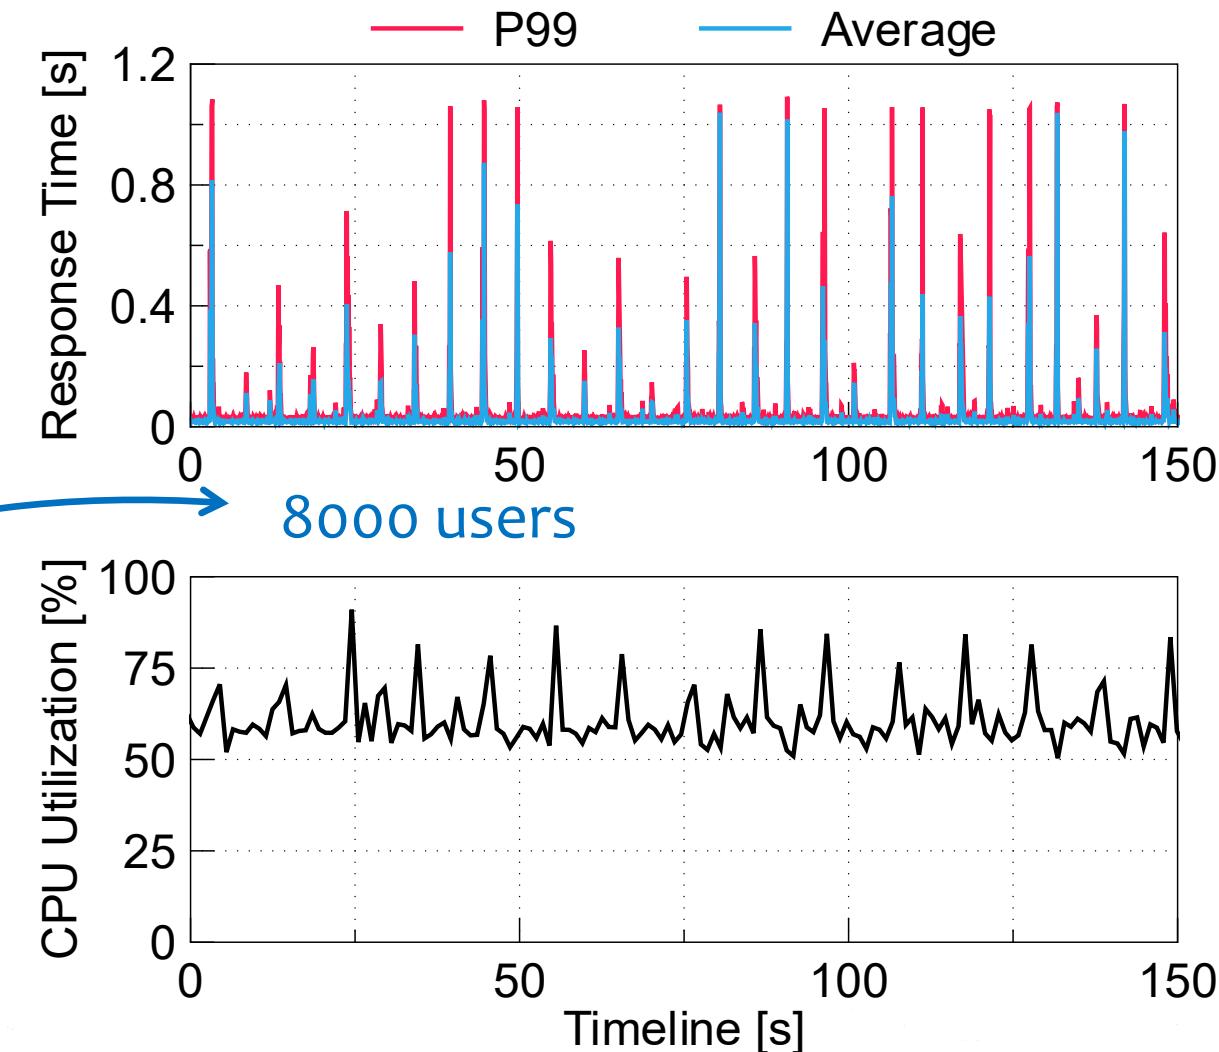
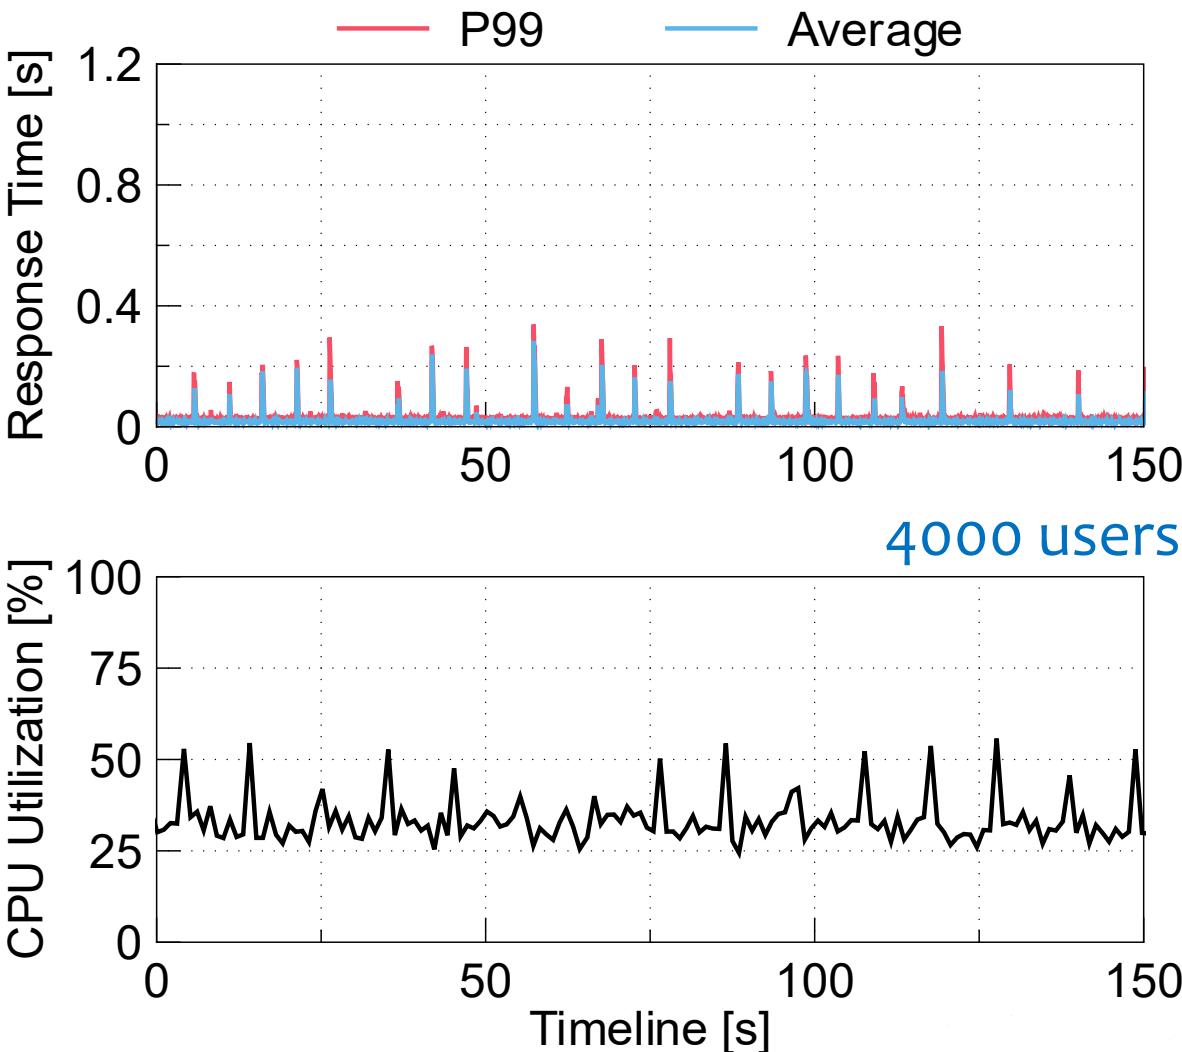
Research Problem

- The challenge: Fragile response time stability in cloud microservices

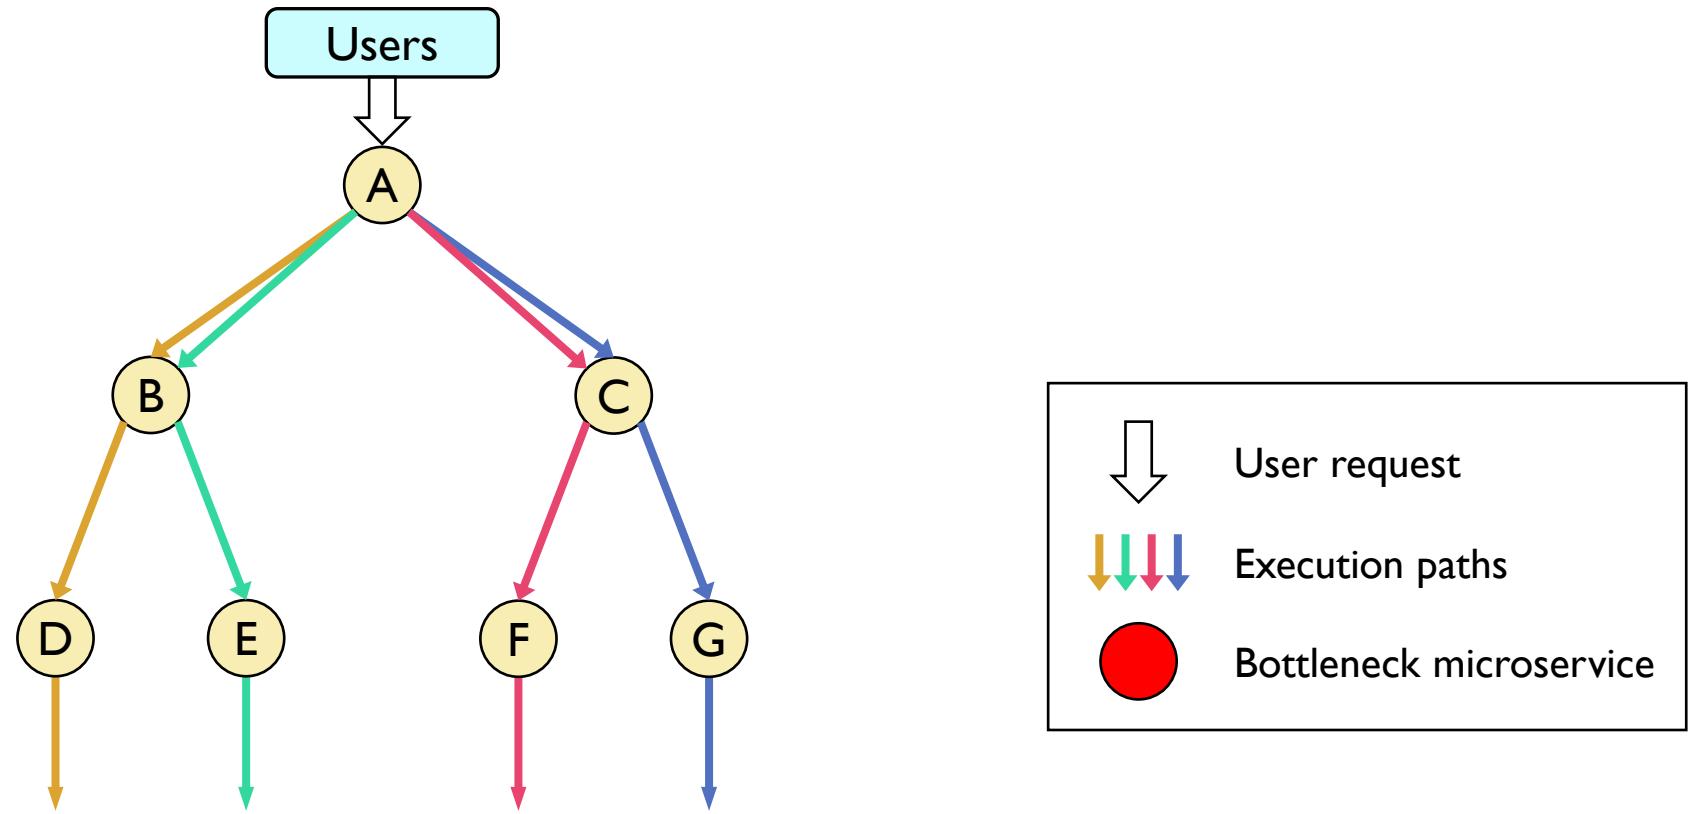


Research Problem

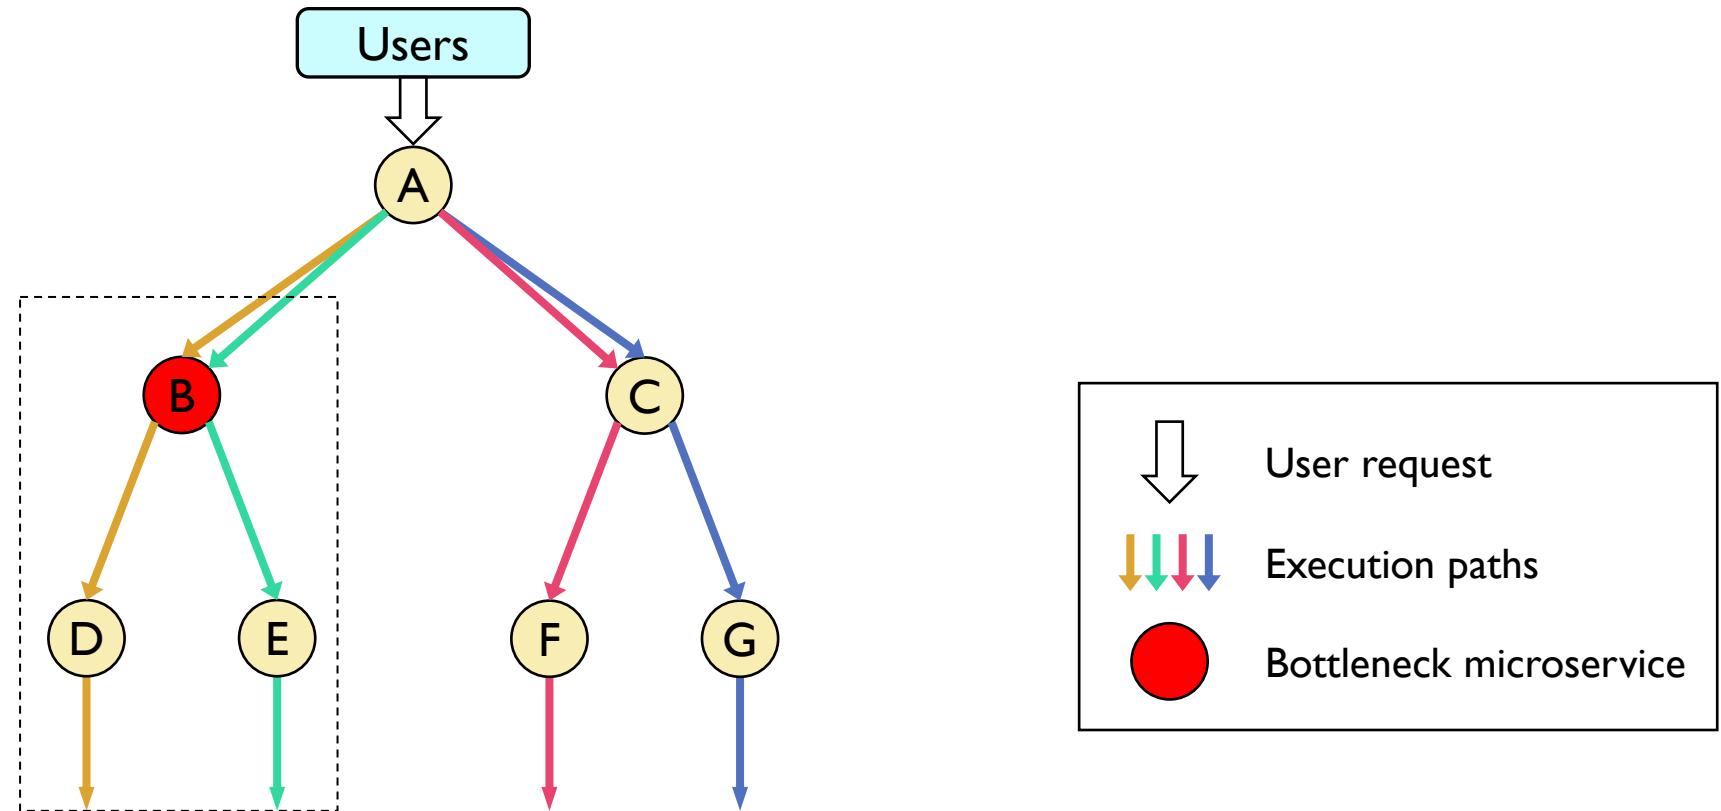
- The challenge: Fragile response time stability in cloud microservices



Cross-Path Dependencies Among Different Execution Paths

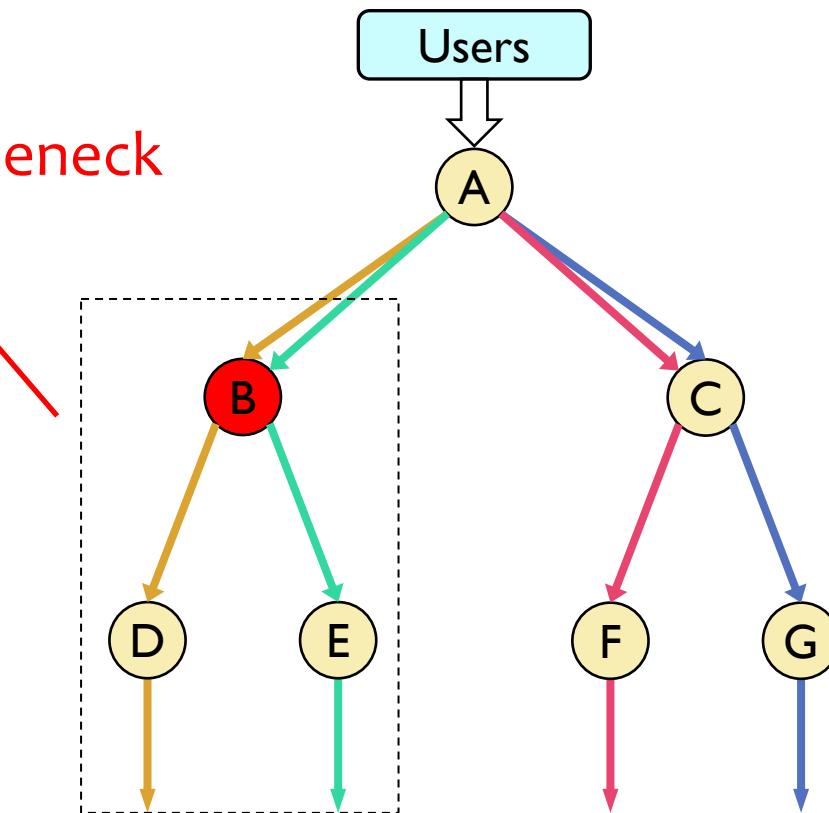


Cross-Path Dependencies Among Different Execution Paths



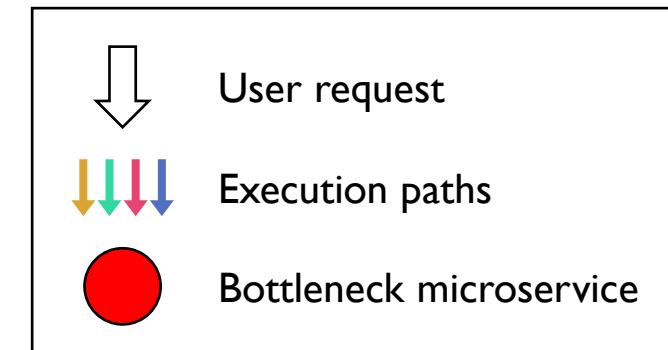
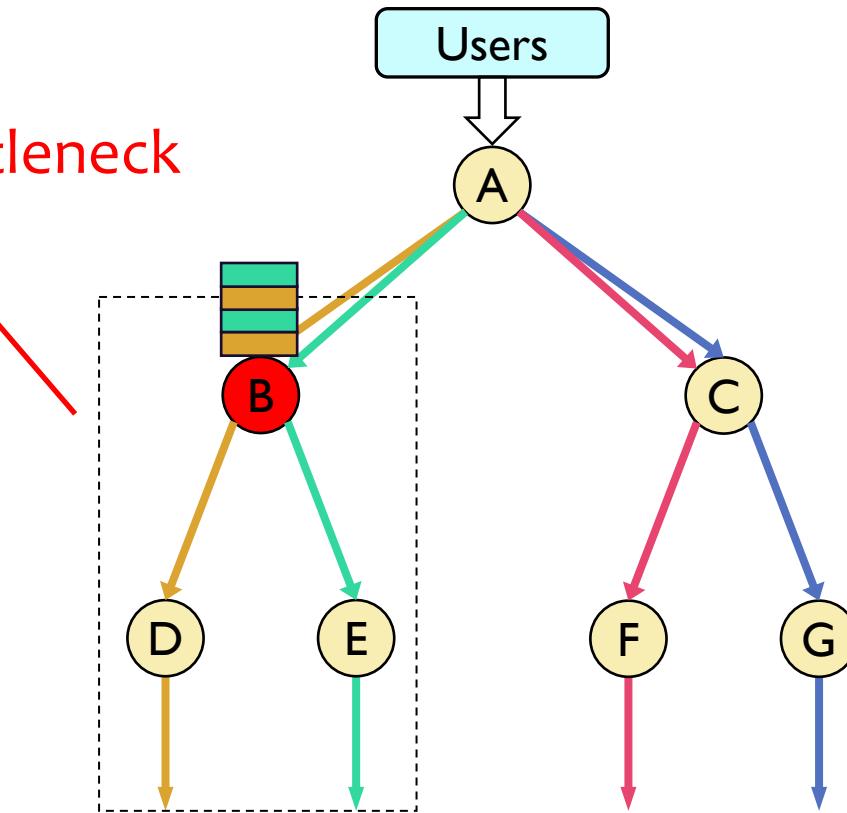
Cross-Path Dependencies Among Different Execution Paths

1. Shared bottleneck microservice



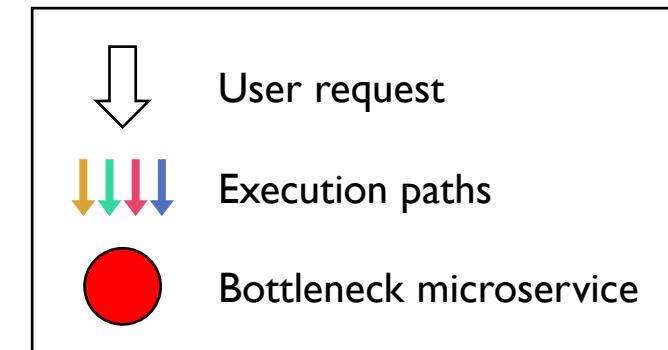
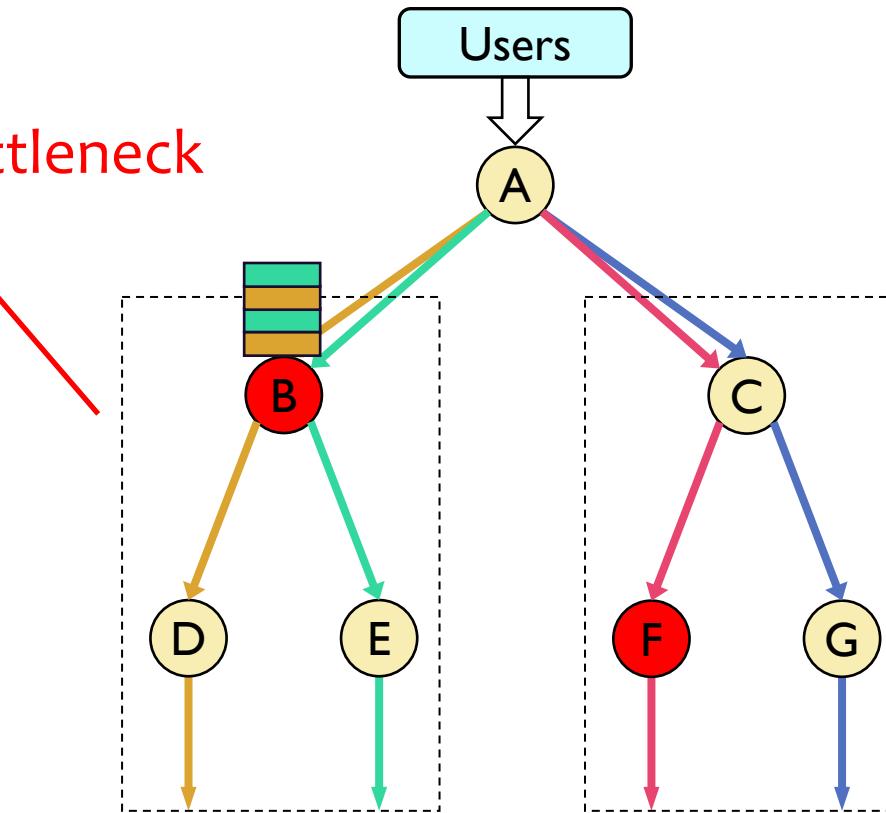
Cross-Path Dependencies Among Different Execution Paths

1. Shared bottleneck microservice



Cross-Path Dependencies Among Different Execution Paths

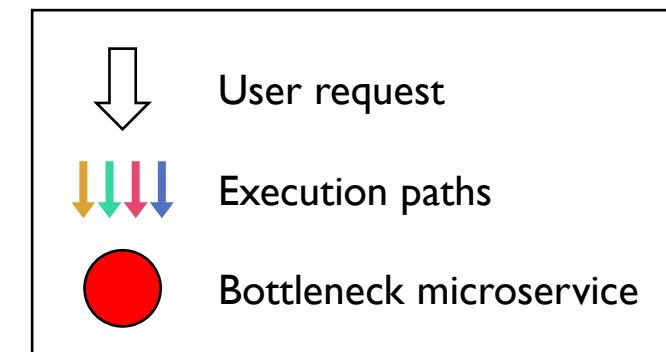
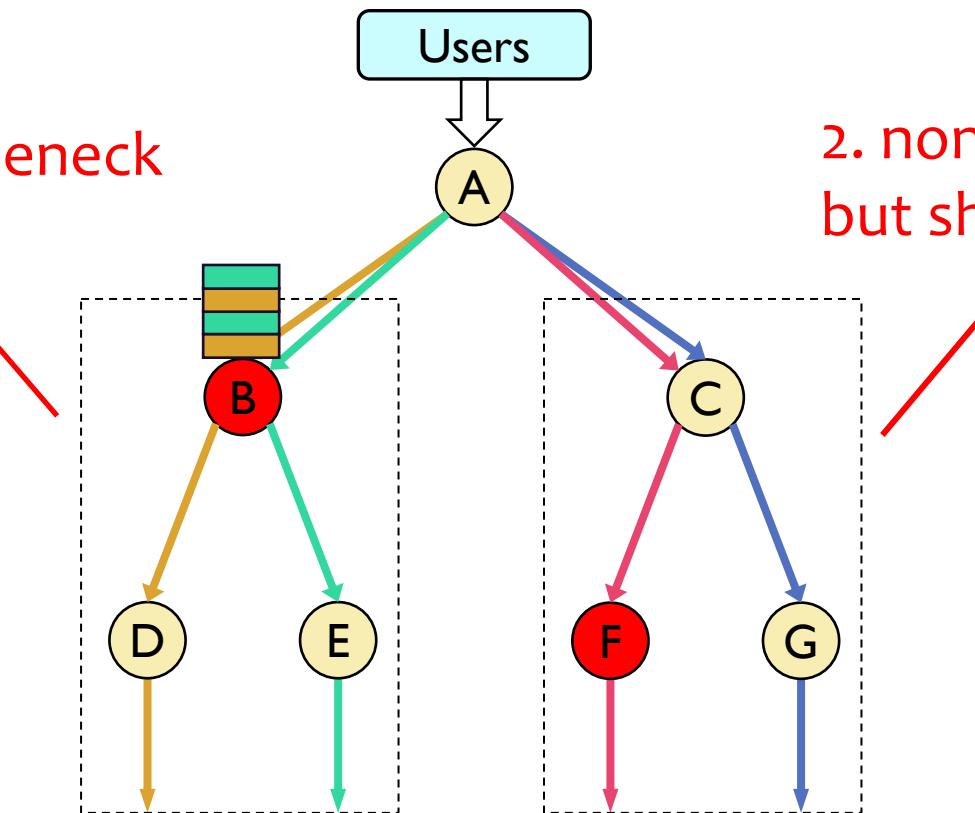
1. Shared bottleneck microservice



Cross-Path Dependencies Among Different Execution Paths

1. Shared bottleneck microservice

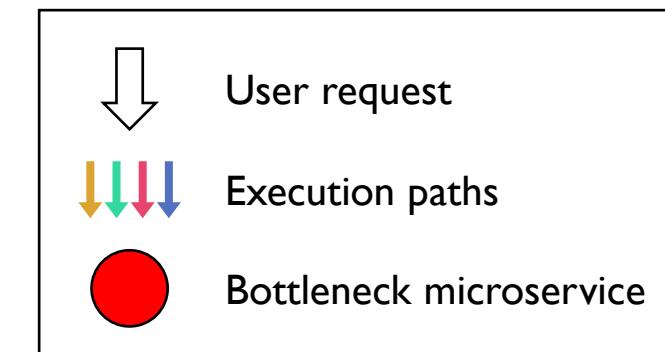
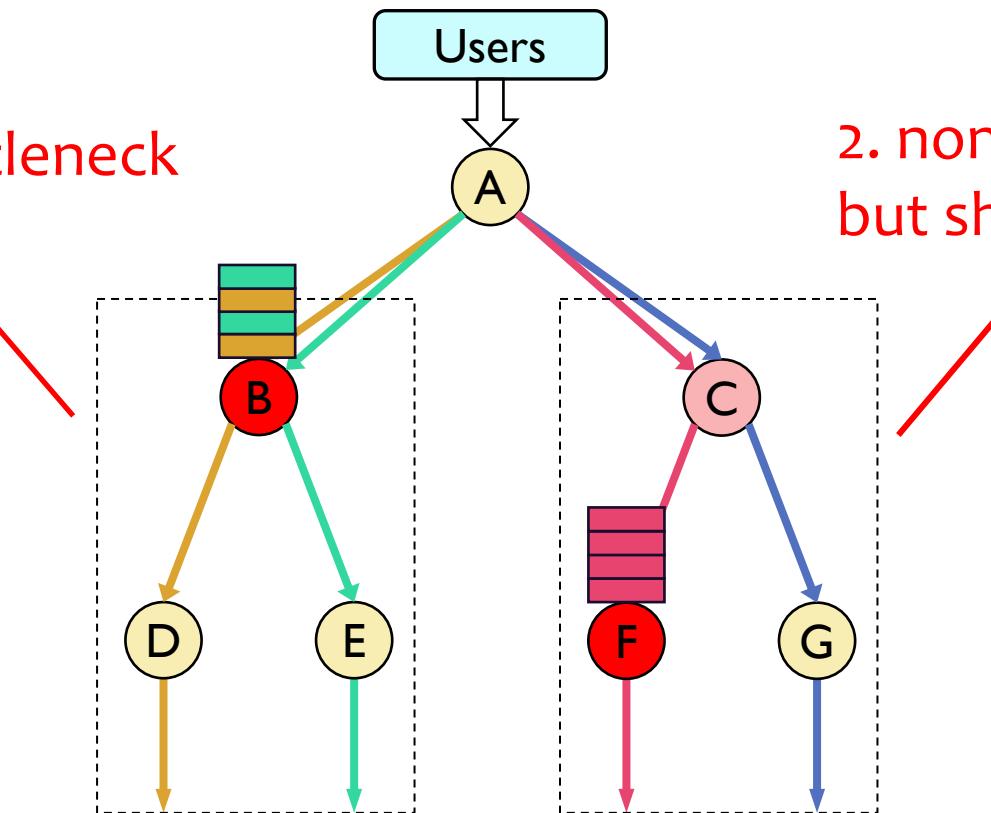
2. non-shared bottleneck microservice, but shared upstream service



Cross-Path Dependencies Among Different Execution Paths

1. Shared bottleneck microservice

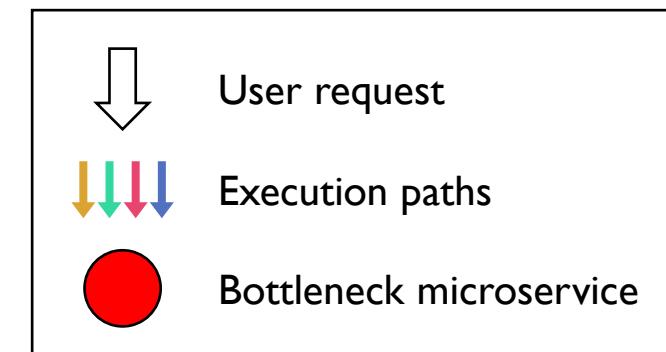
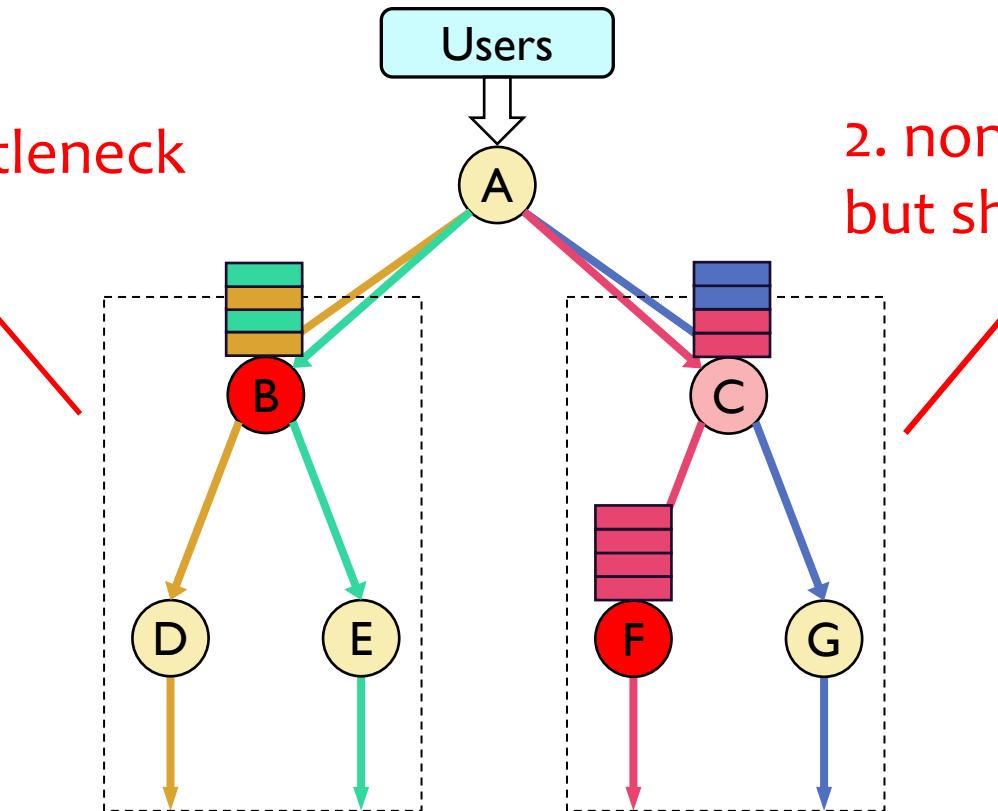
2. non-shared bottleneck microservice, but shared upstream service



Cross-Path Dependencies Among Different Execution Paths

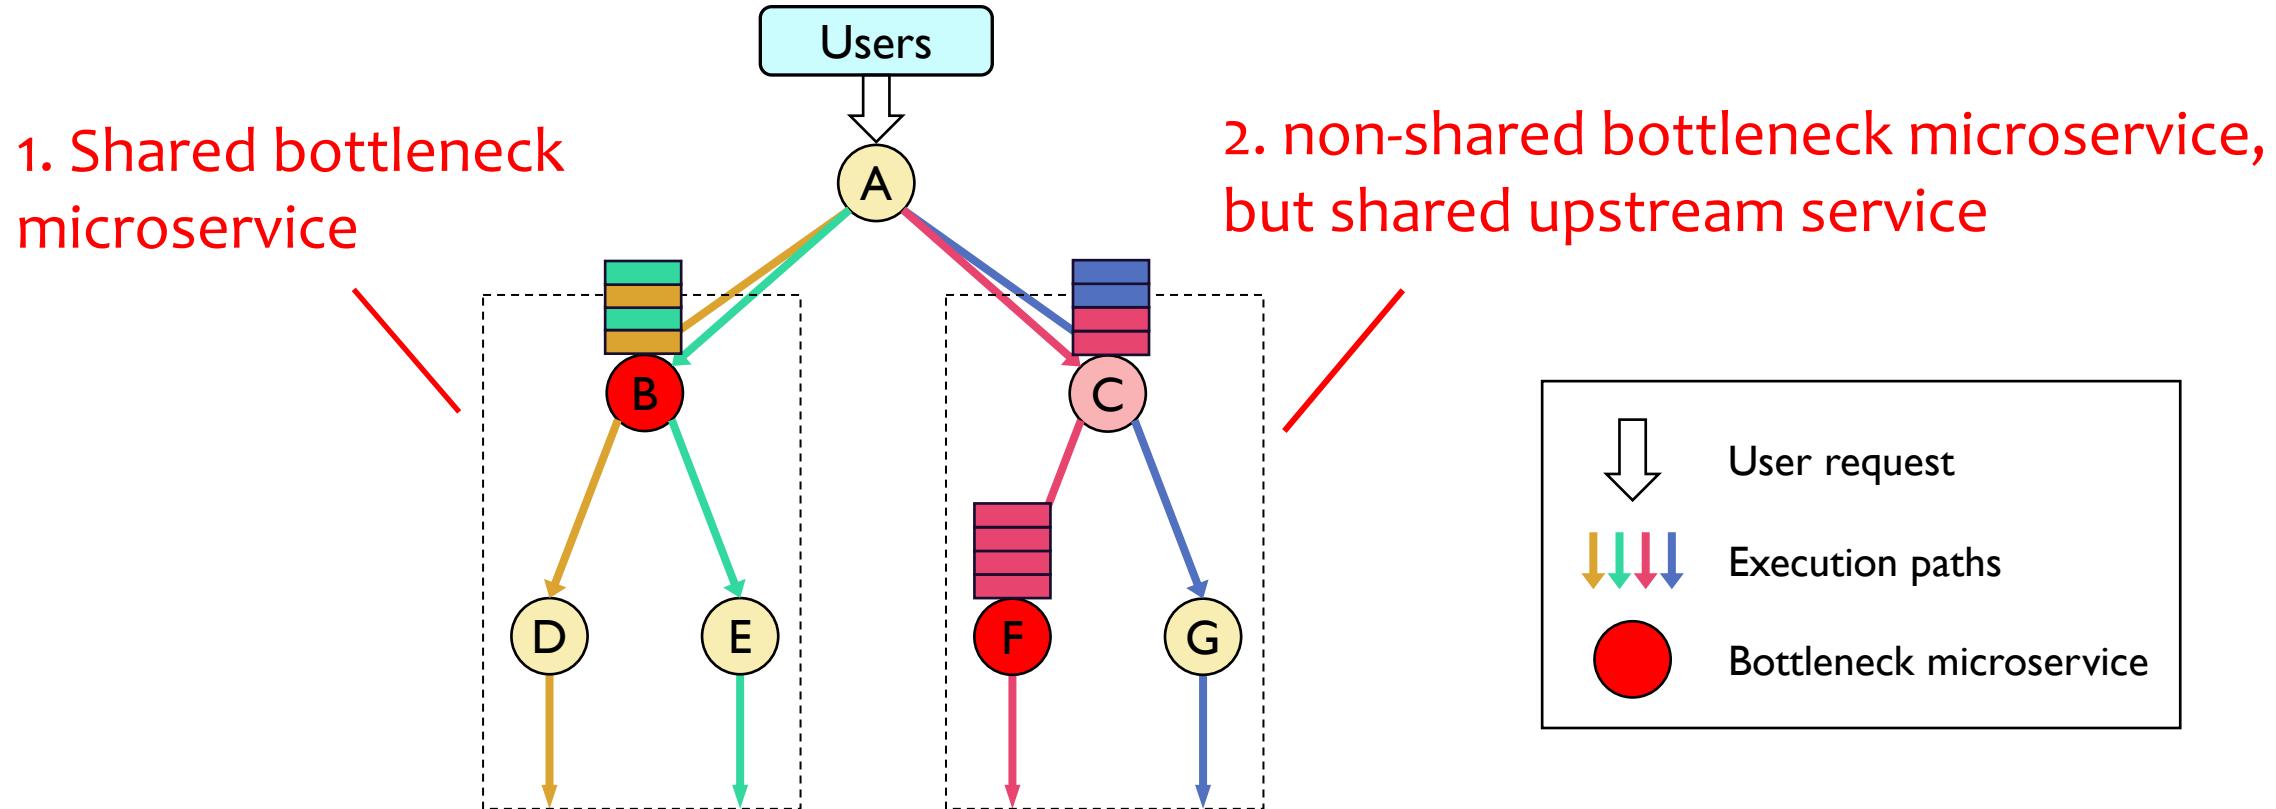
1. Shared bottleneck microservice

2. non-shared bottleneck microservice, but shared upstream service



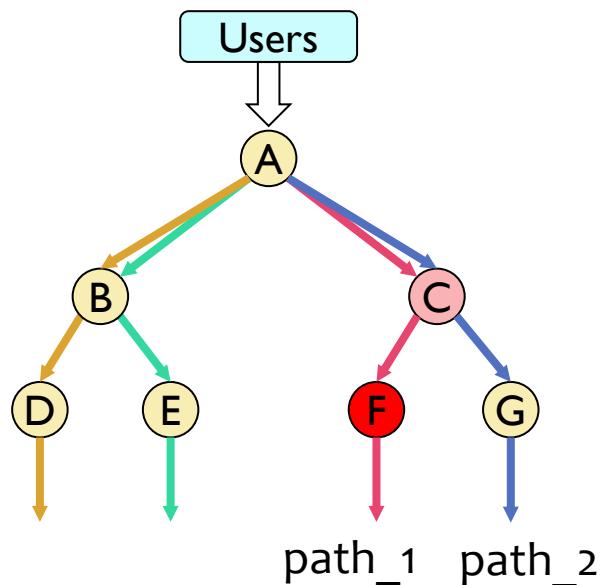
Cross-Path Dependencies Among Different Execution Paths

- Two typical cross-path dependencies



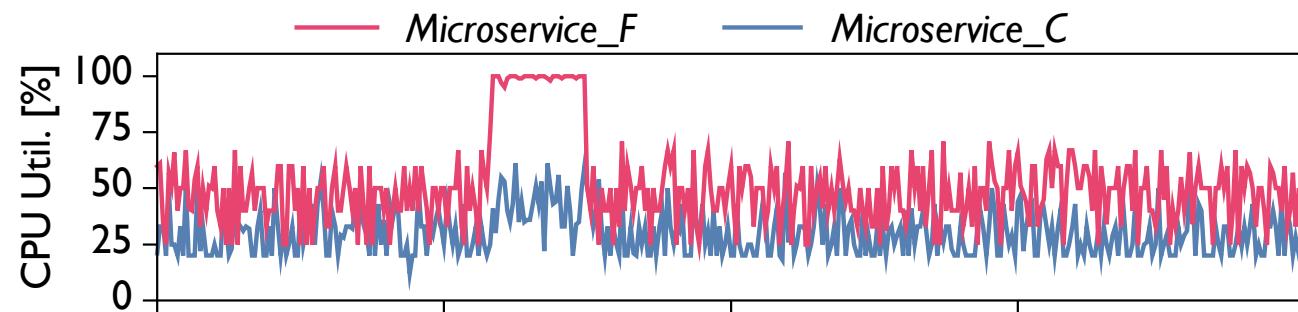
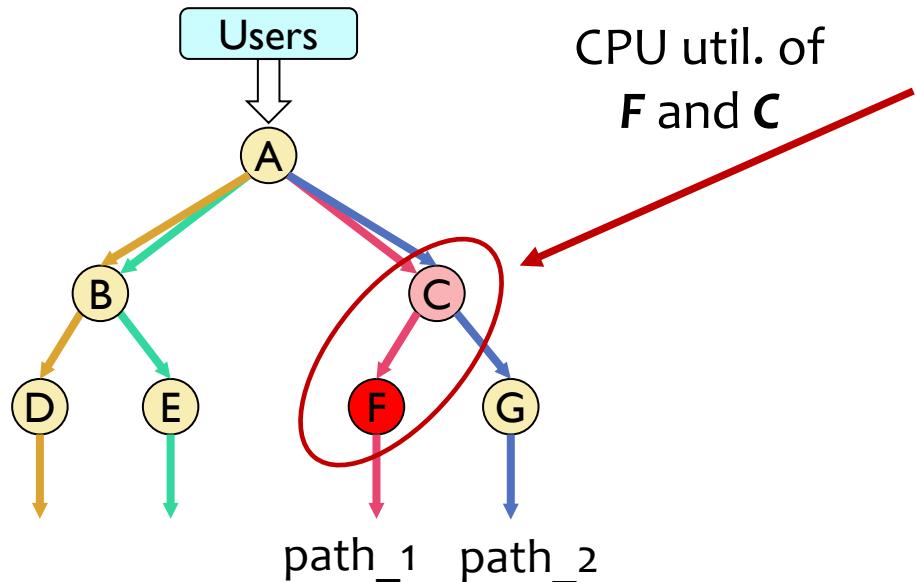
Cross-Path Dependencies → Long Response Time

- Cross-path dependencies exist due to shared microservice
 - bottleneck on **one path** → performance interference to **many other paths**



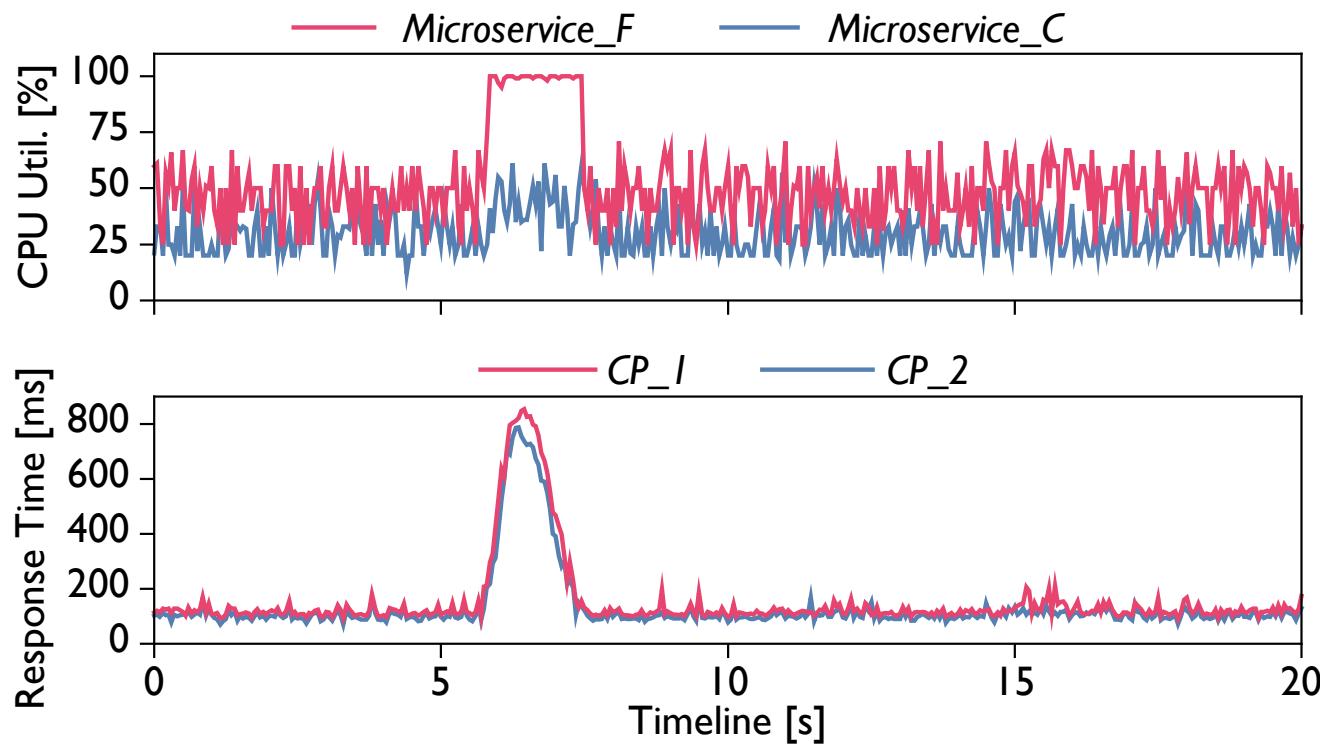
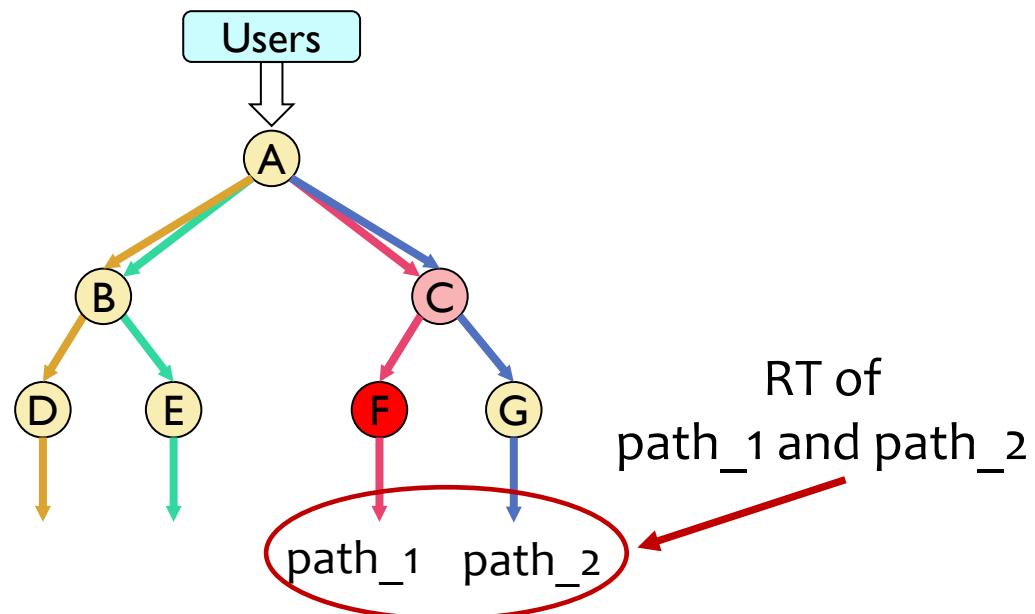
Cross-Path Dependencies → Long Response Time

- Cross-path dependencies exist due to shared microservice
 - bottleneck on one path → performance interference to many other paths



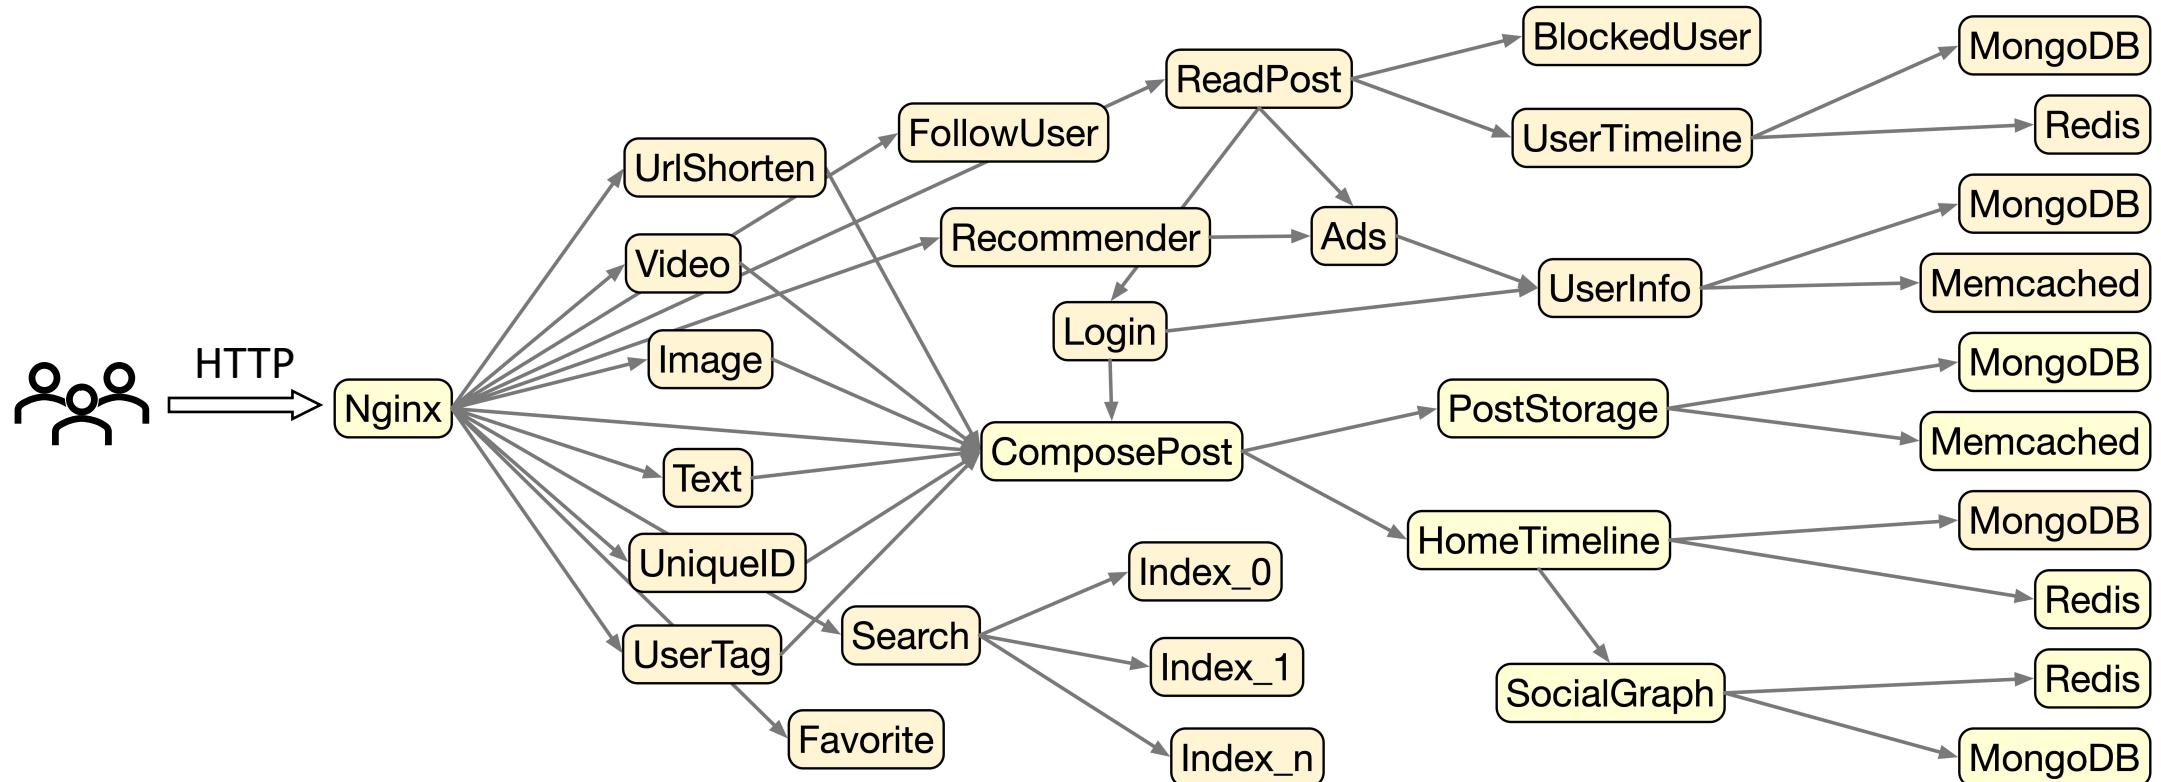
Cross-Path Dependencies → Long Response Time

- Cross-path dependencies exist due to shared microservice
 - bottleneck on one path → performance interference to many other paths



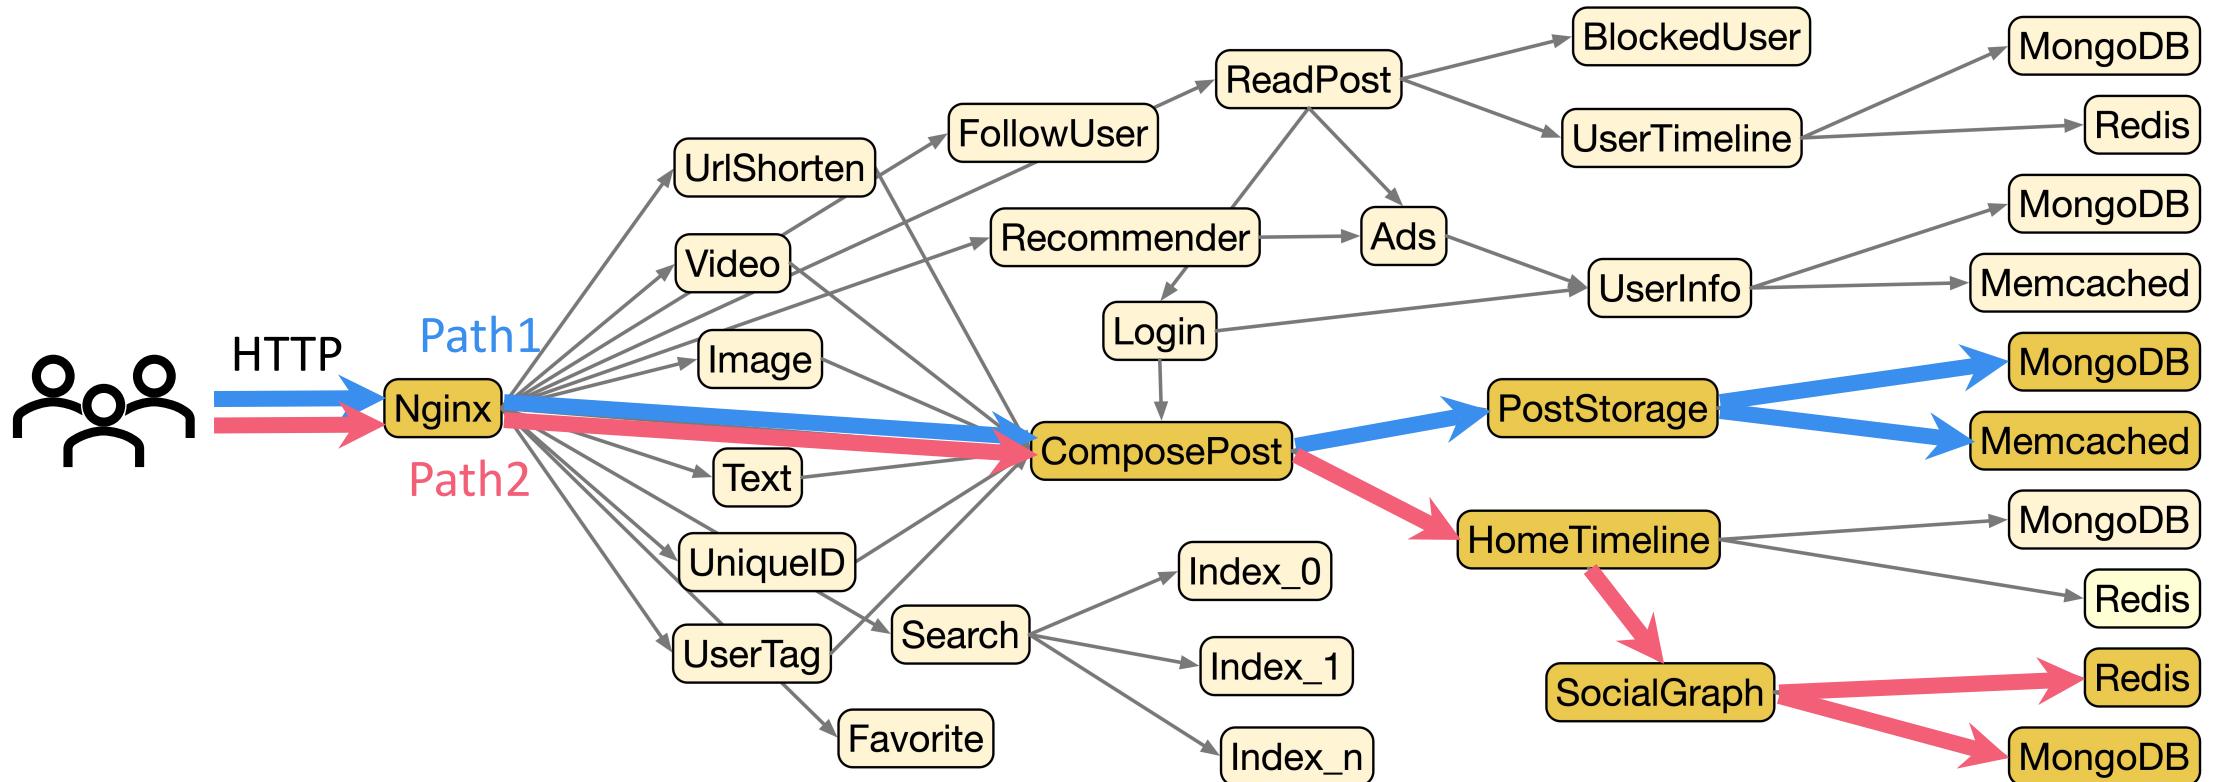
Case Study: Experimental Setup

- ❑ SocialNetwork microservices – [Gan et al. ASPLOS' 19]
- ❑ 36 unique component microservices
- ❑ More than 20 different execution paths supported



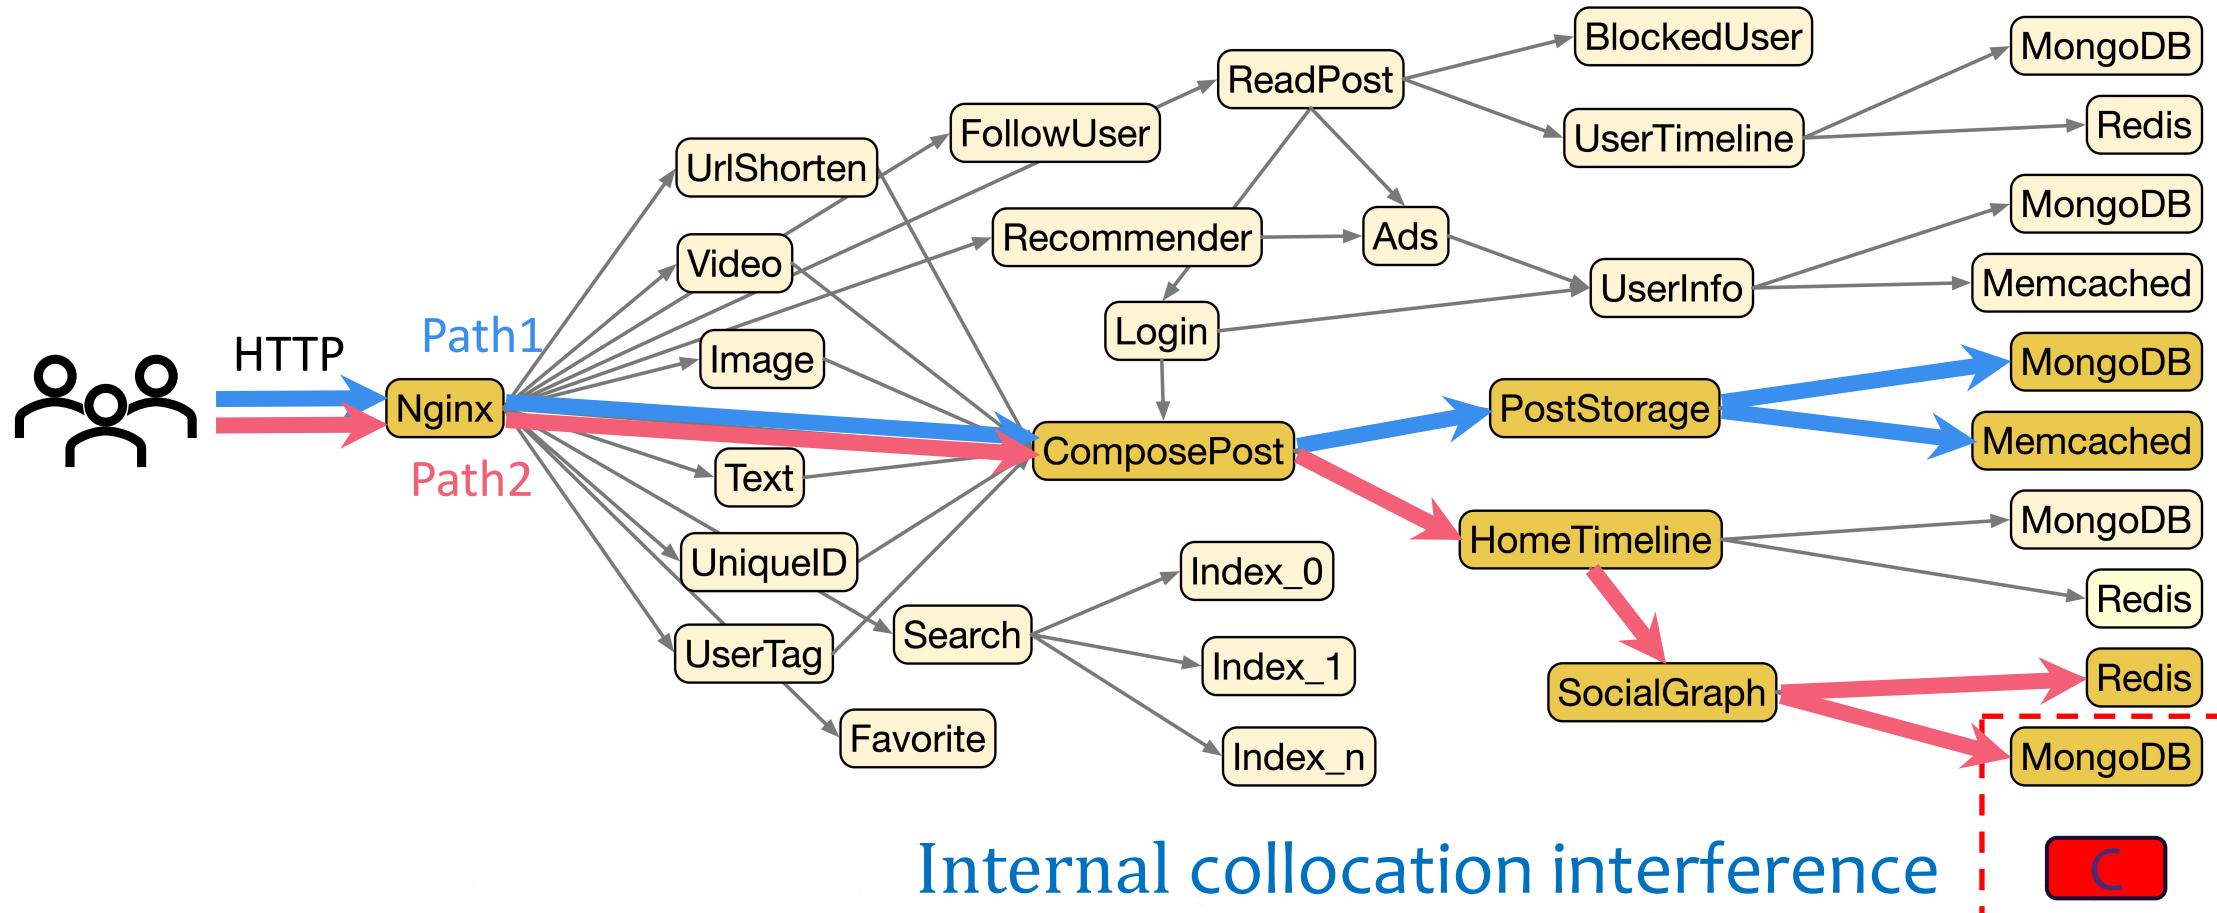
Case Study: Bottleneck Caused by Noisy Neighbor

- Two representative execution paths
- Collocated CPU-intensive neighbor



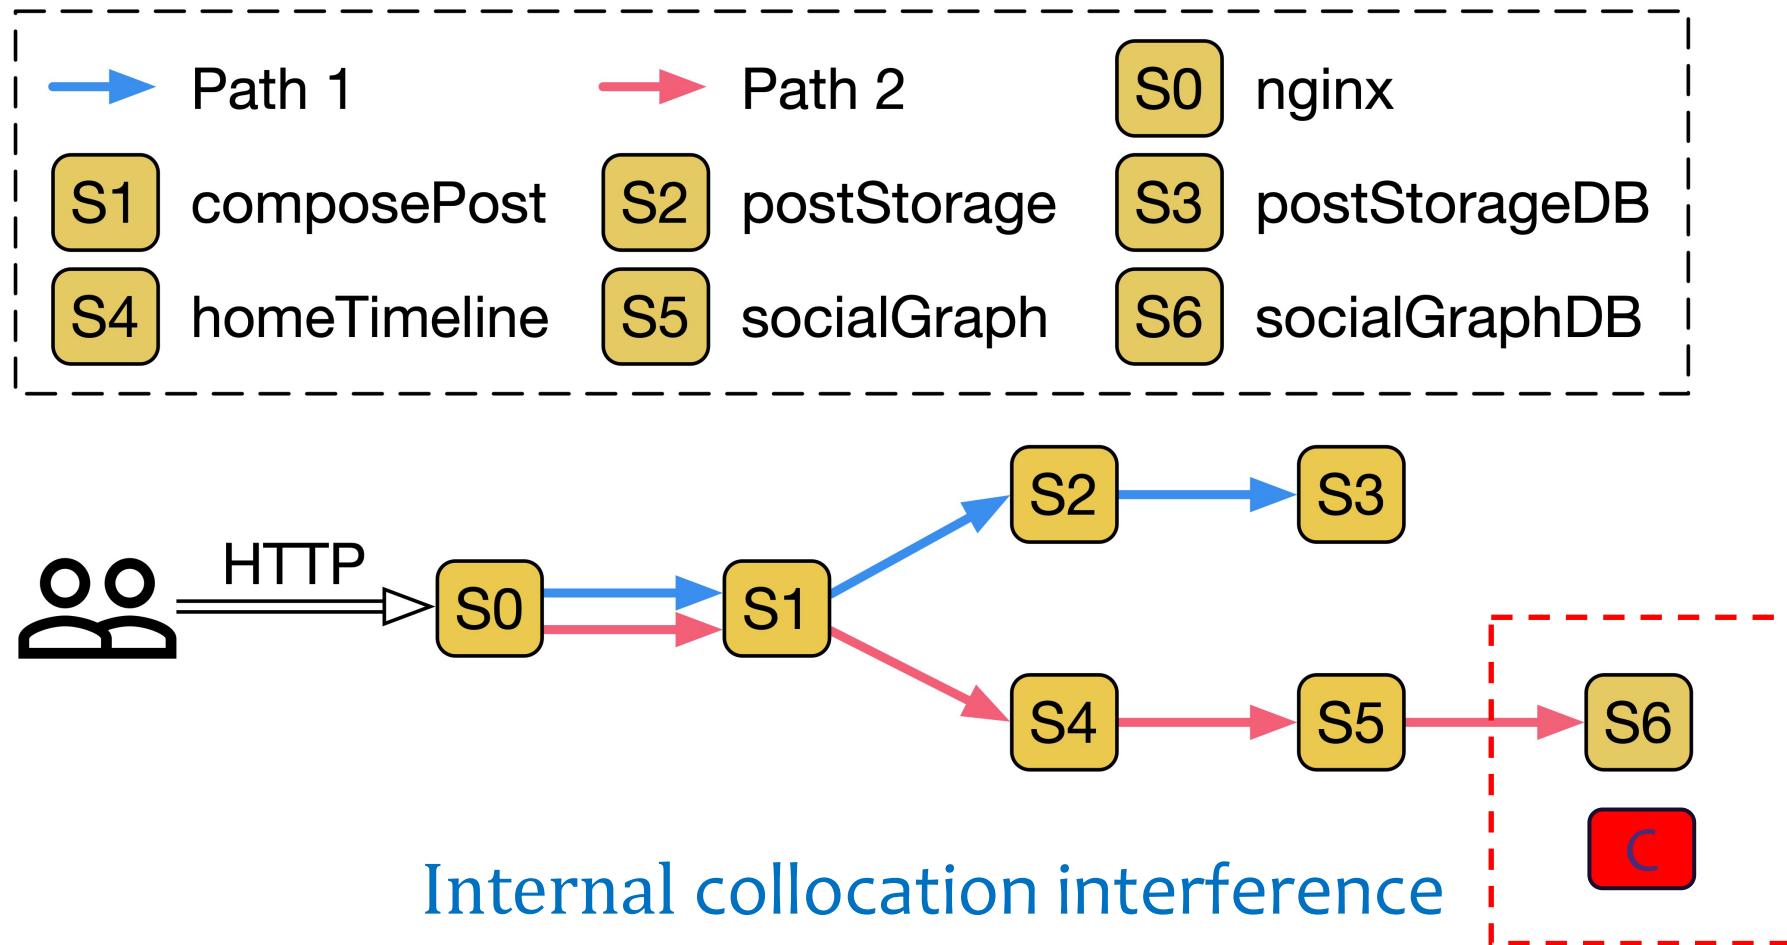
Case Study: Bottleneck Caused by Noisy Neighbor

- Two representative execution paths
- Collocated CPU-intensive neighbor

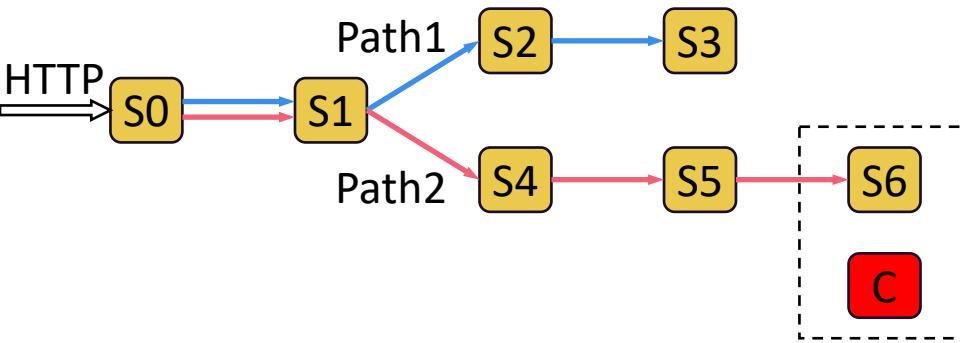


Case Study: Bottleneck Caused by Noisy Neighbor

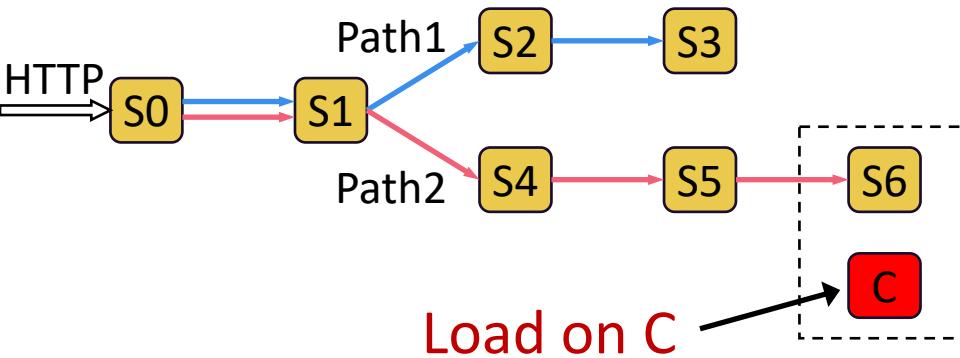
- Two representative execution paths
- Collocated CPU-intensive neighbor



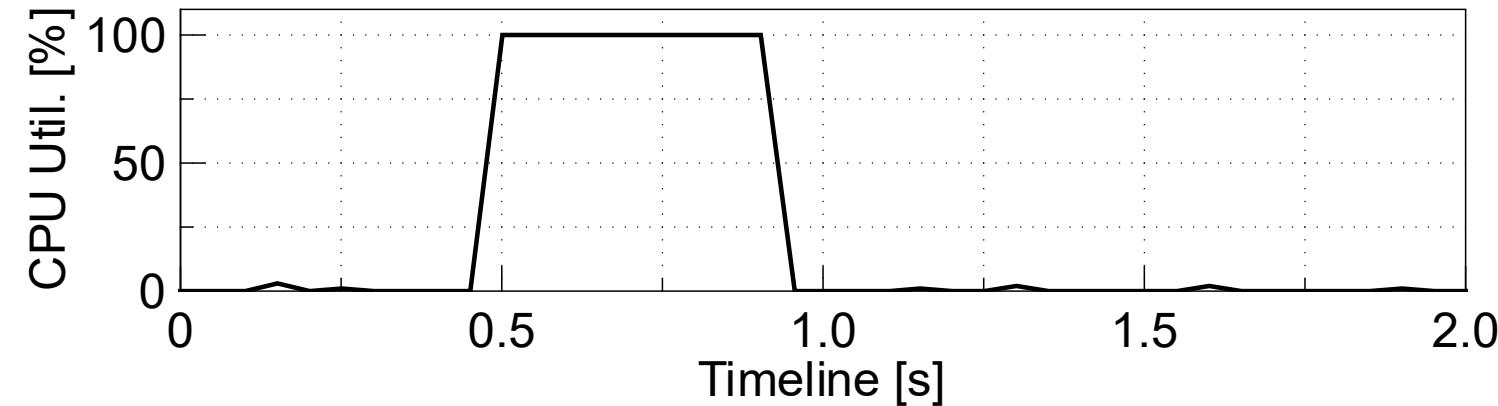
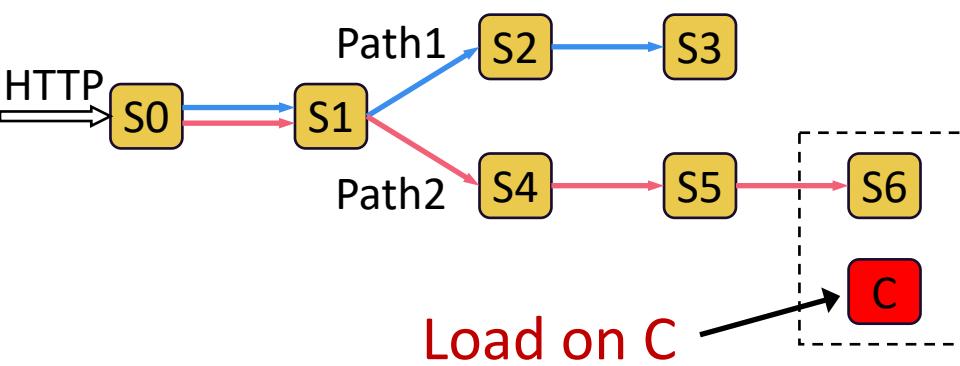
Noisy Neighbor Interference → CPU Saturation on S6



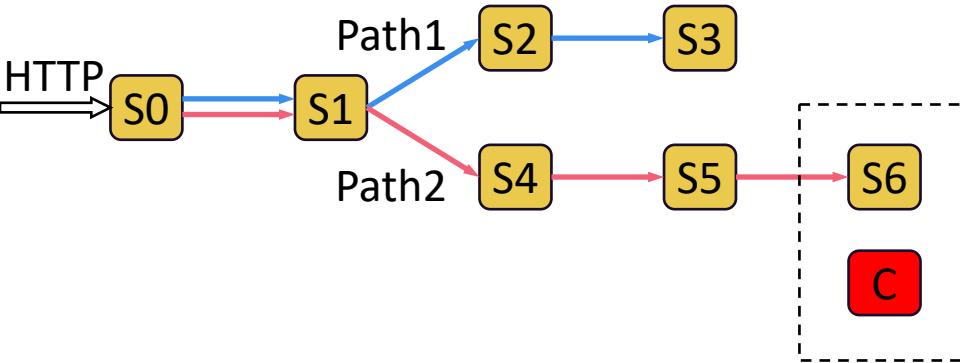
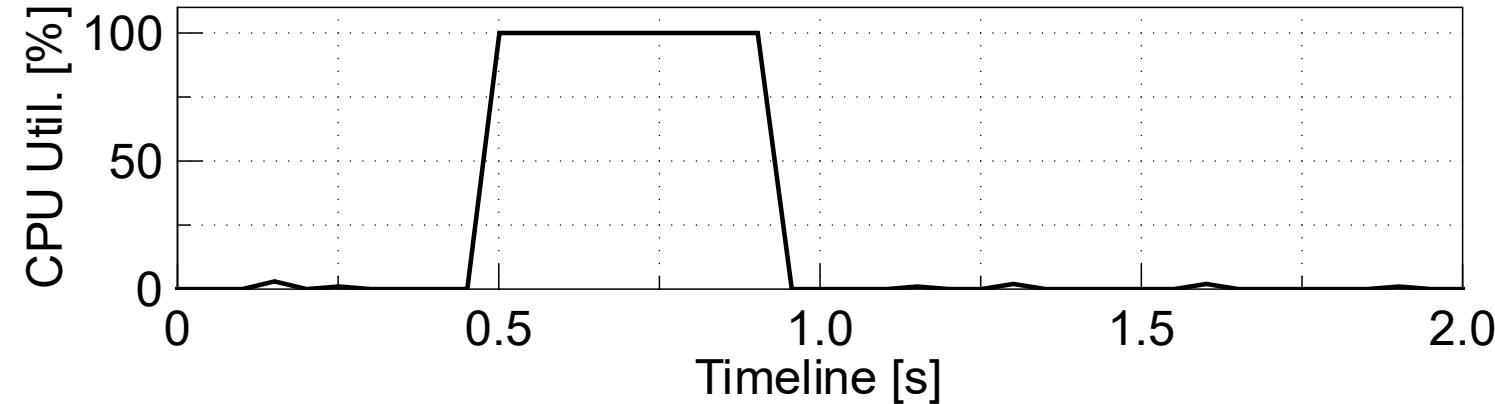
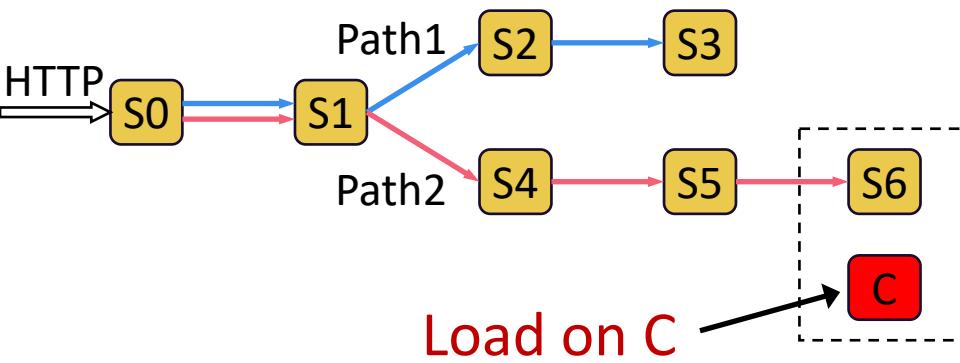
Noisy Neighbor Interference → CPU Saturation on S6



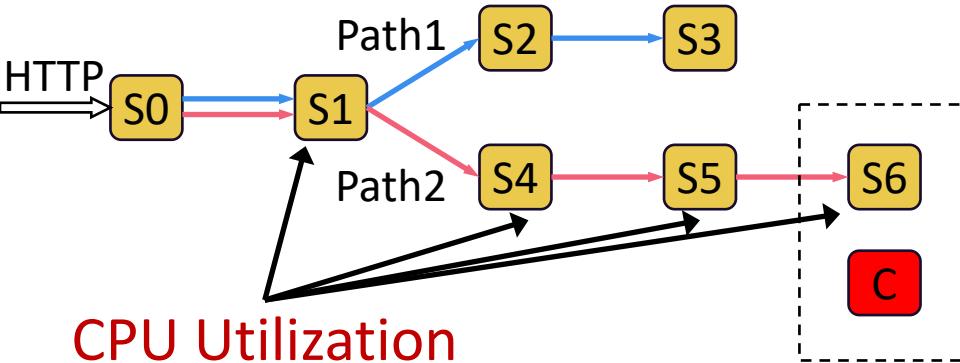
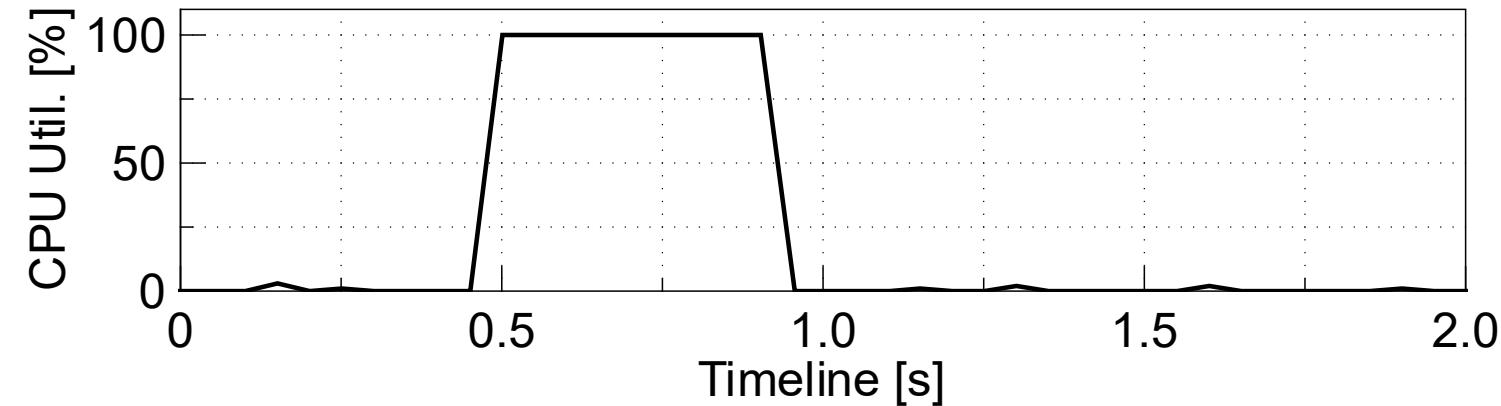
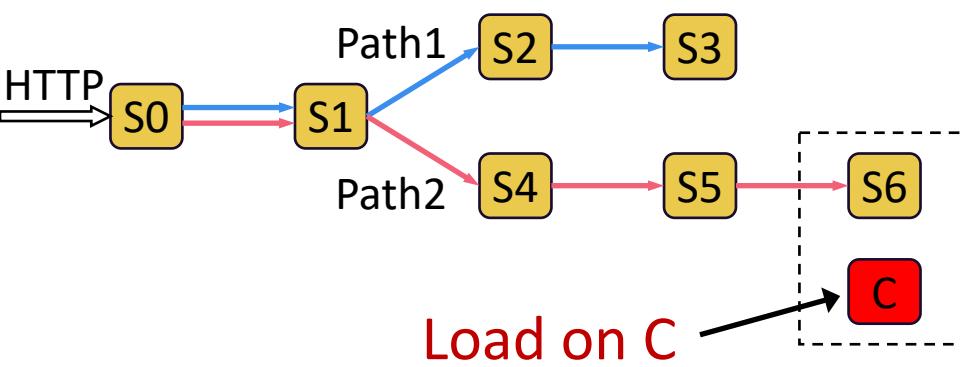
Noisy Neighbor Interference → CPU Saturation on S6



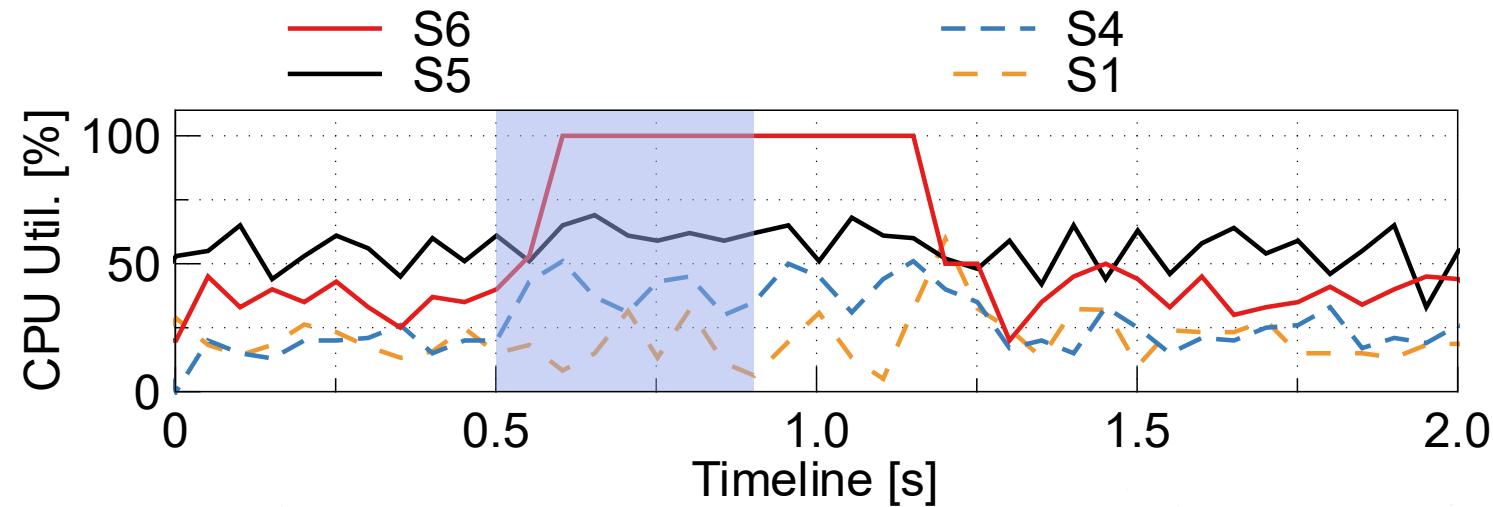
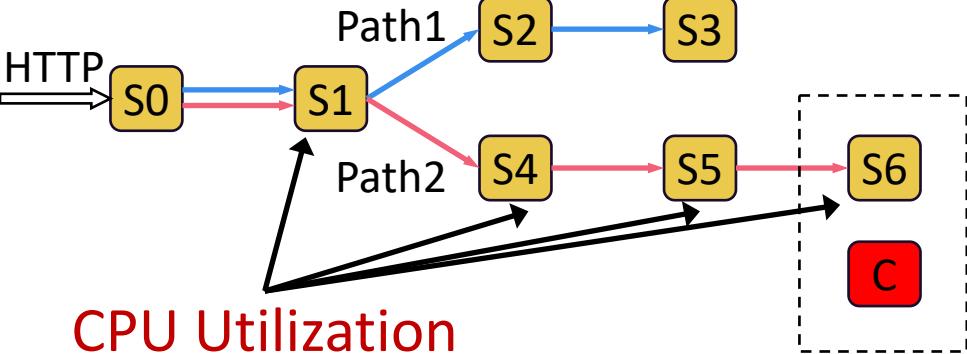
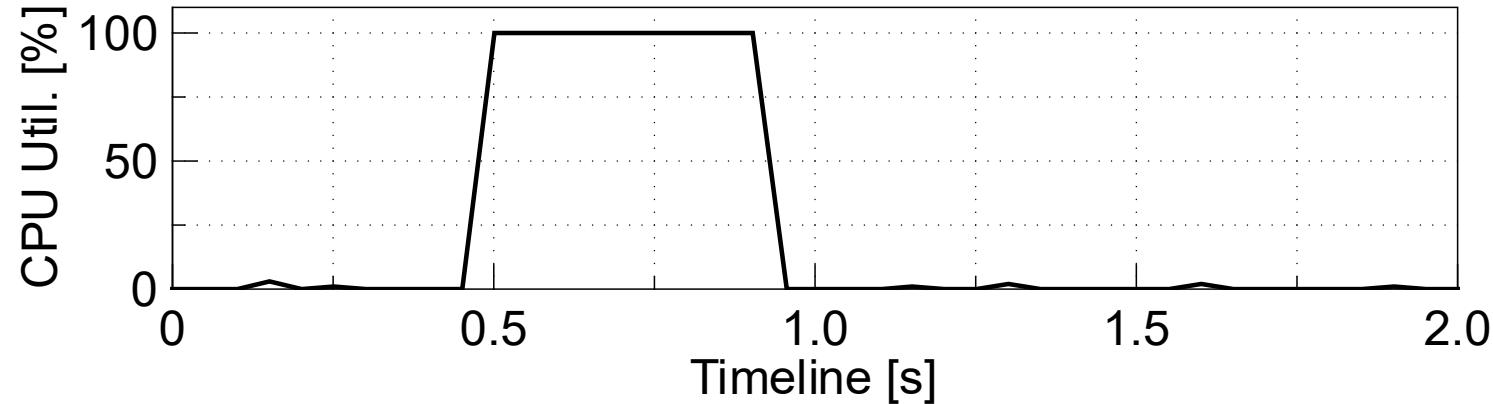
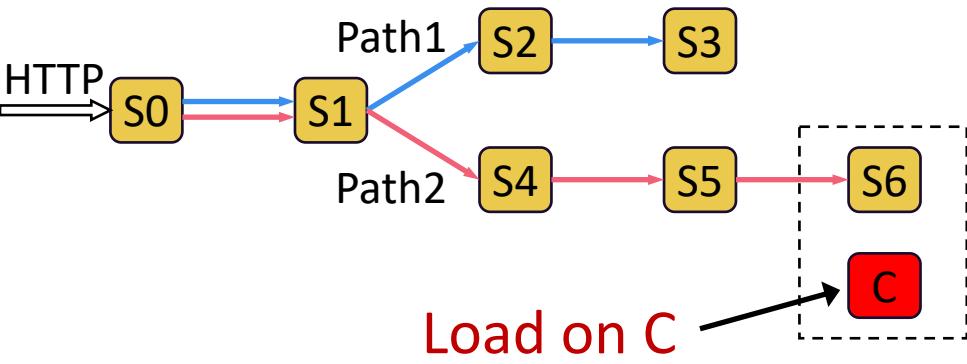
Noisy Neighbor Interference → CPU Saturation on S6



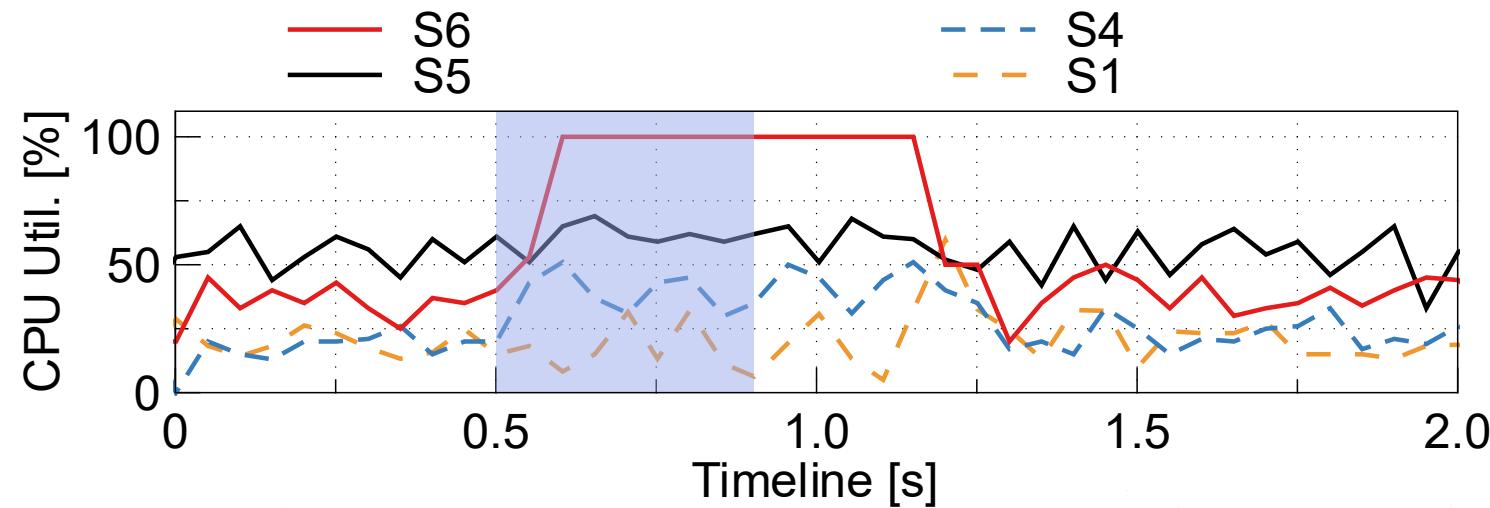
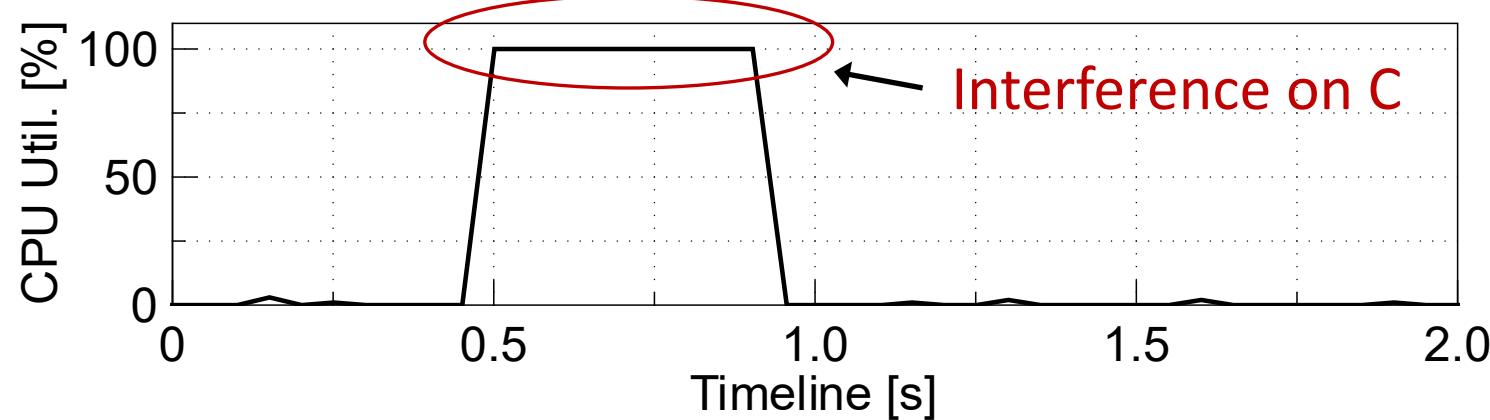
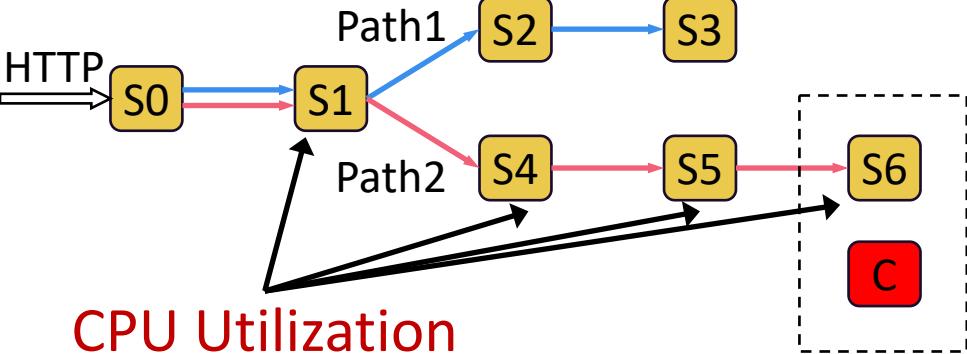
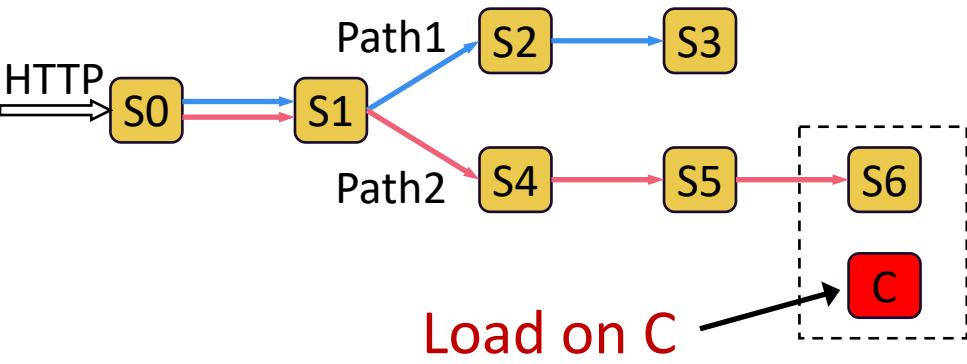
Noisy Neighbor Interference → CPU Saturation on S6



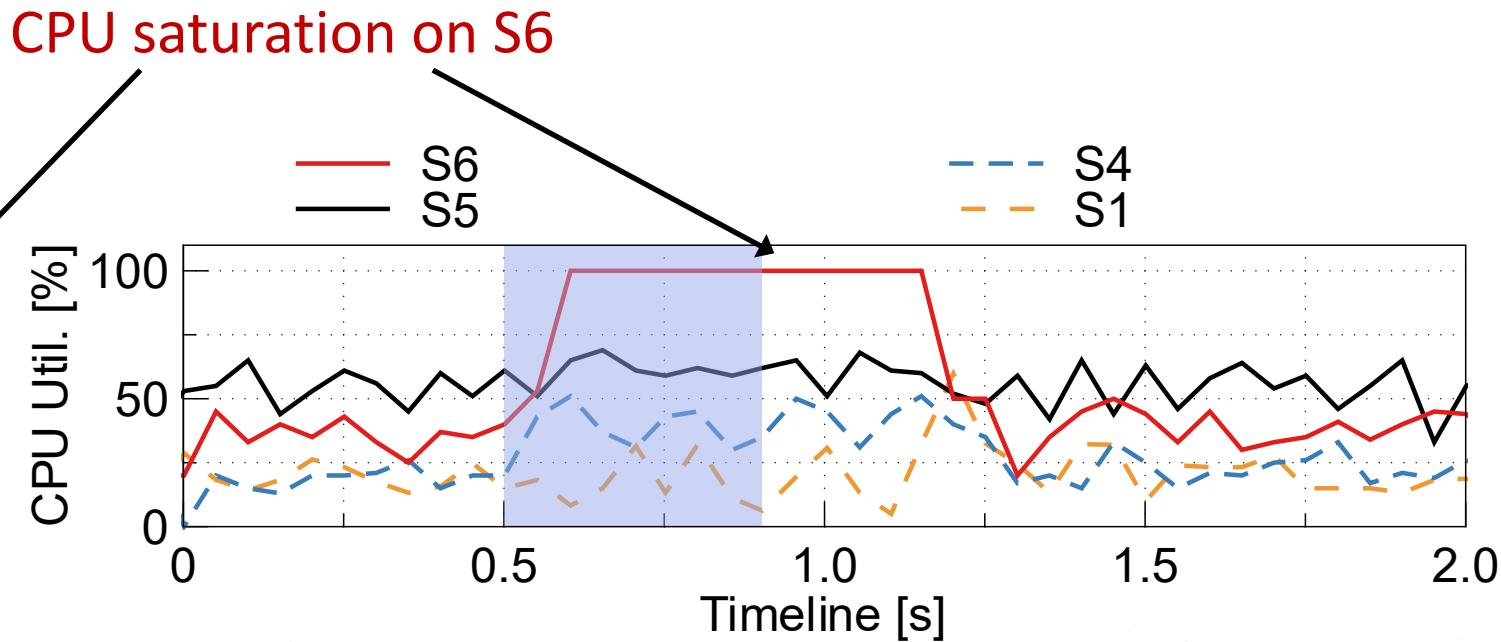
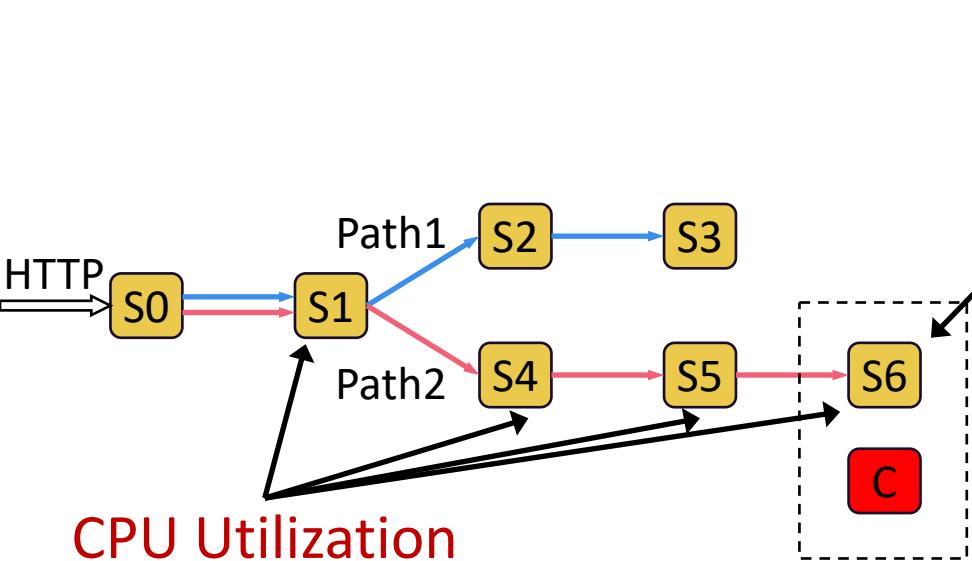
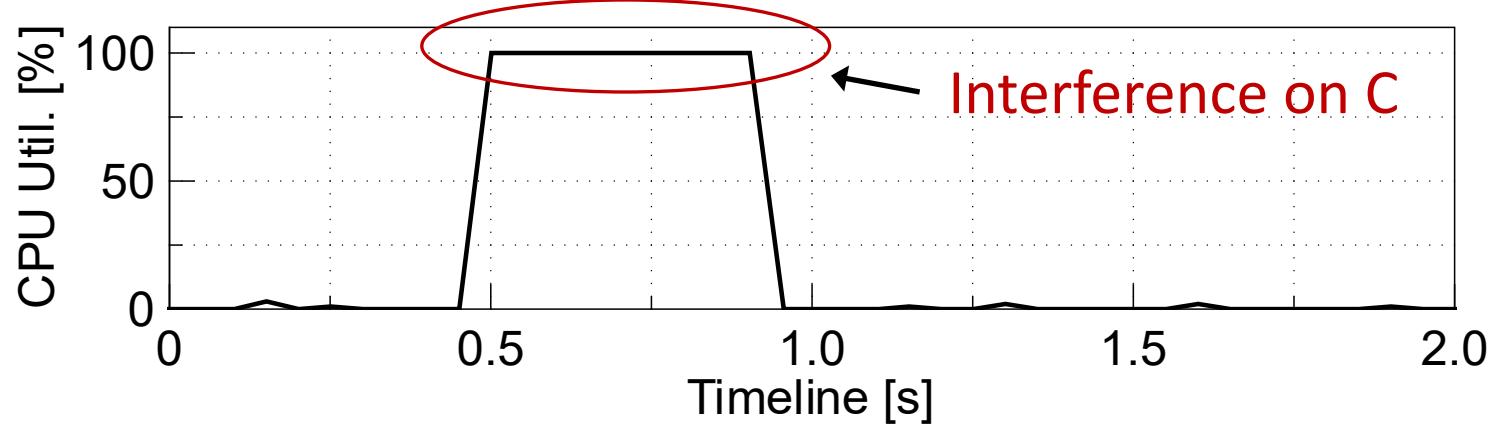
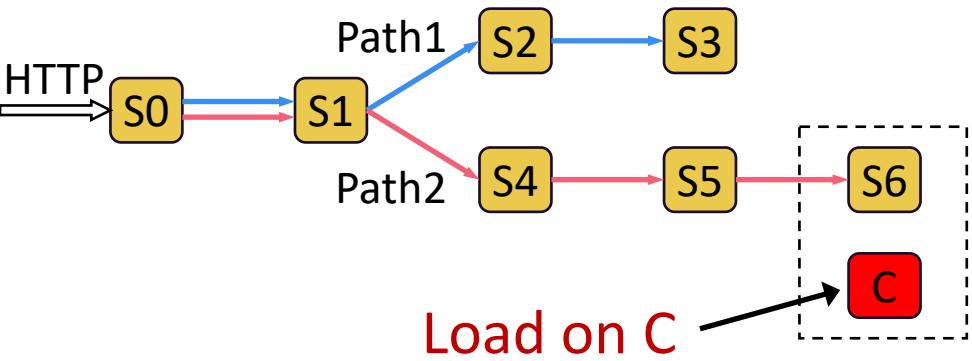
Noisy Neighbor Interference → CPU Saturation on S6



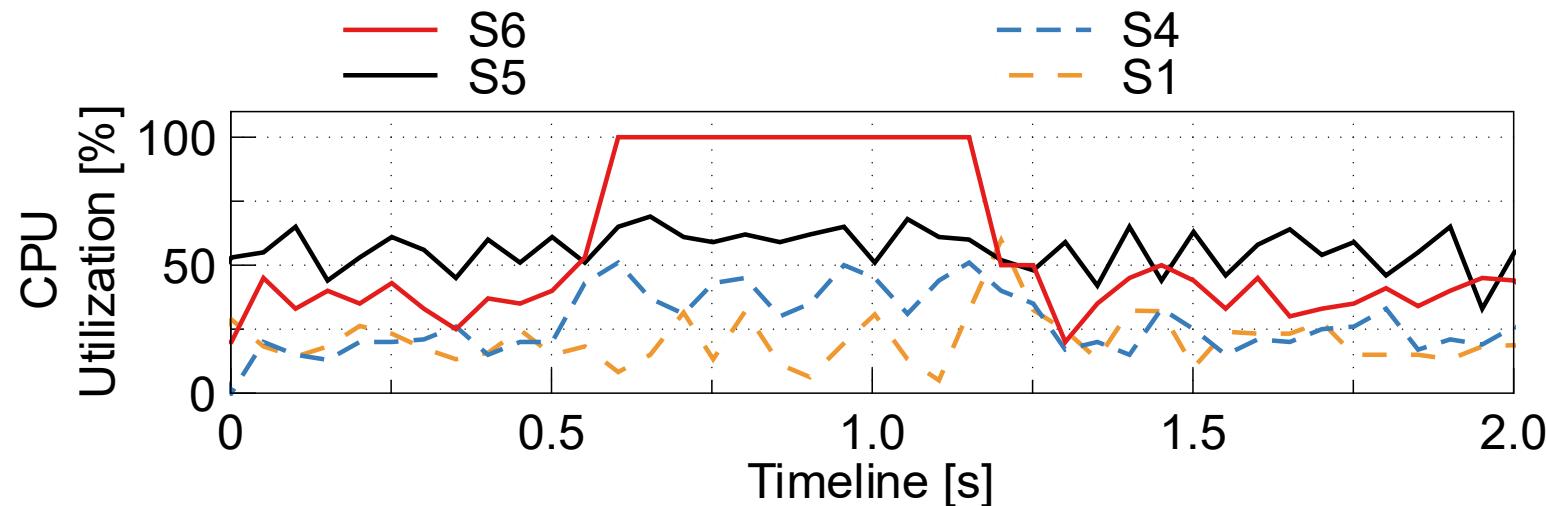
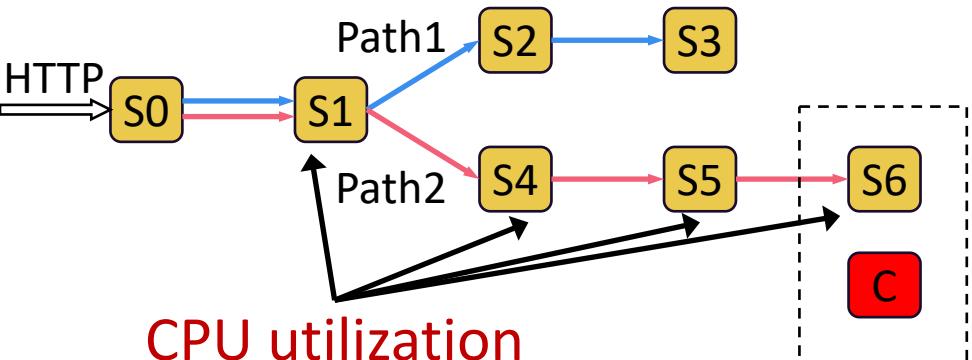
Noisy Neighbor Interference → CPU Saturation on S6



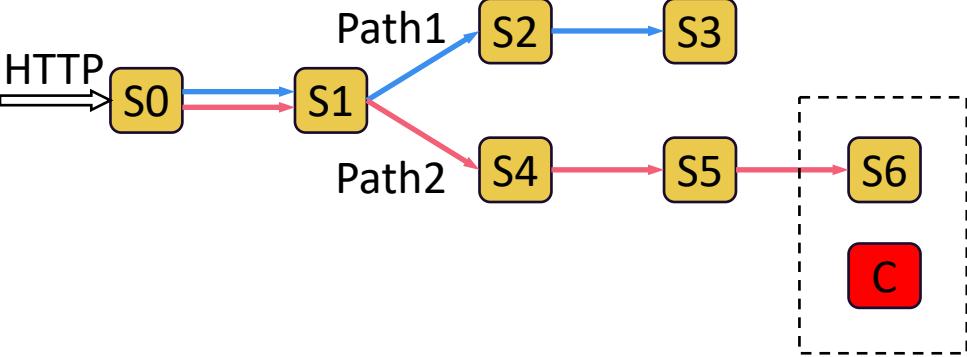
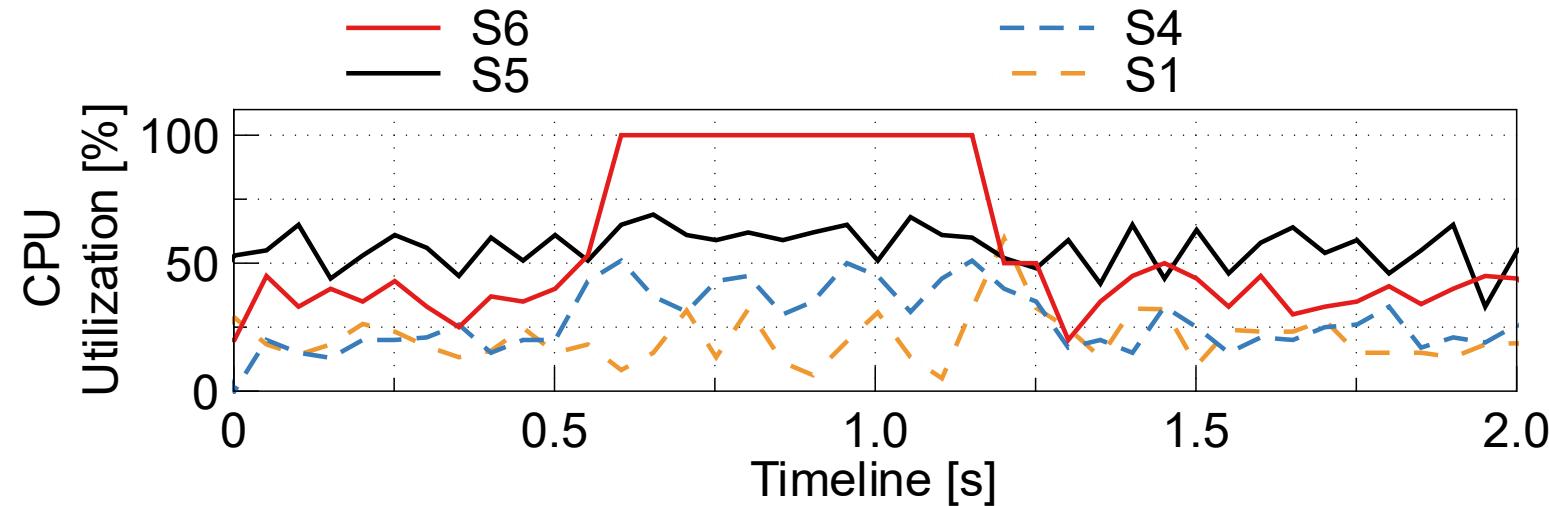
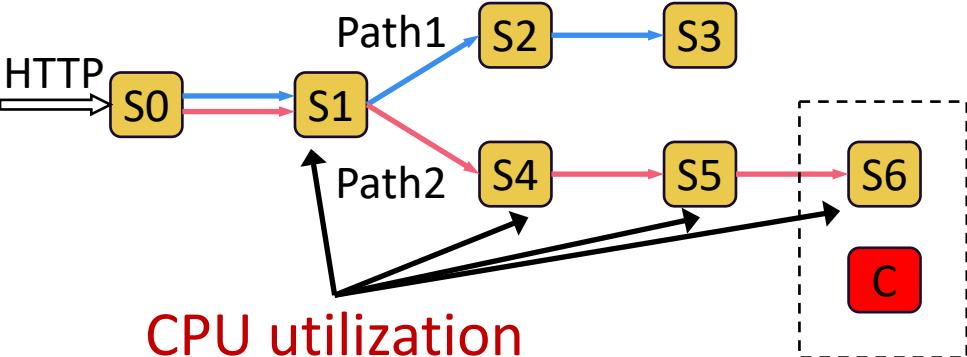
Noisy Neighbor Interference → CPU Saturation on S6



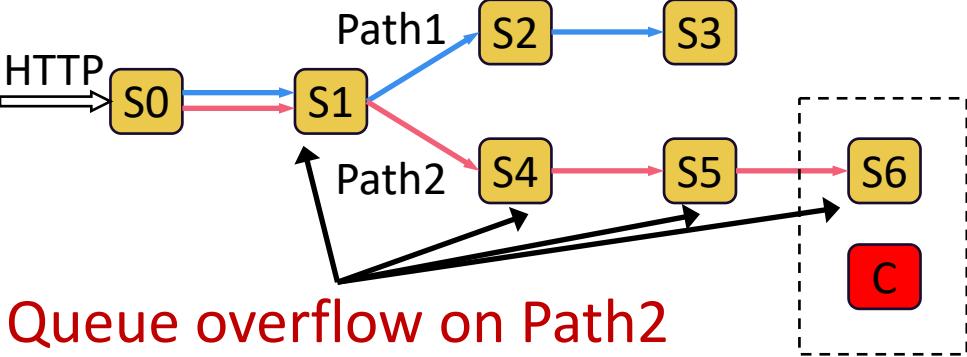
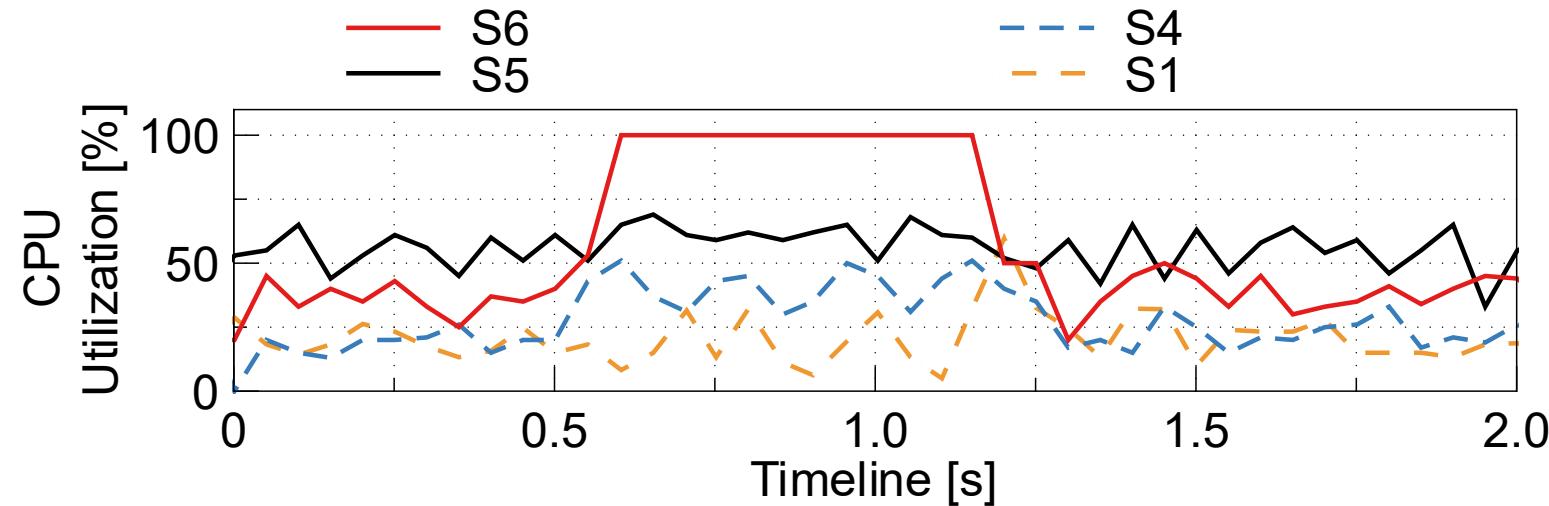
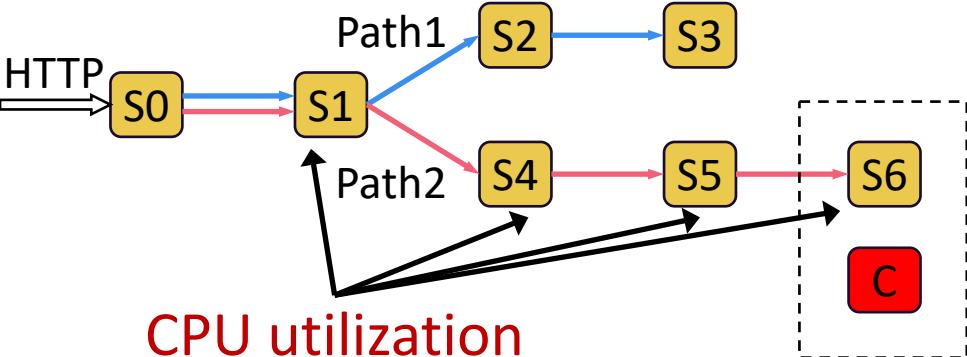
CPU Saturation on S6 → Cross-Service Queue Overflow



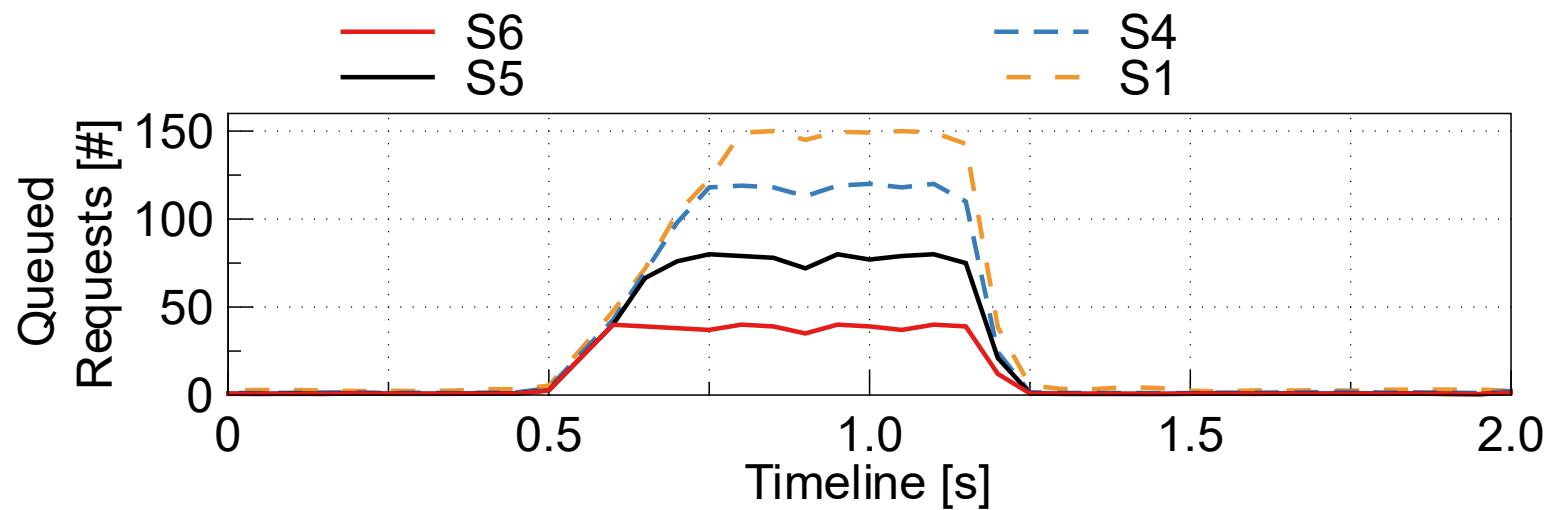
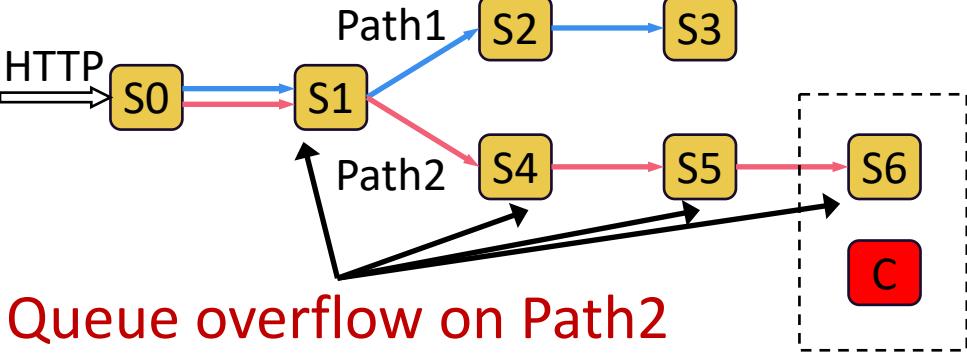
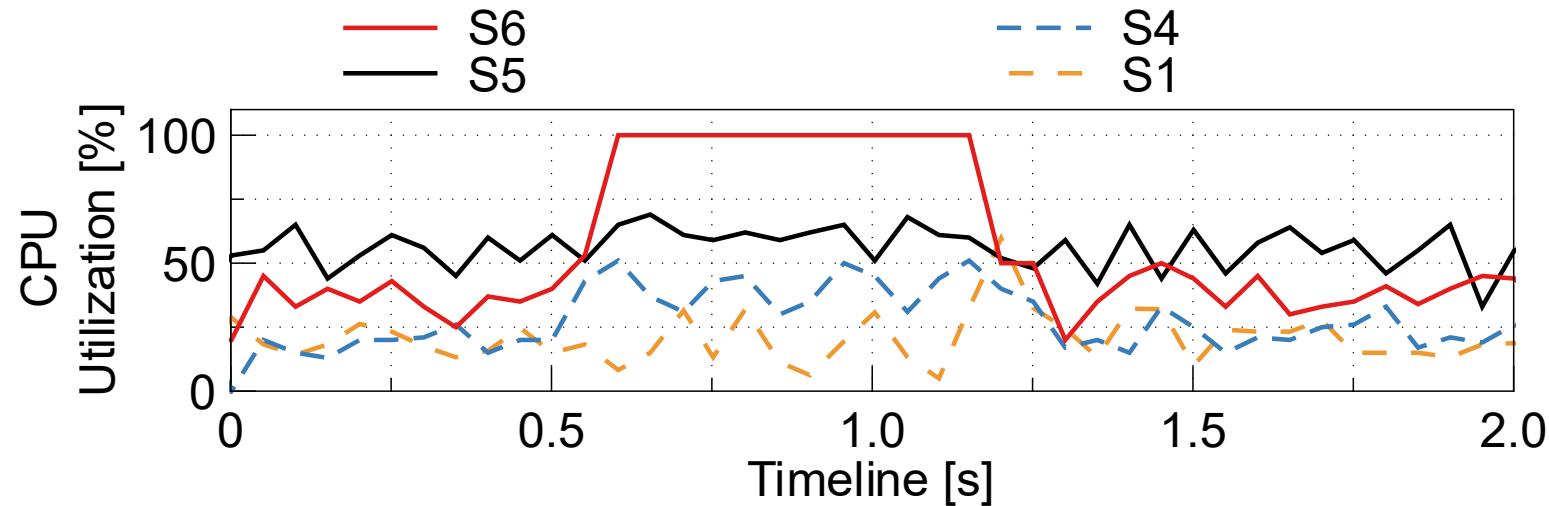
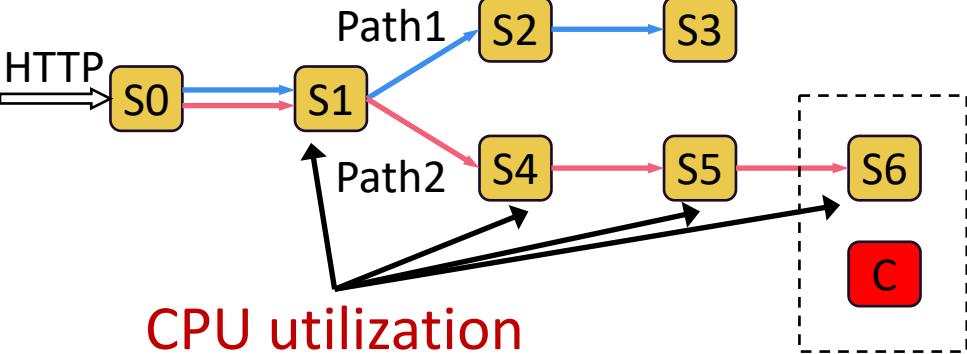
CPU Saturation on S6 → Cross-Service Queue Overflow



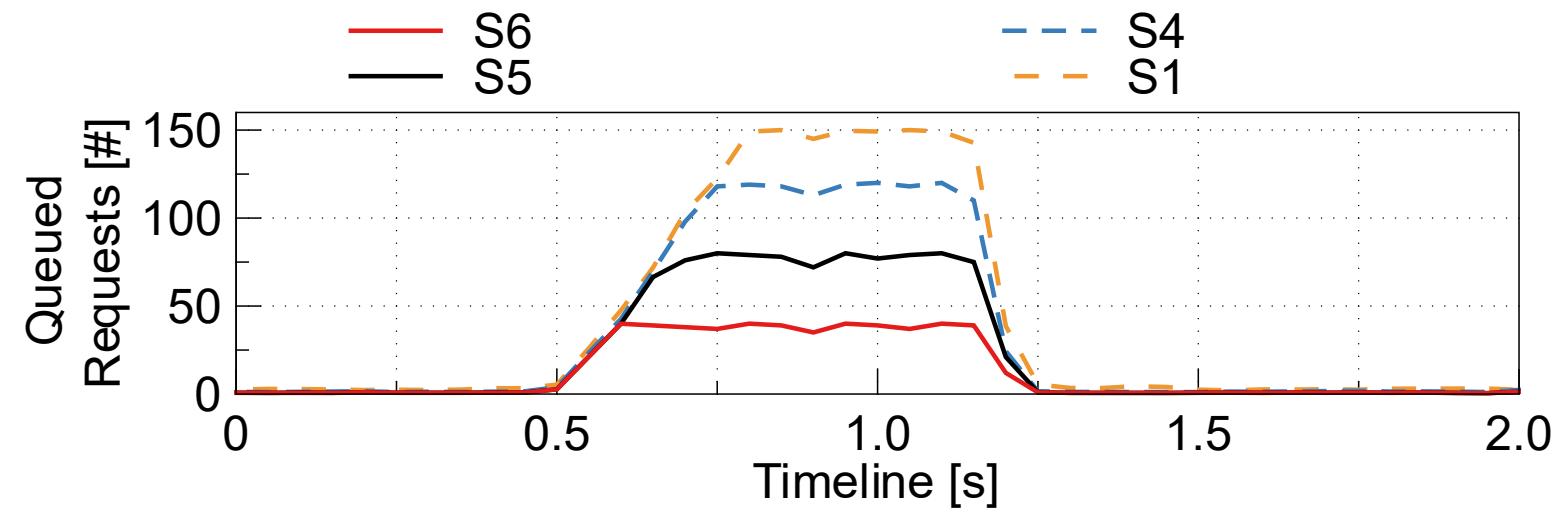
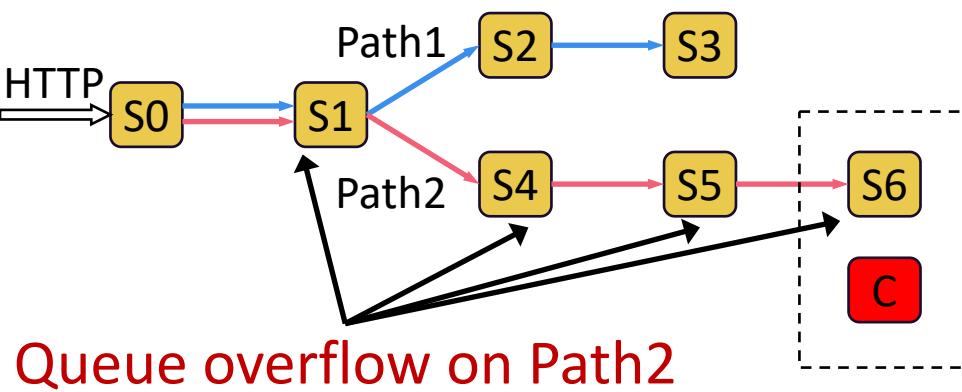
CPU Saturation on S6 → Cross-Service Queue Overflow



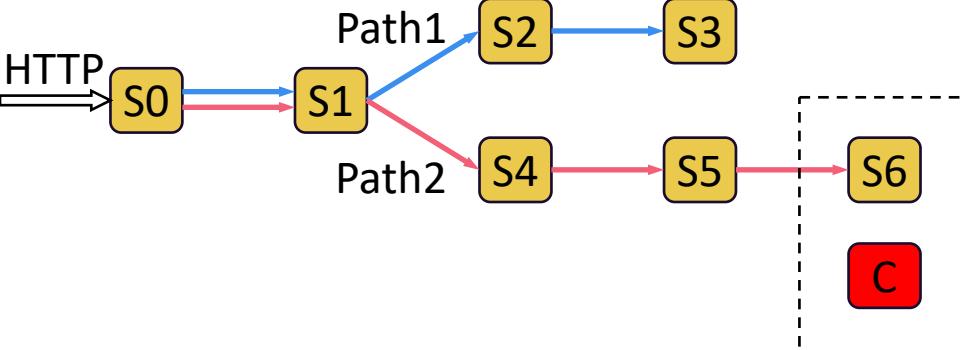
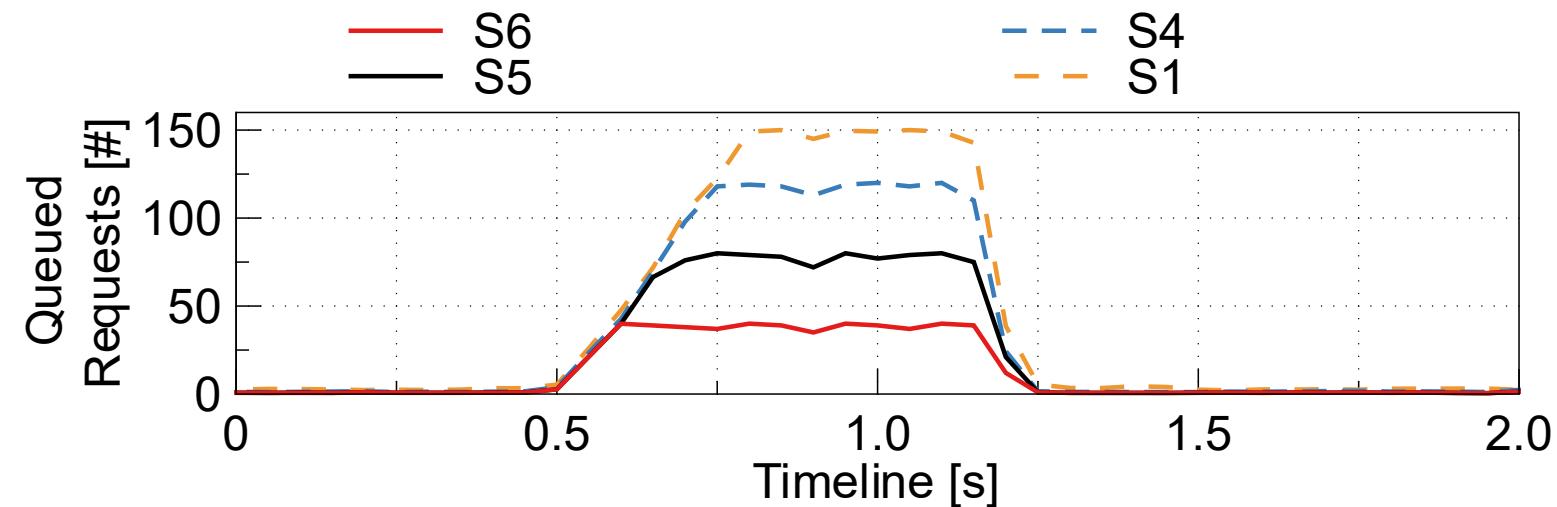
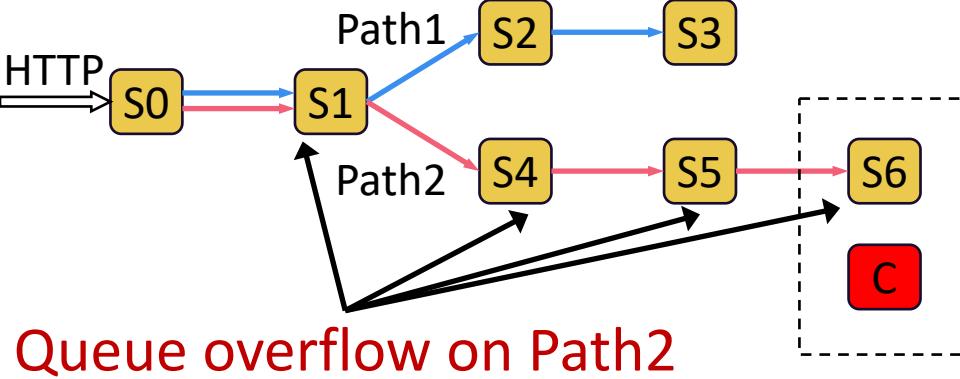
CPU Saturation on S6 → Cross-Service Queue Overflow



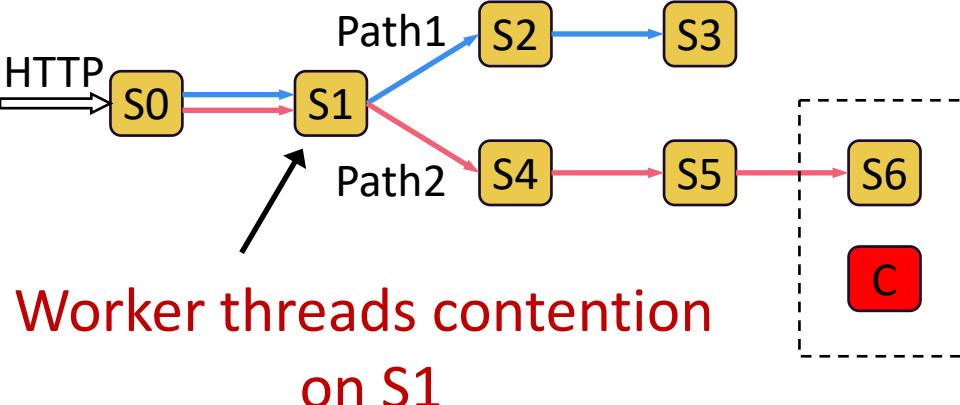
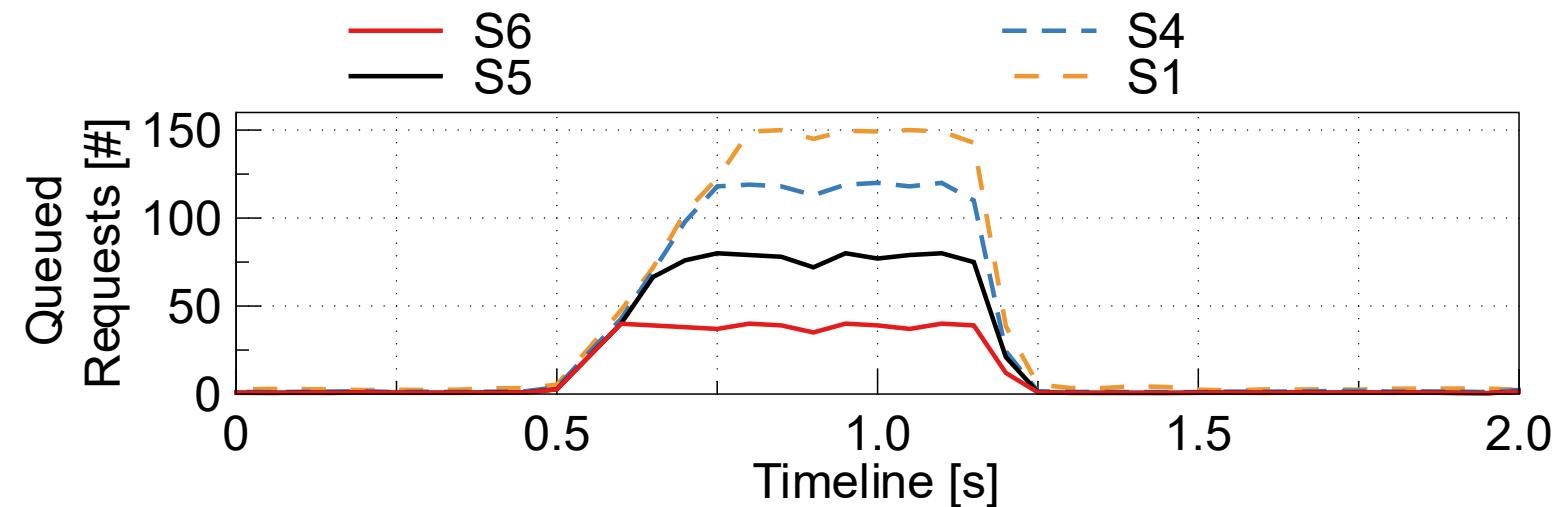
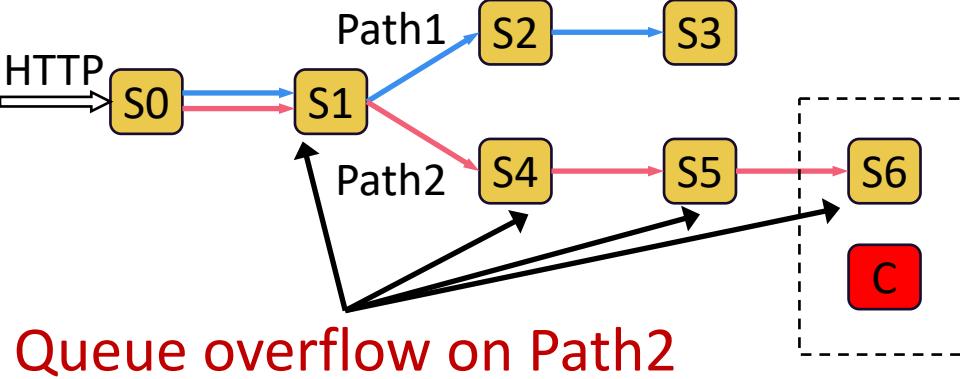
Cross-Service Queue Overflow → Worker Thread Contention



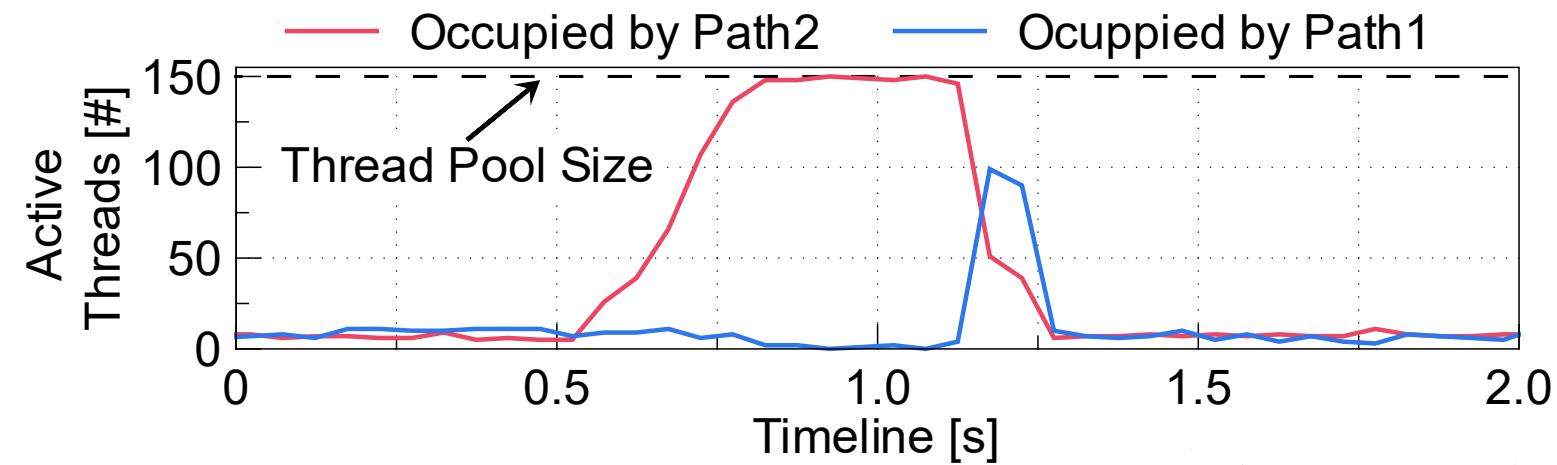
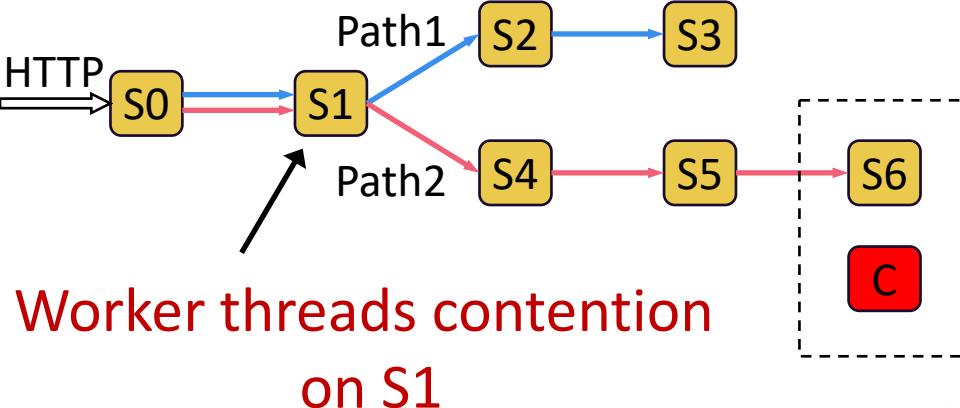
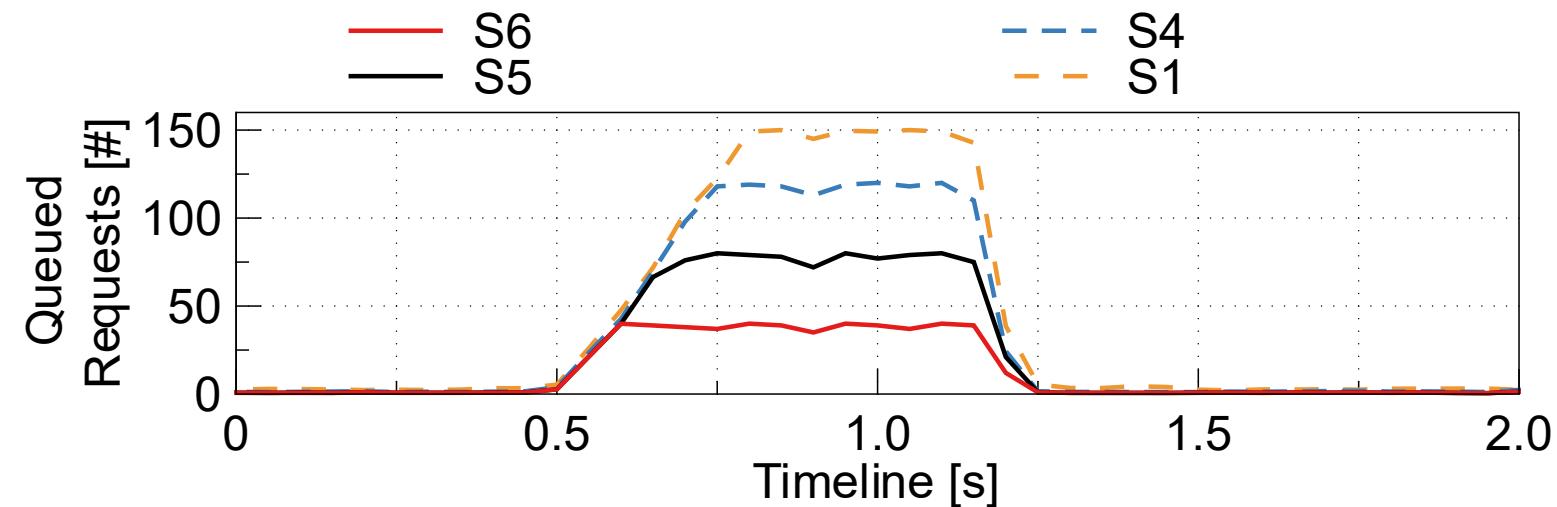
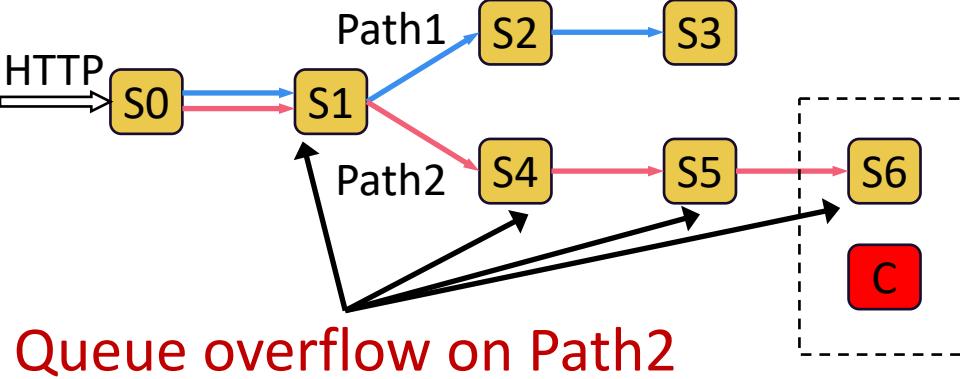
Cross-Service Queue Overflow → Worker Thread Contention



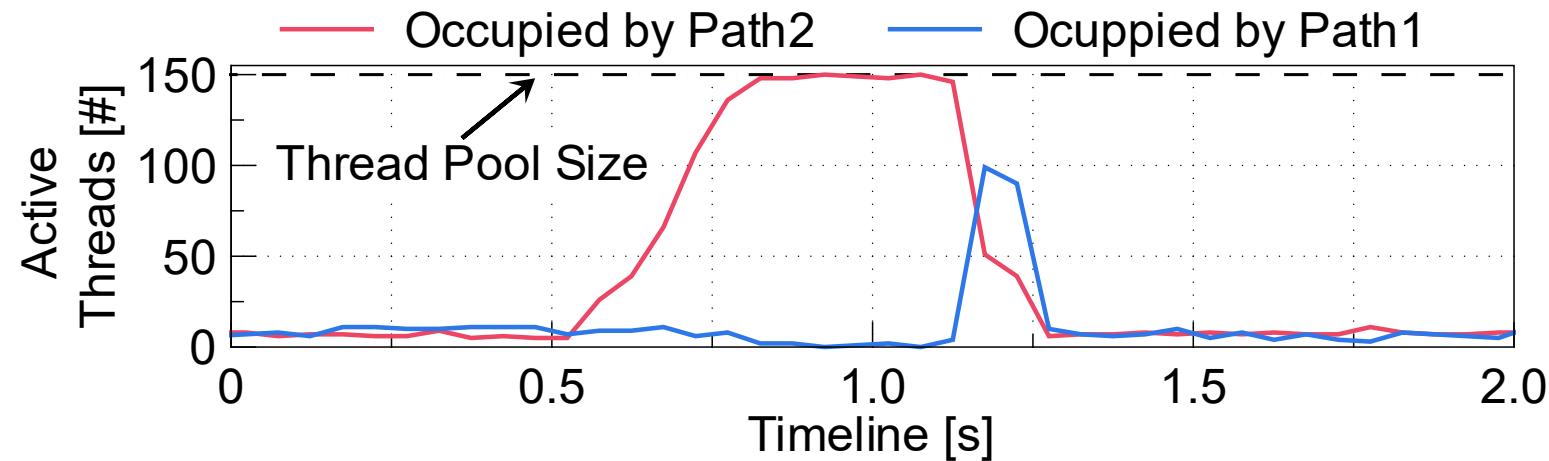
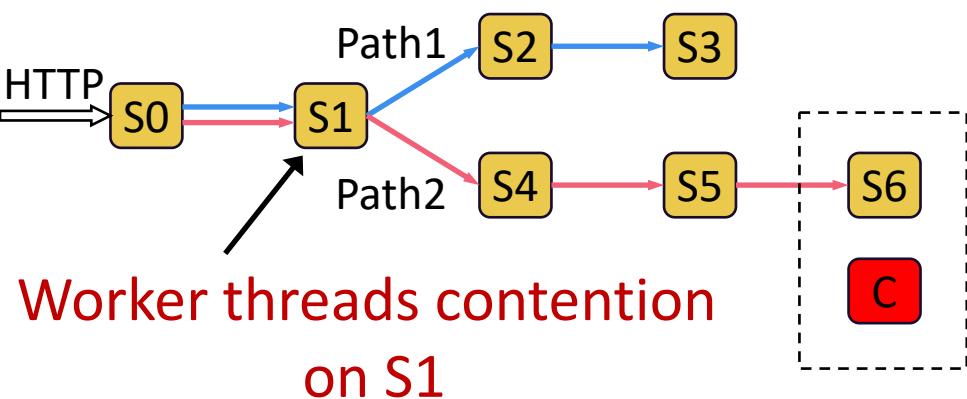
Cross-Service Queue Overflow → Worker Thread Contention



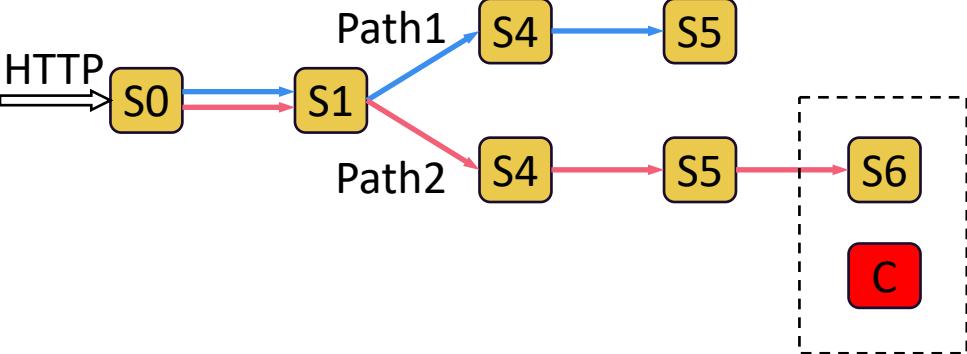
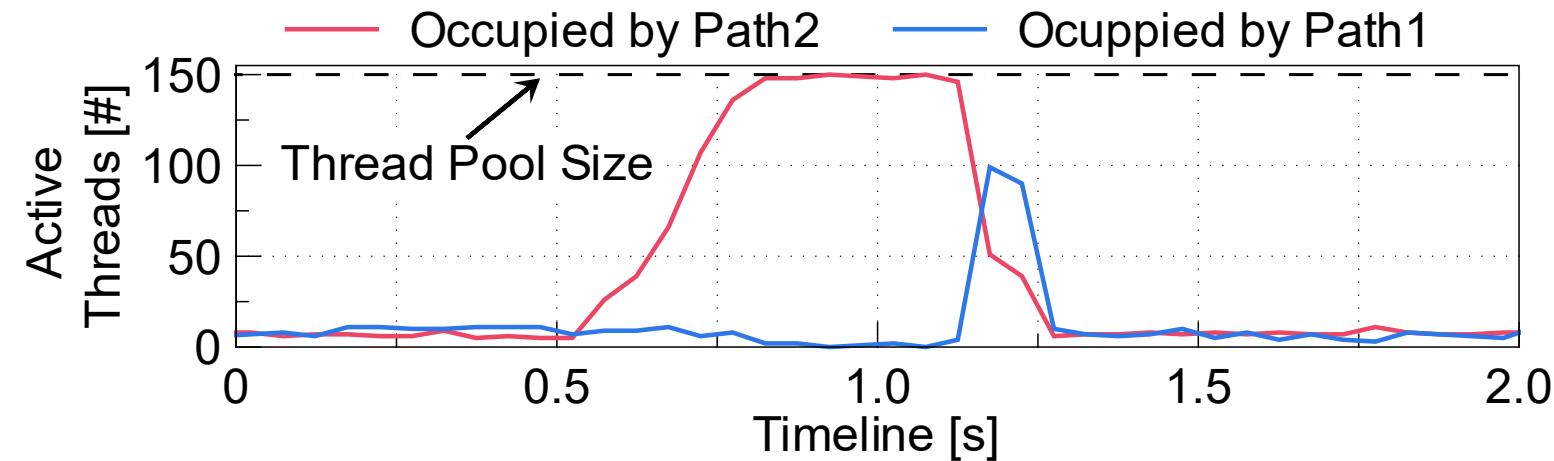
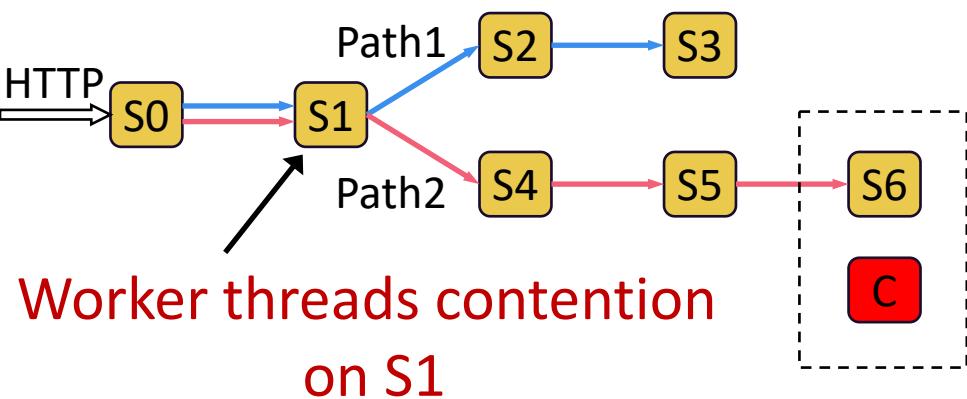
Cross-Service Queue Overflow → Worker Thread Contention



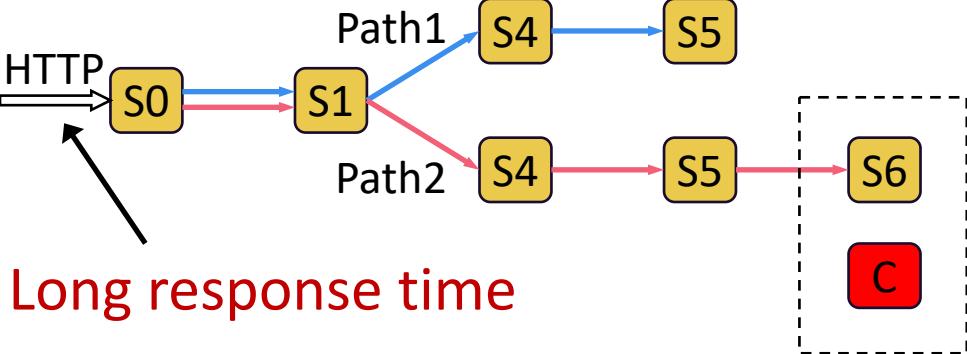
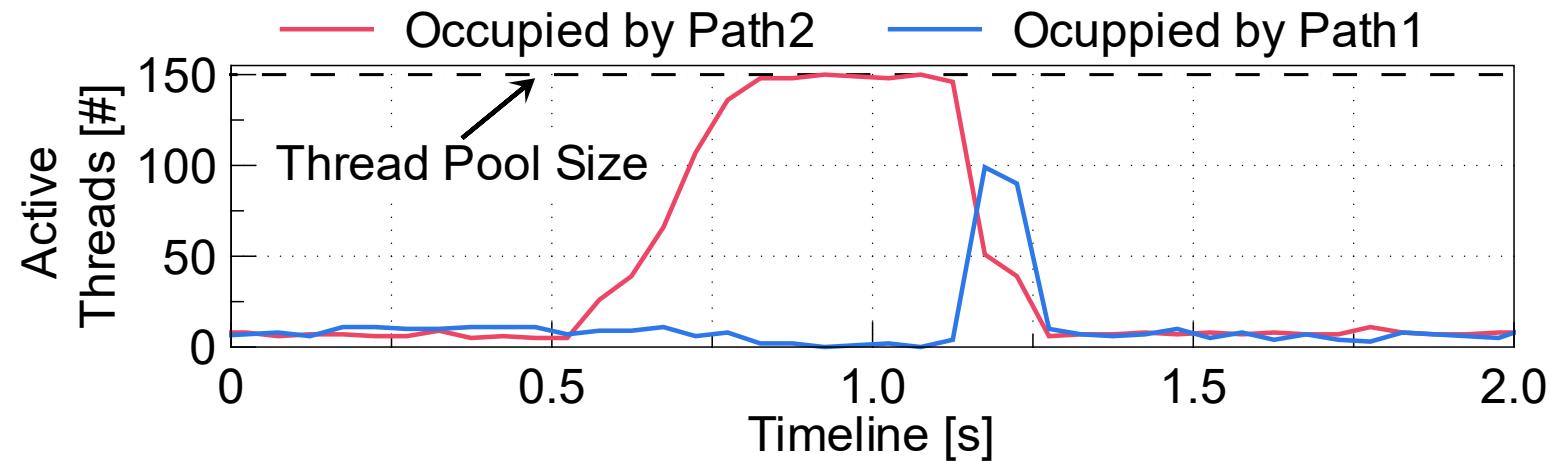
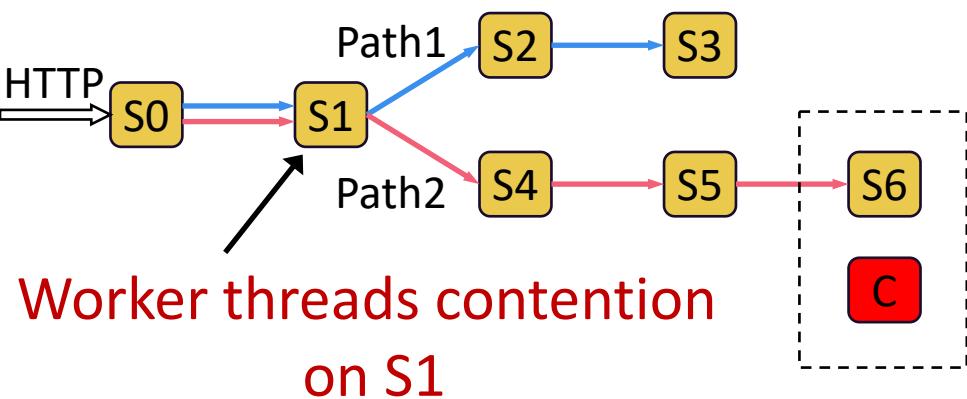
Worker Thread Contention → Long Response Time



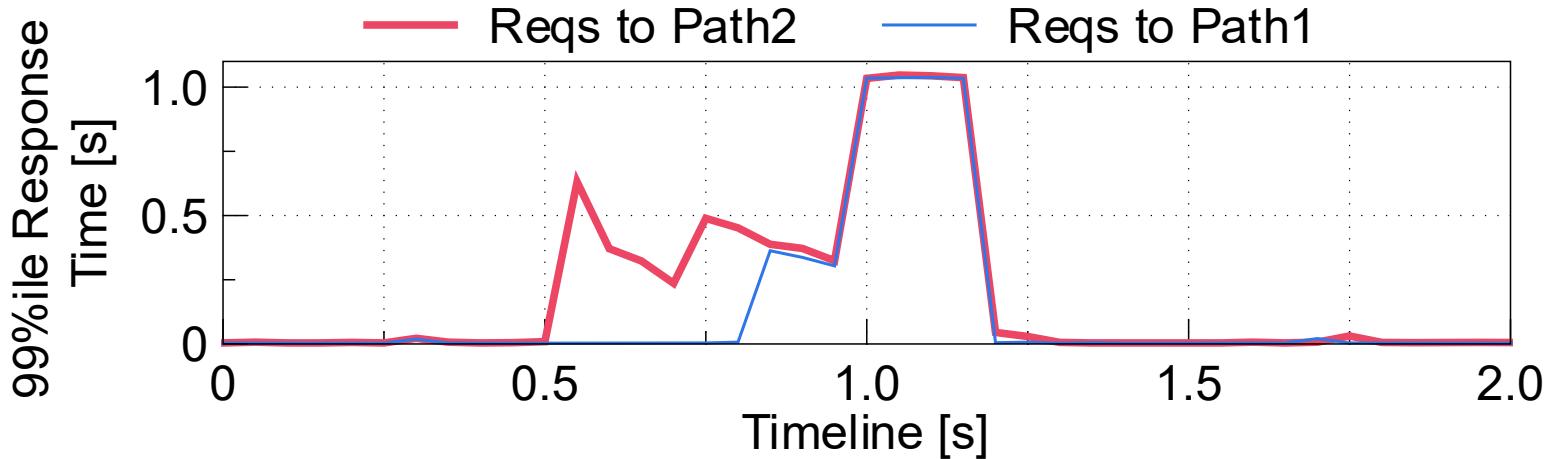
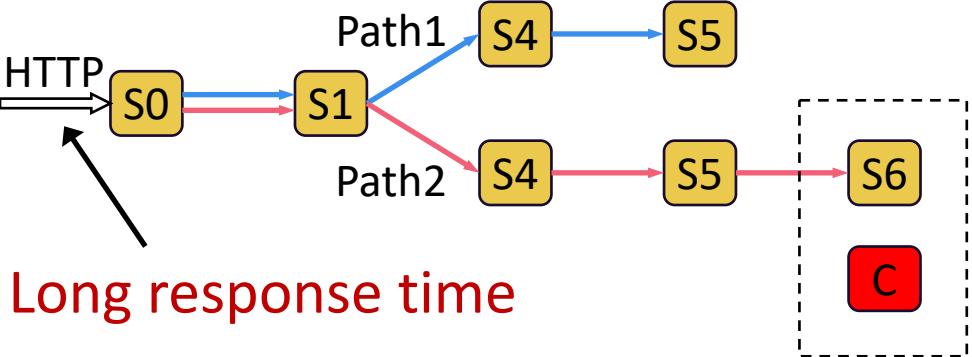
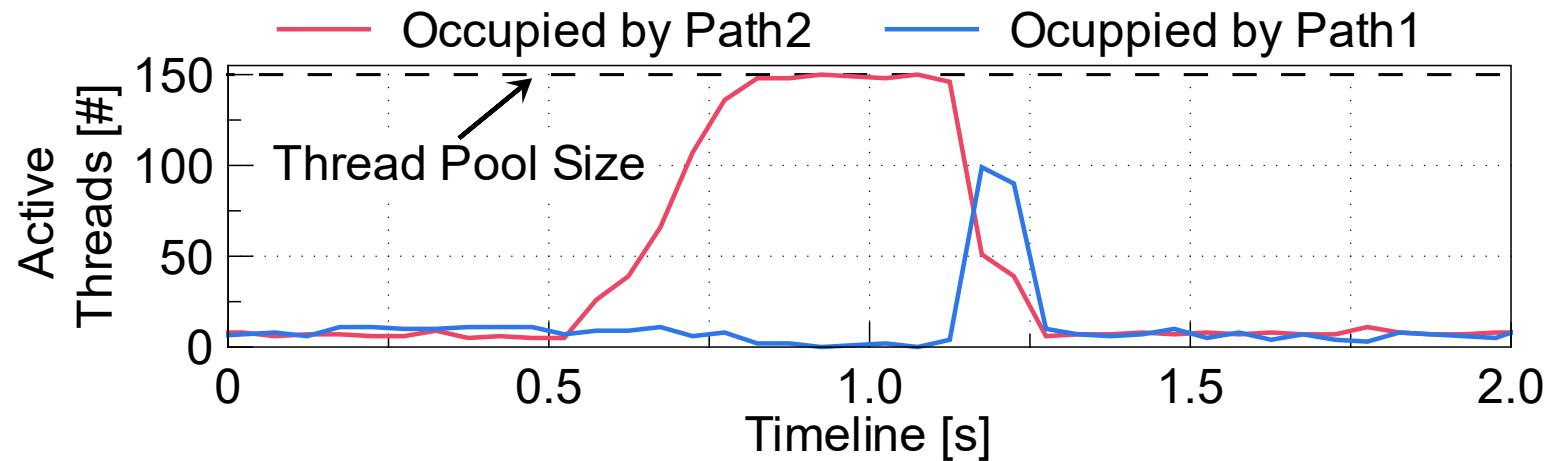
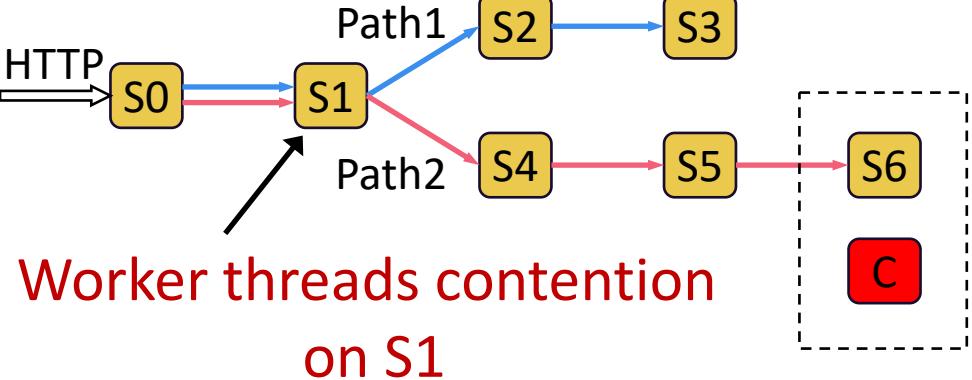
Worker Thread Contention → Long Response Time



Worker Thread Contention → Long Response Time



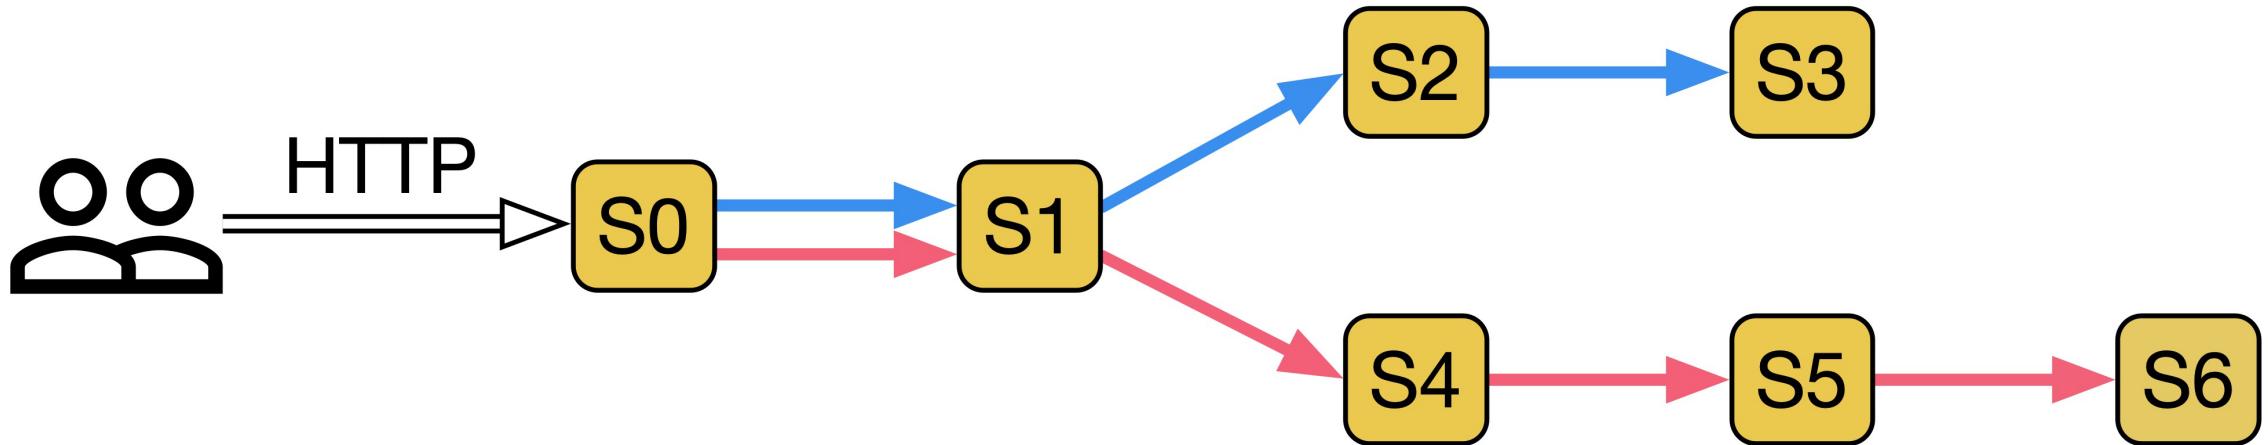
Worker Thread Contention → Long Response Time



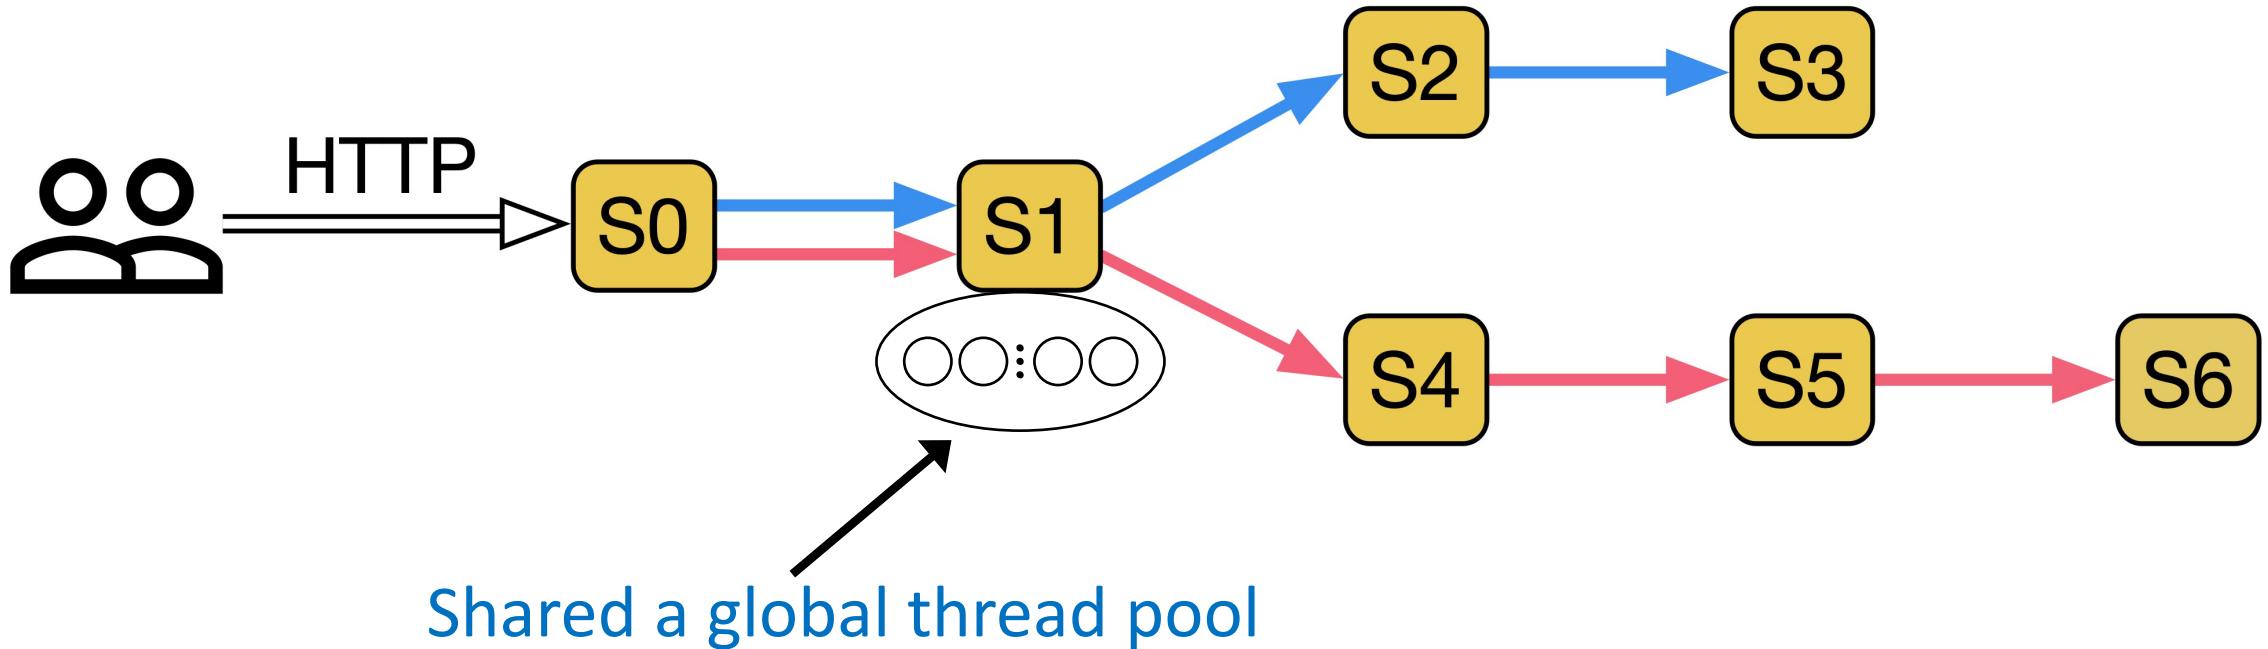
Other Sources of Bottleneck in Microservices

- ❑ **Bursty workload** from external users
- ❑ **Collocation interference** from internal containers or processes
- ❑ **Self-maintenance activities** within services
 - e.g., garbage collection (GC), log flushing
- ❑ **Potential bottleneck resources:**
 - Hardware resources (e.g., CPU, disk, memory bandwidth)
 - Soft resources (e.g., threads, database connection pool)

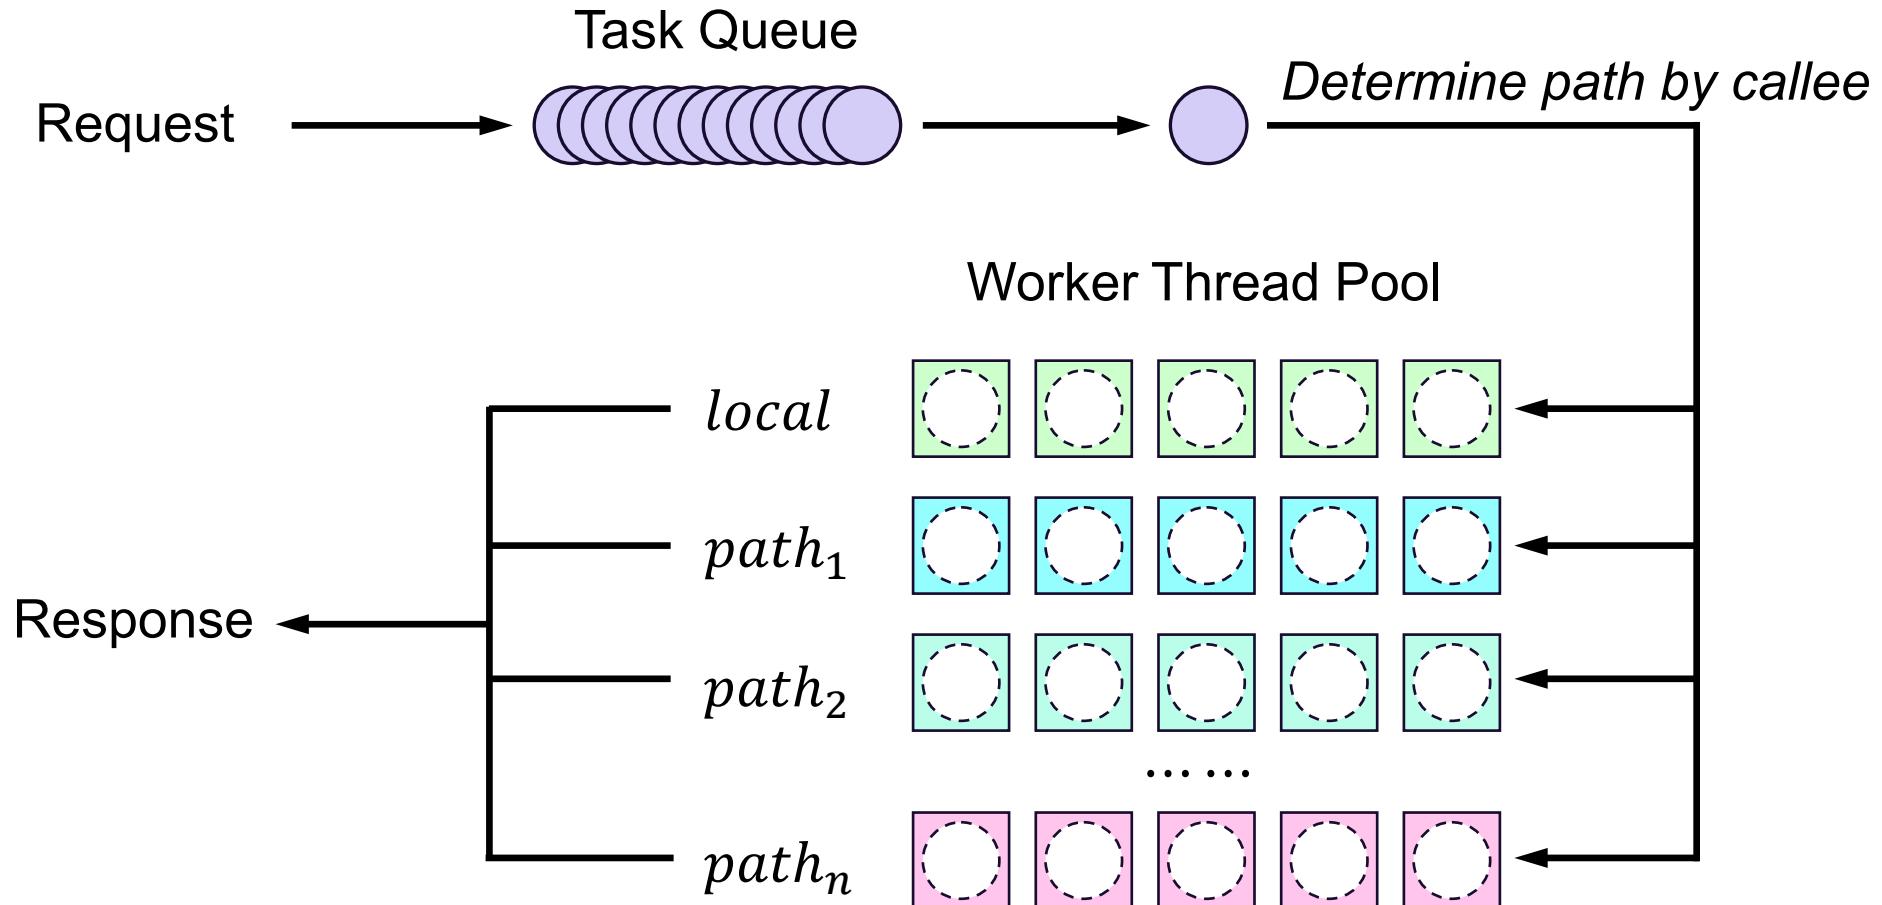
Root Cause: Shared Global Thread Pool



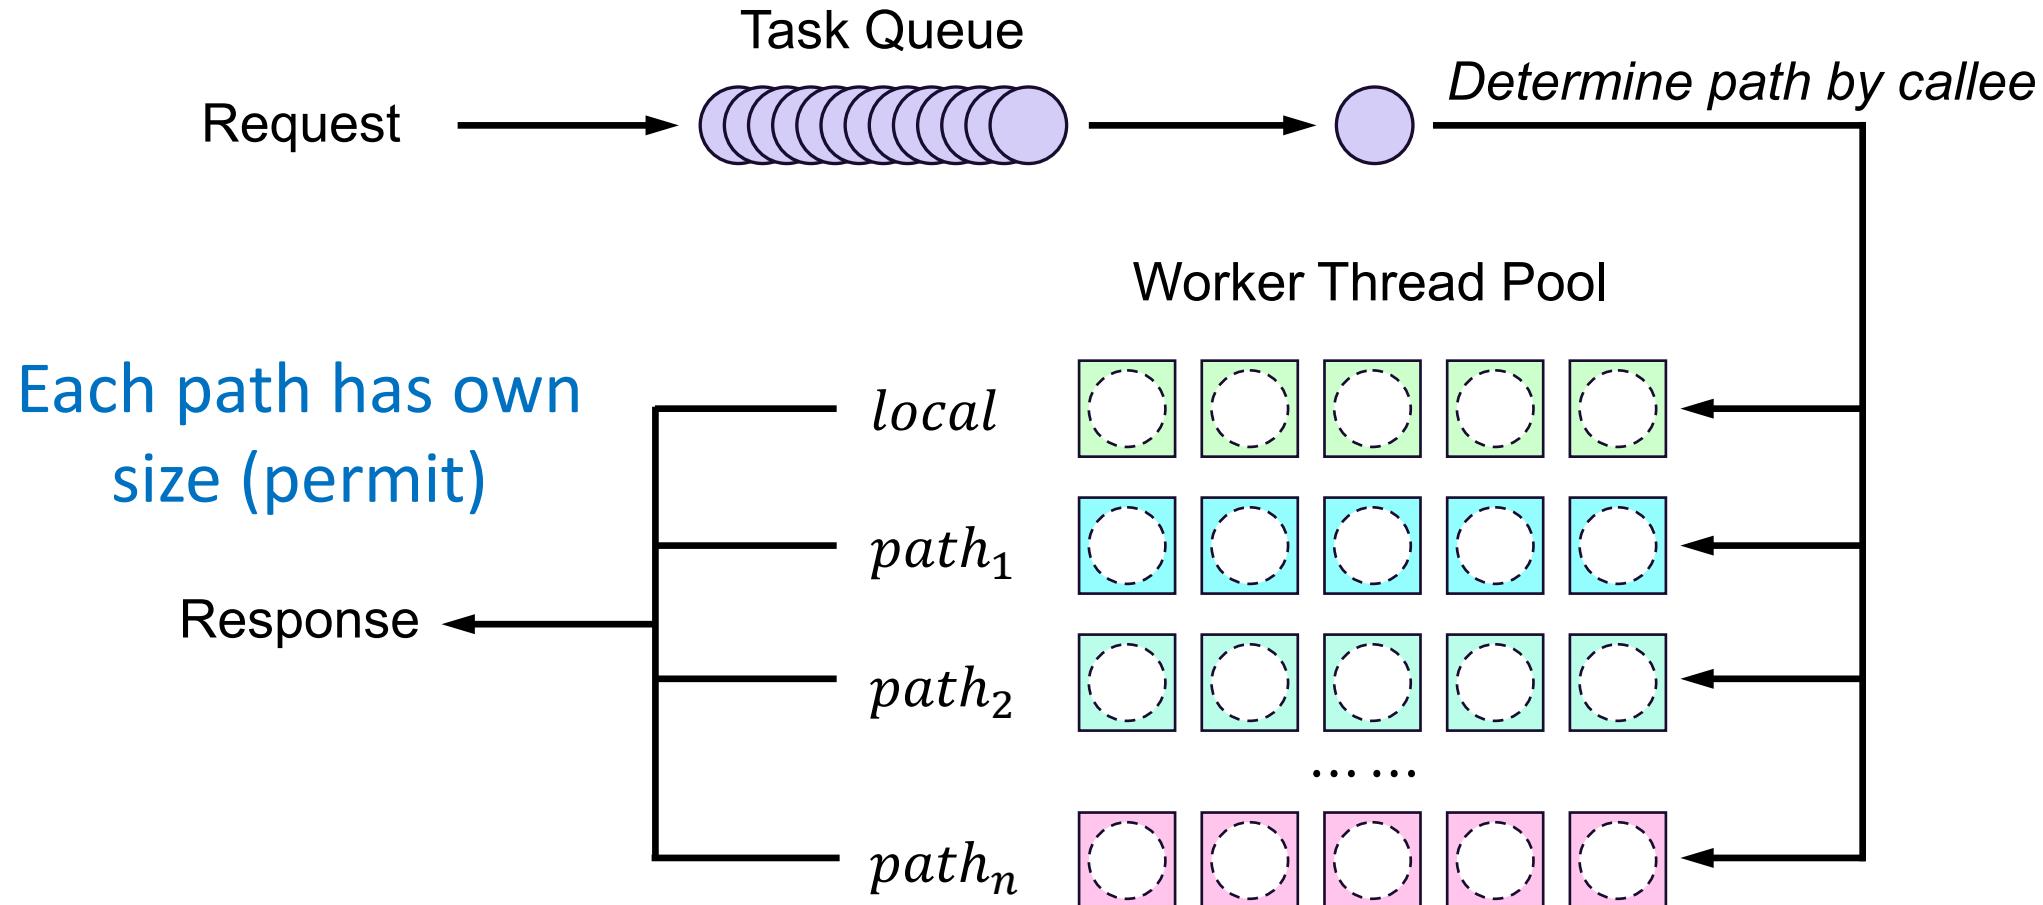
Root Cause: Shared Global Thread Pool



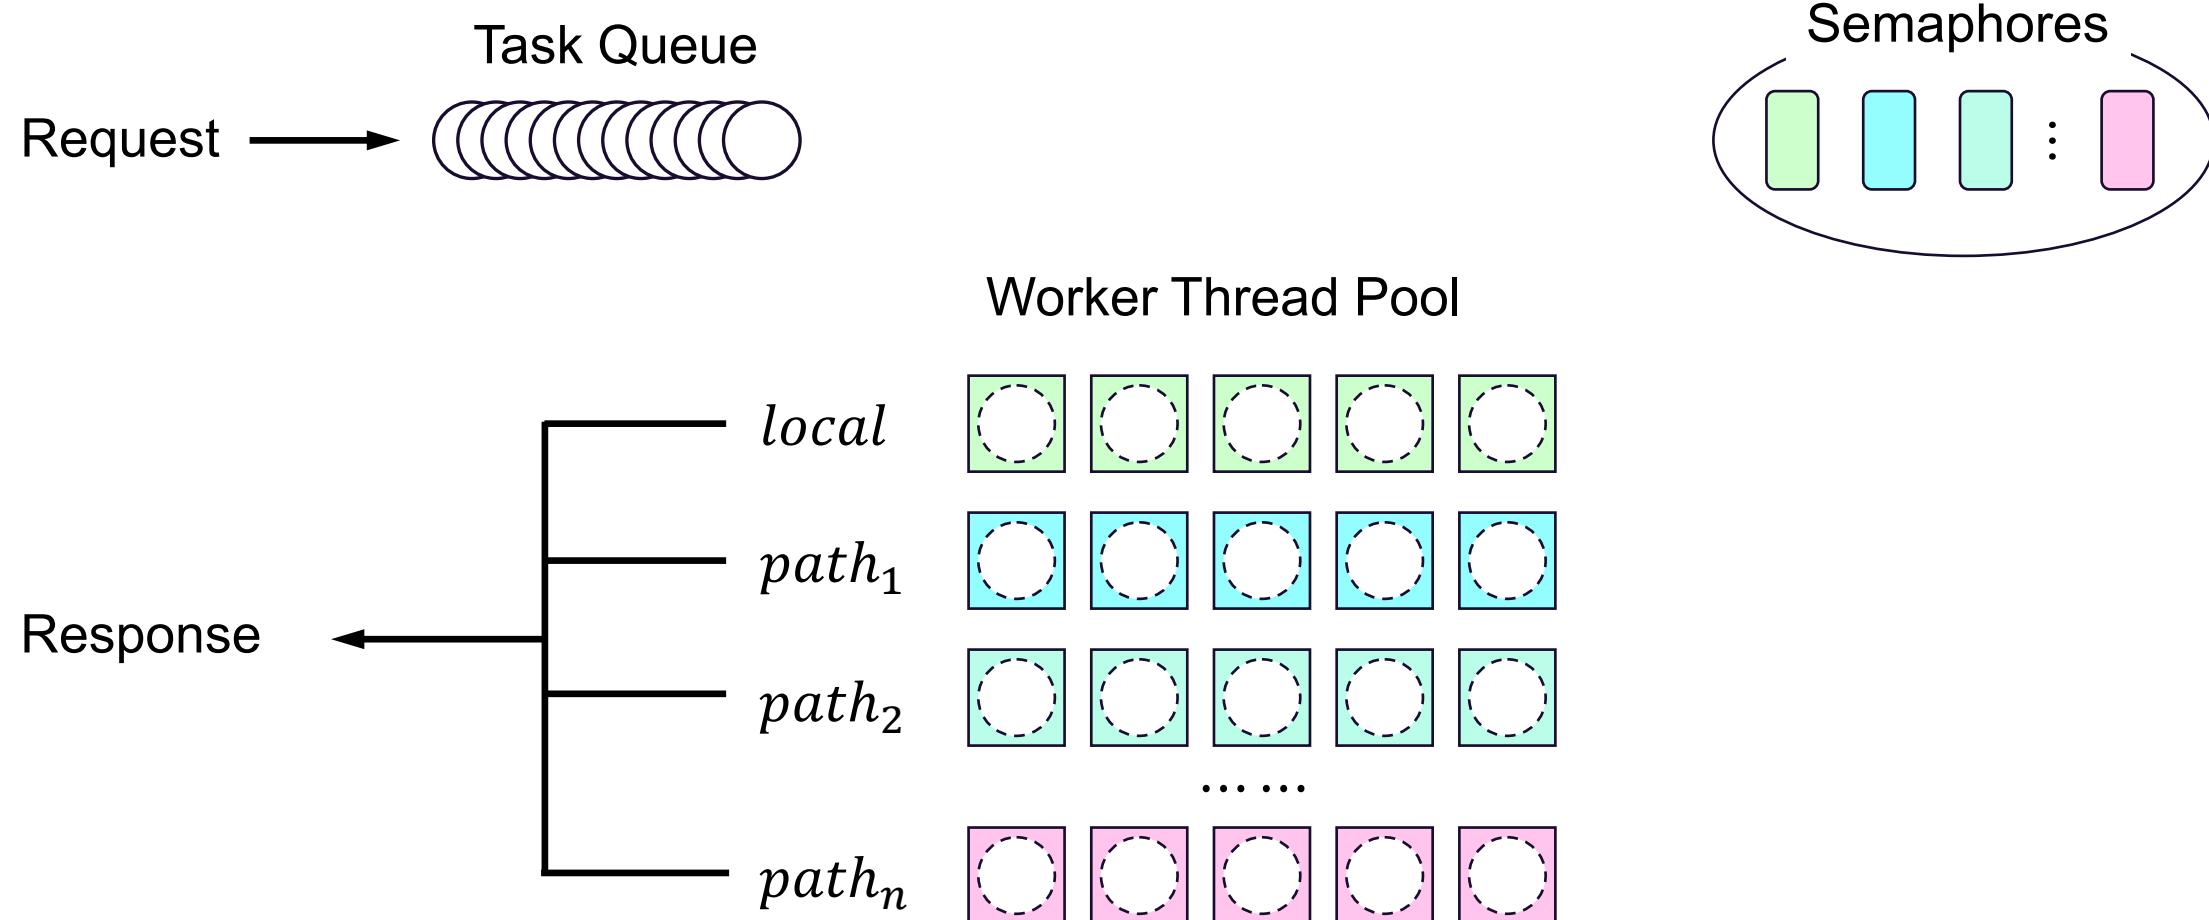
PathFence: Path-Aware Thread Scheduling



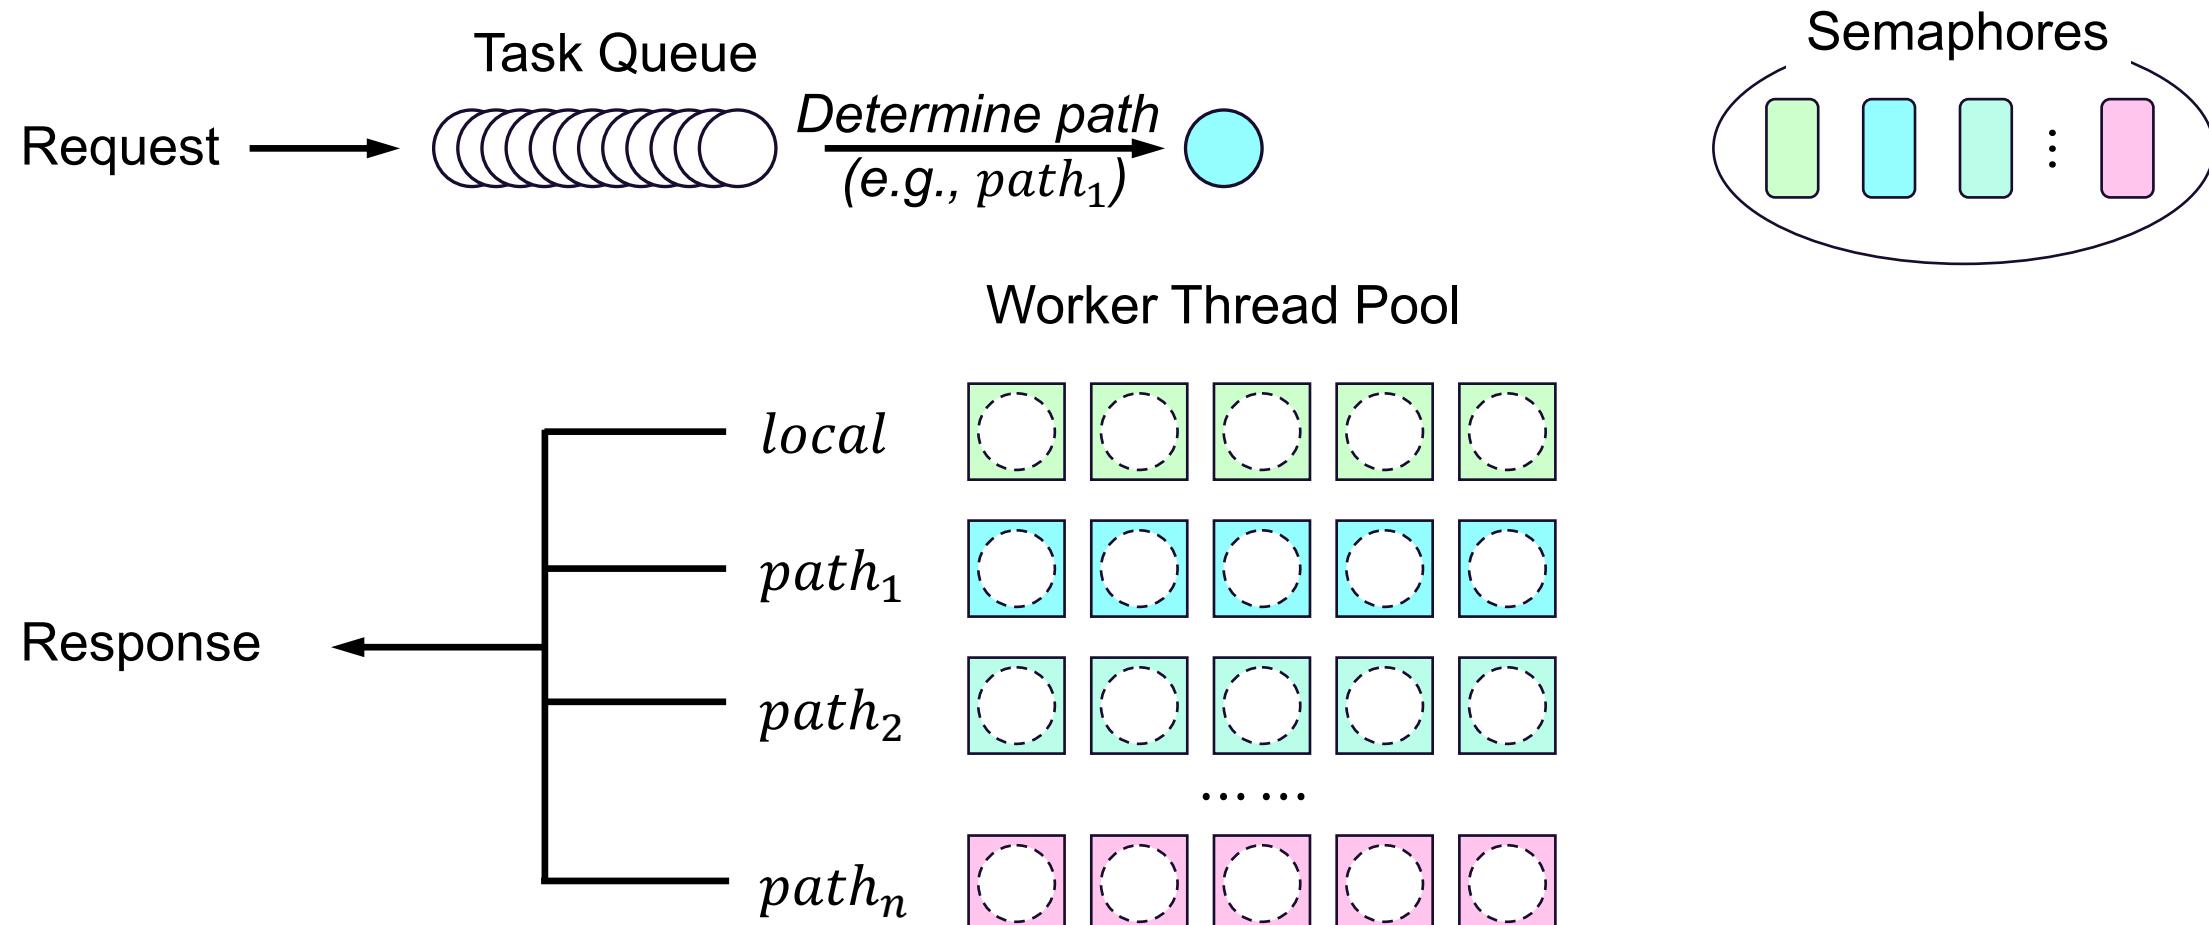
PathFence: Path-Aware Thread Scheduling



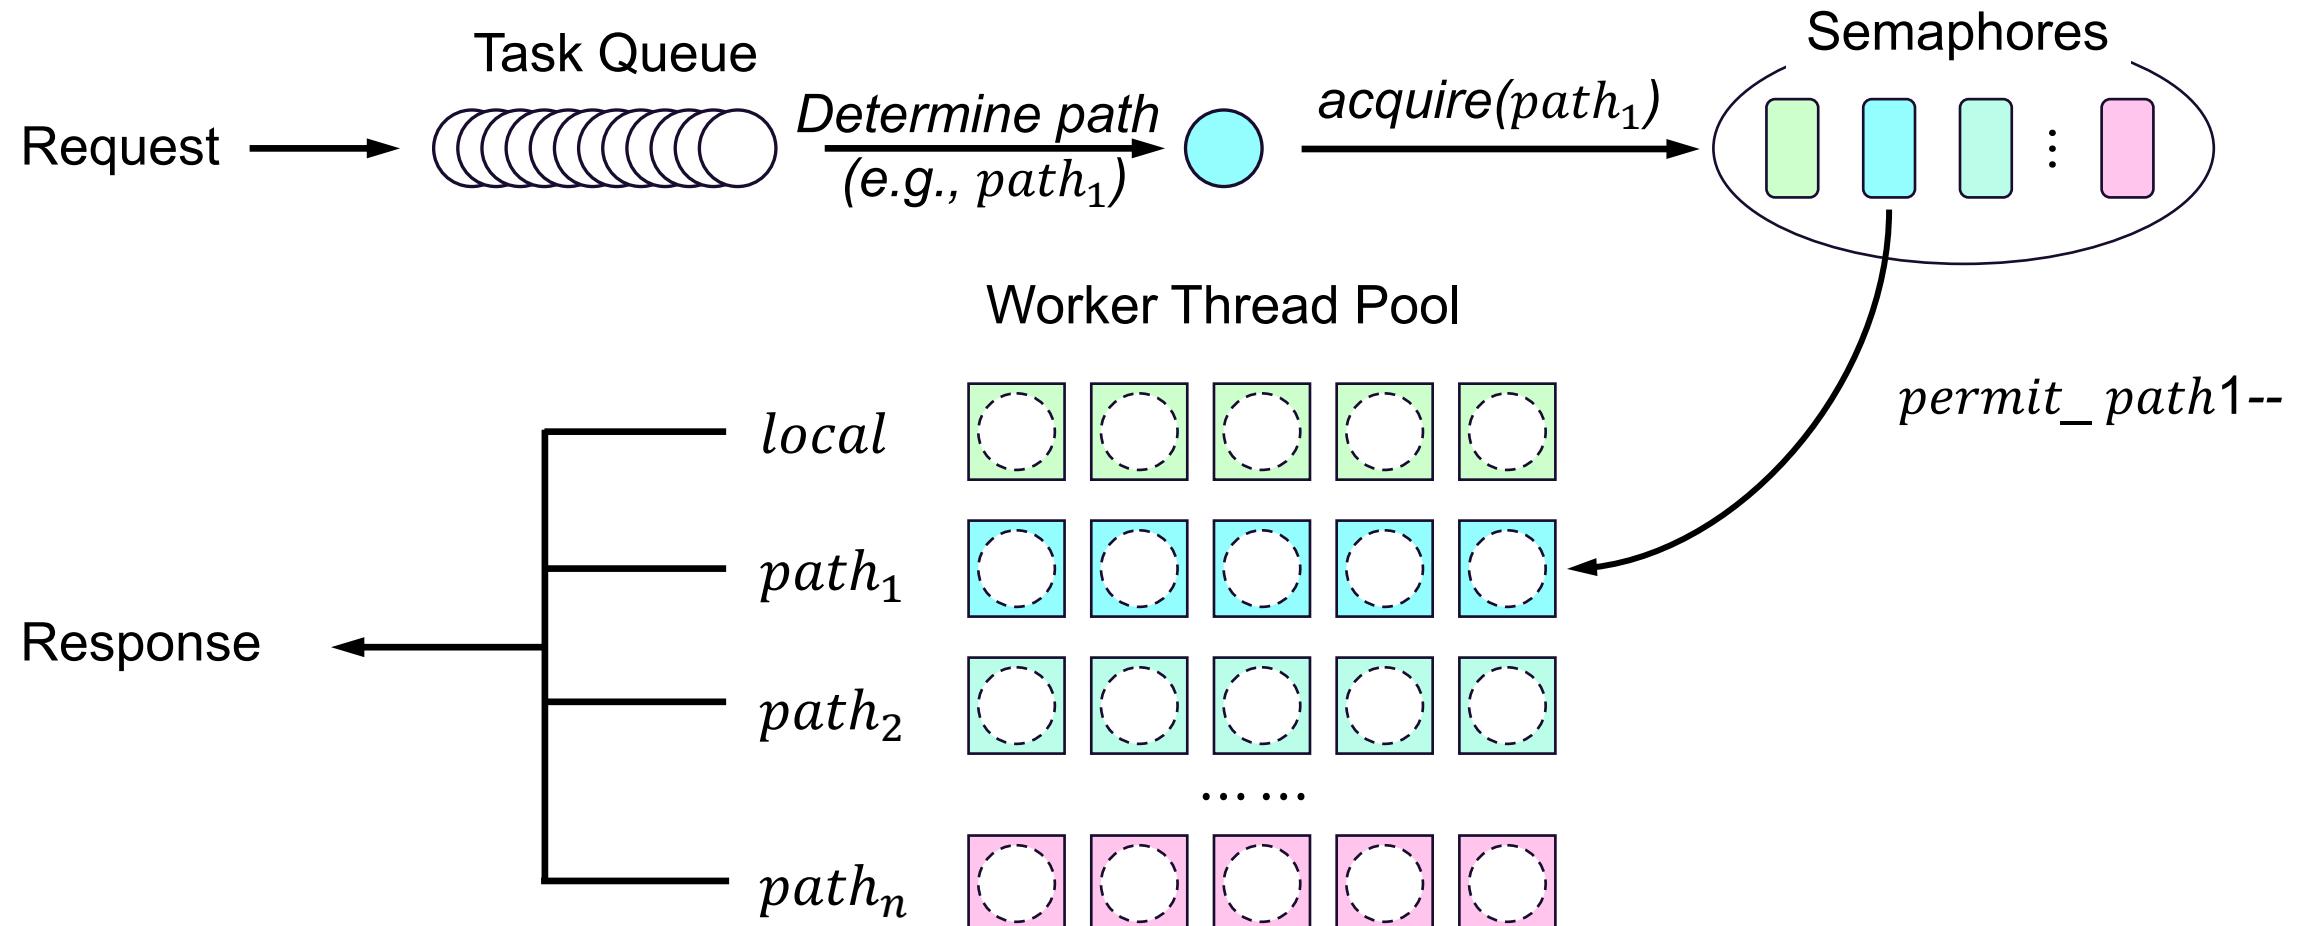
Semaphores-Based Thread Management



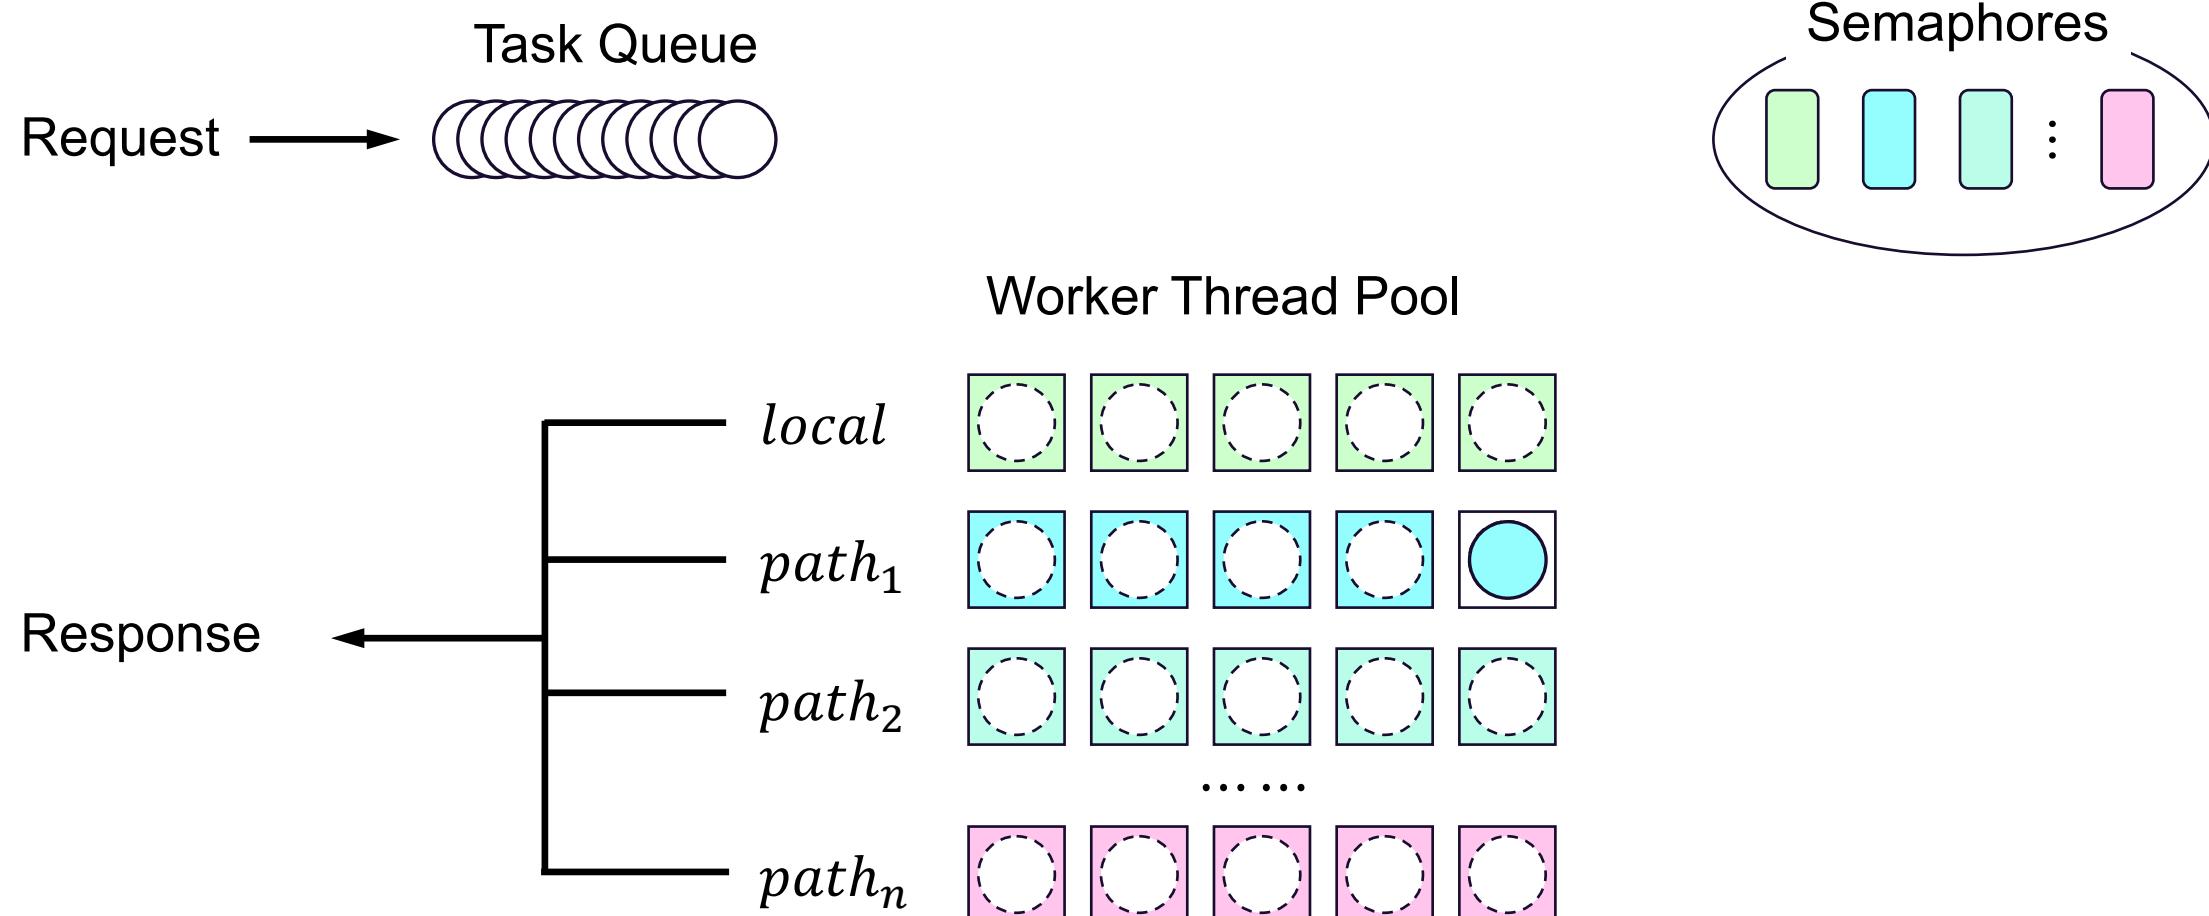
Semaphores-Based Thread Management



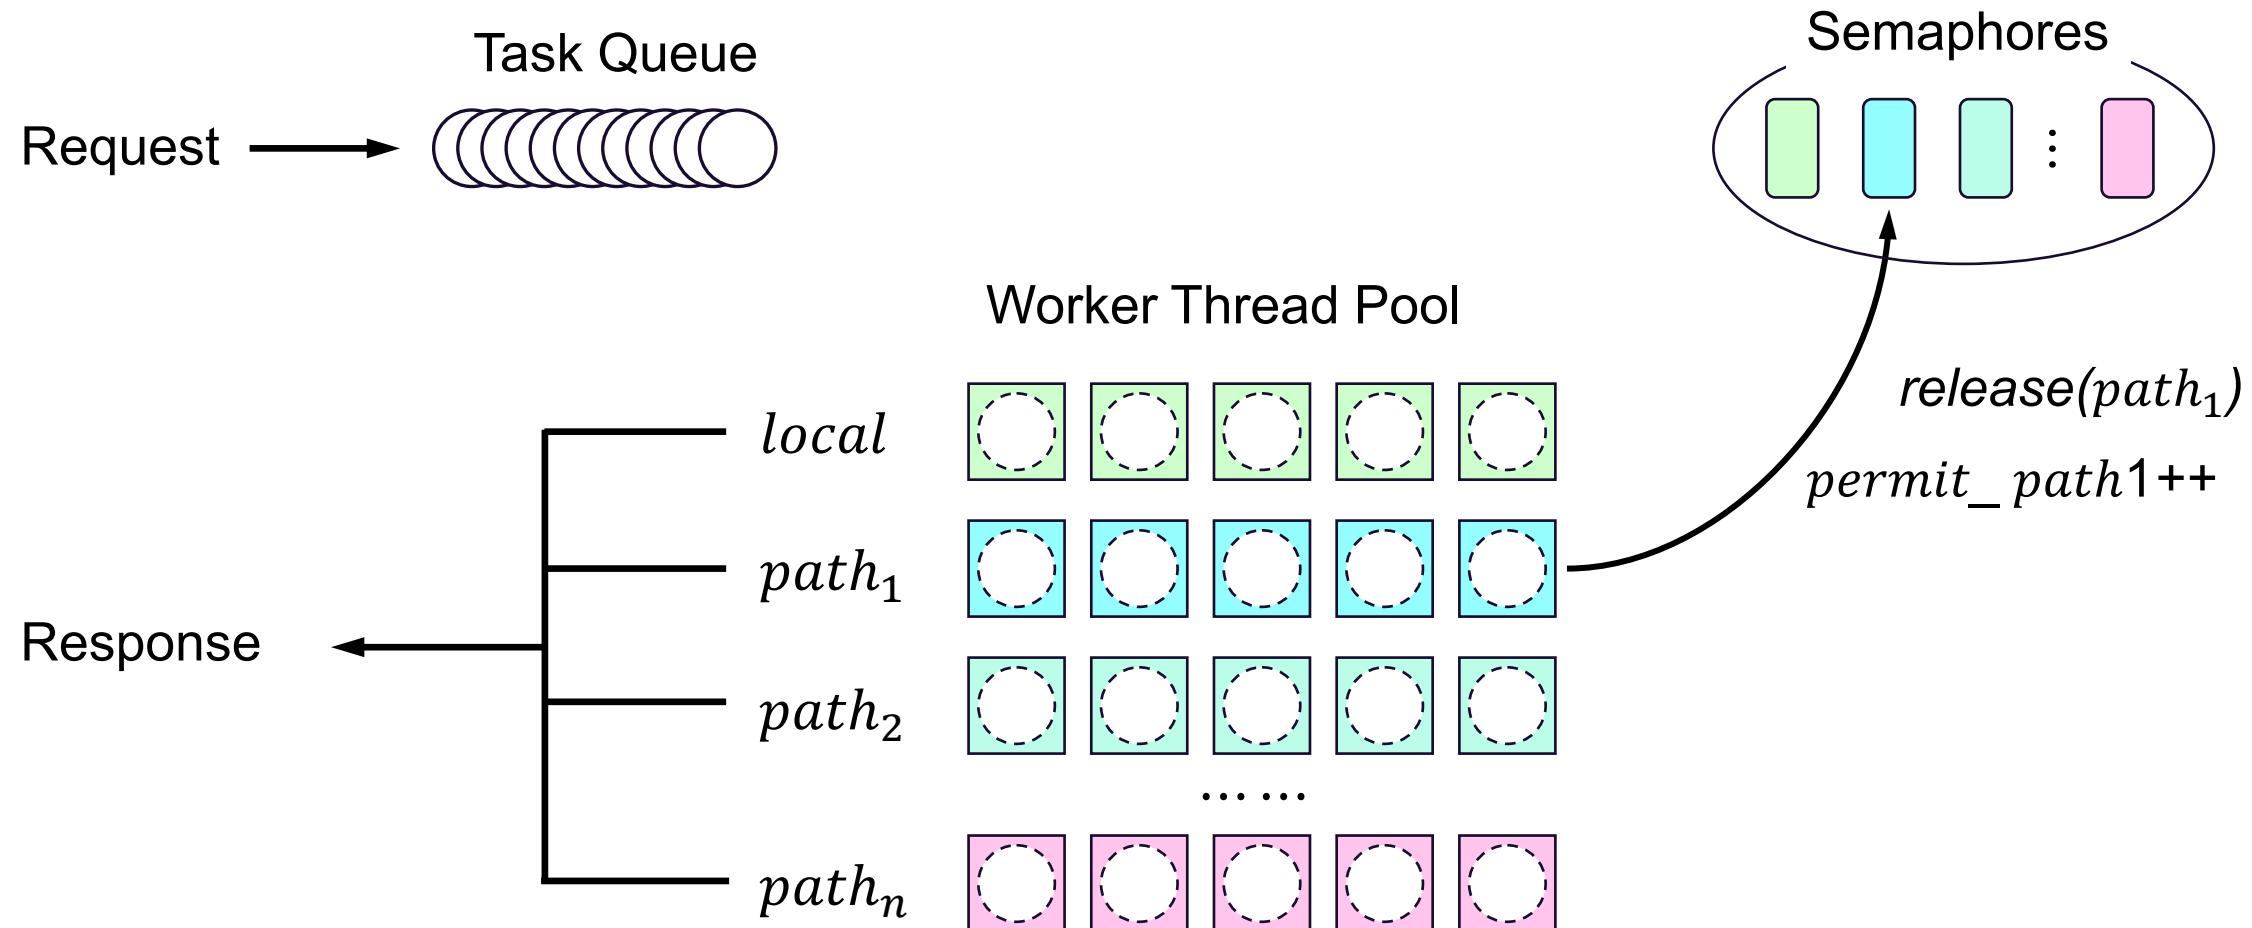
Semaphores-Based Thread Management



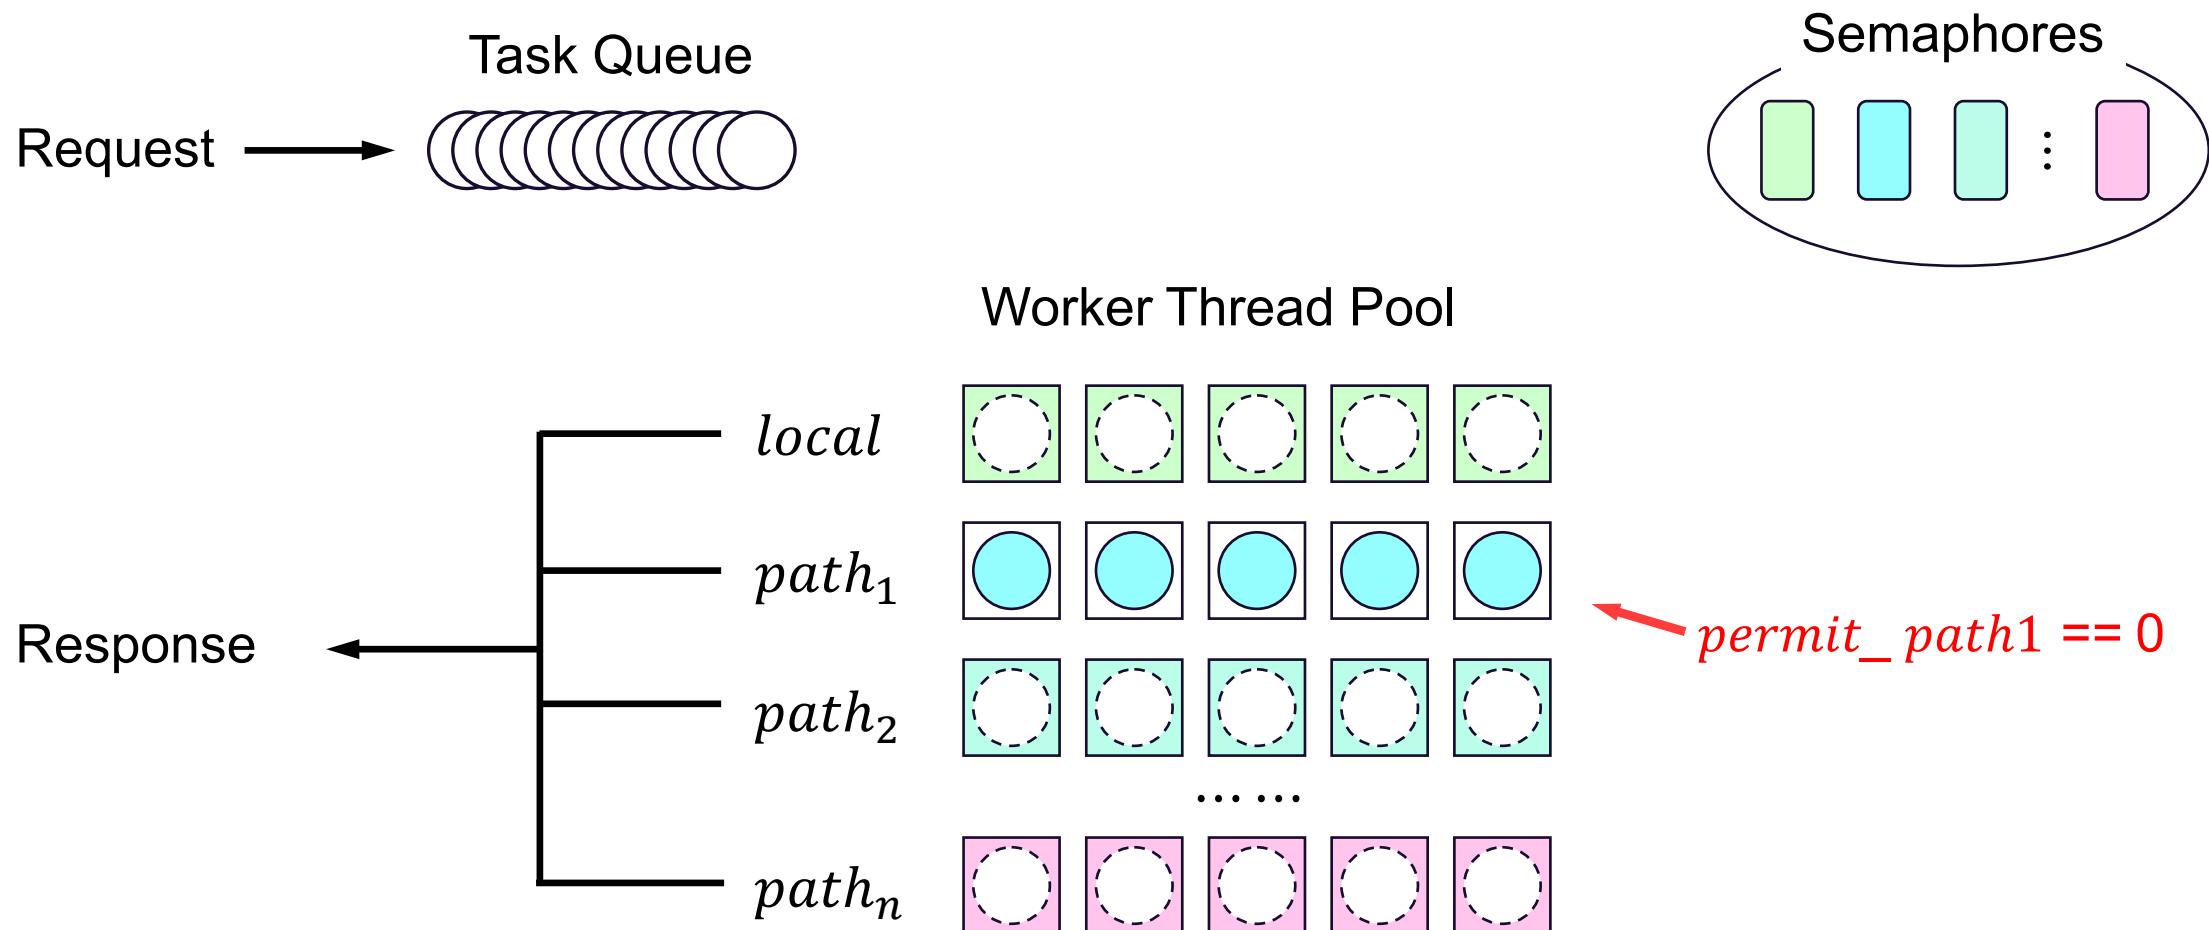
Semaphores-Based Thread Management



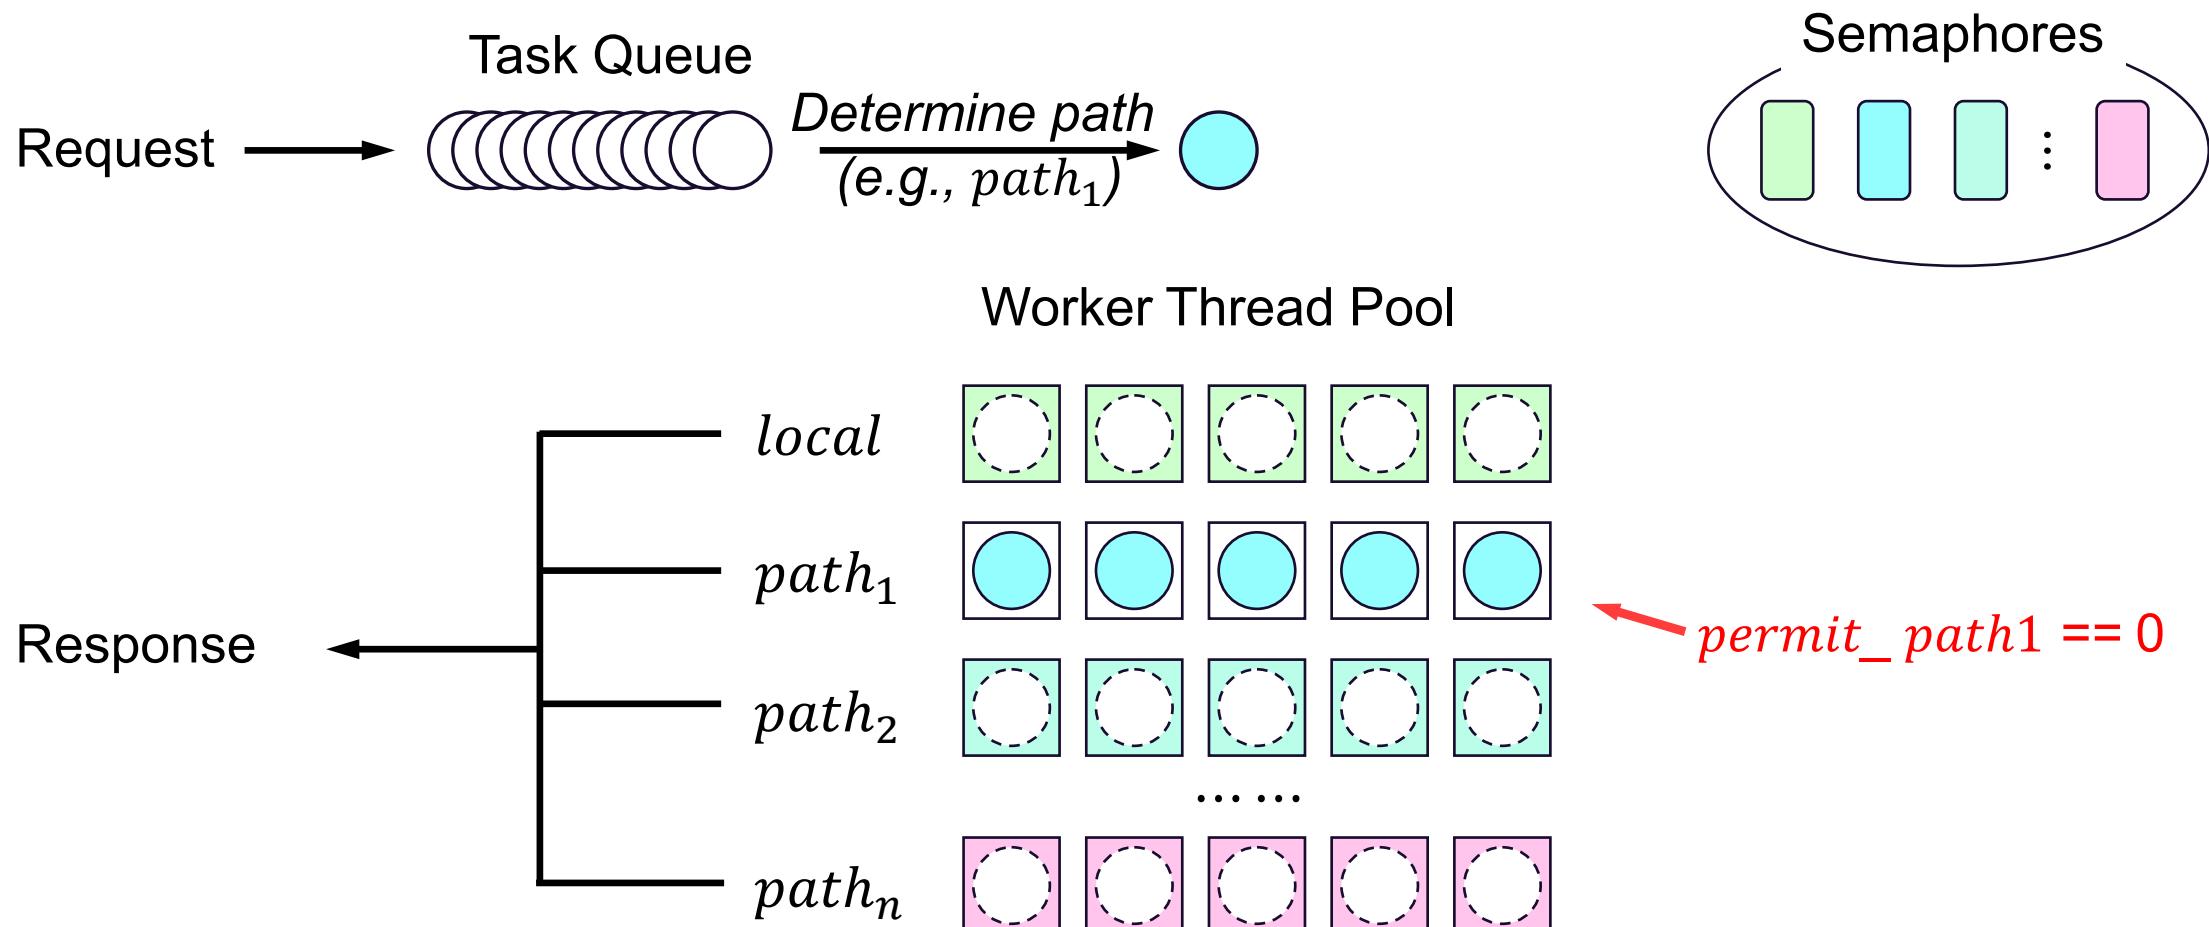
Semaphores-Based Thread Management



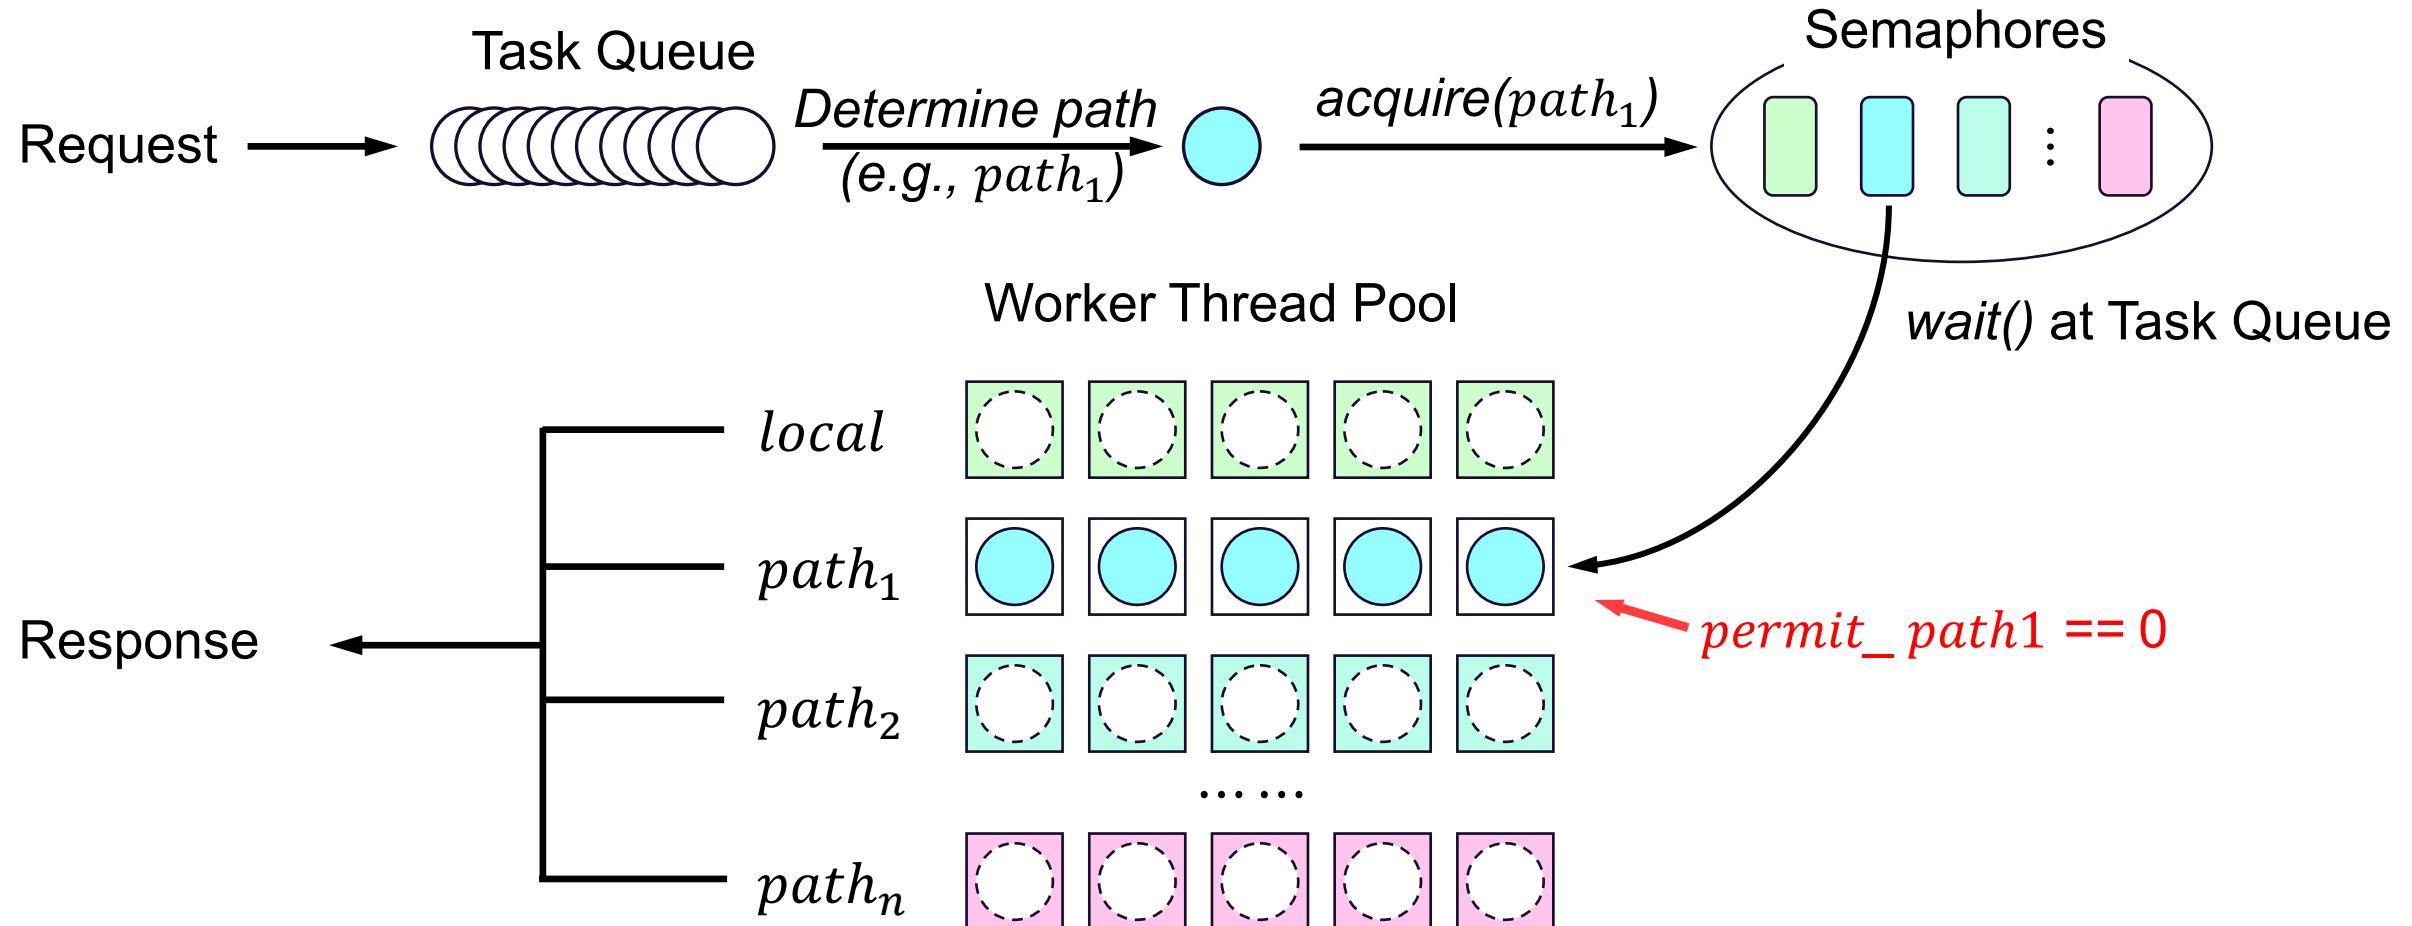
Eliminating Cross-Path Blocking



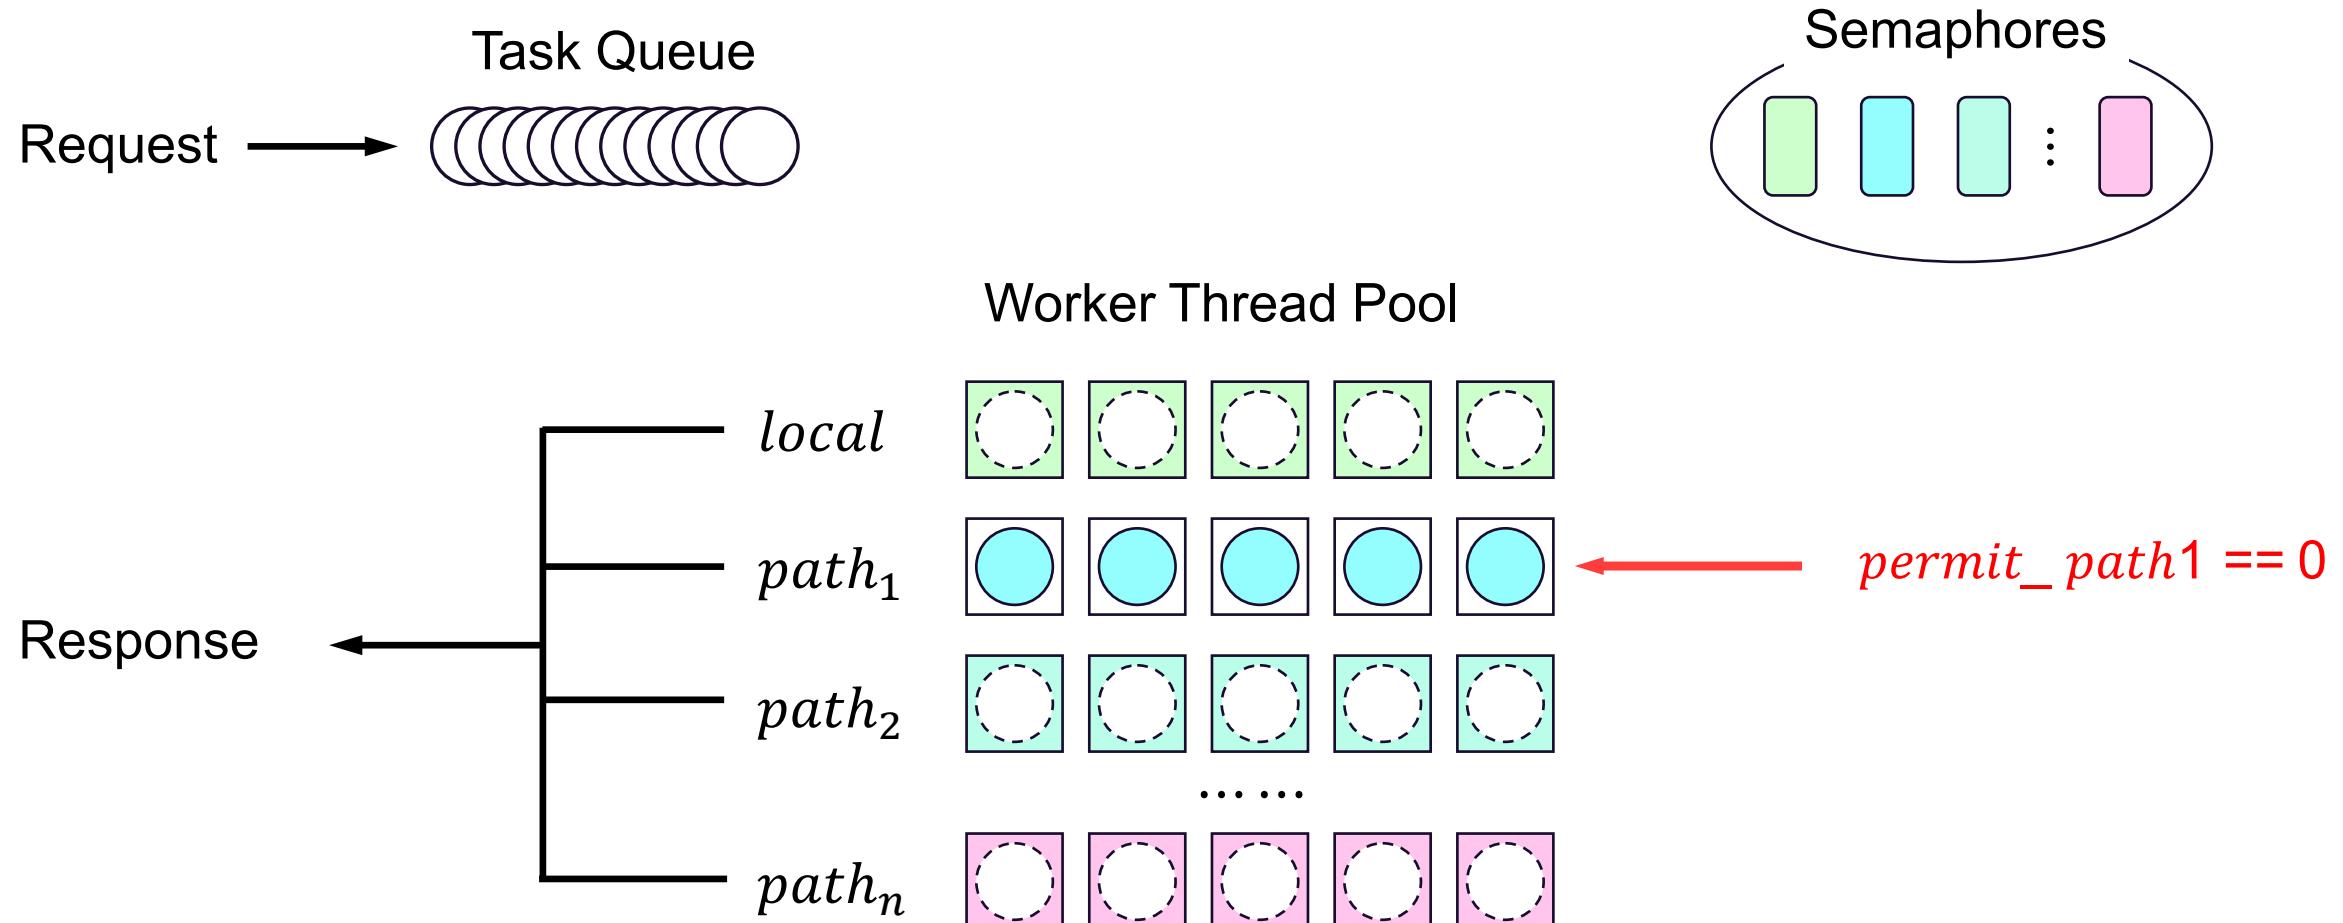
Eliminating Cross-Path Blocking



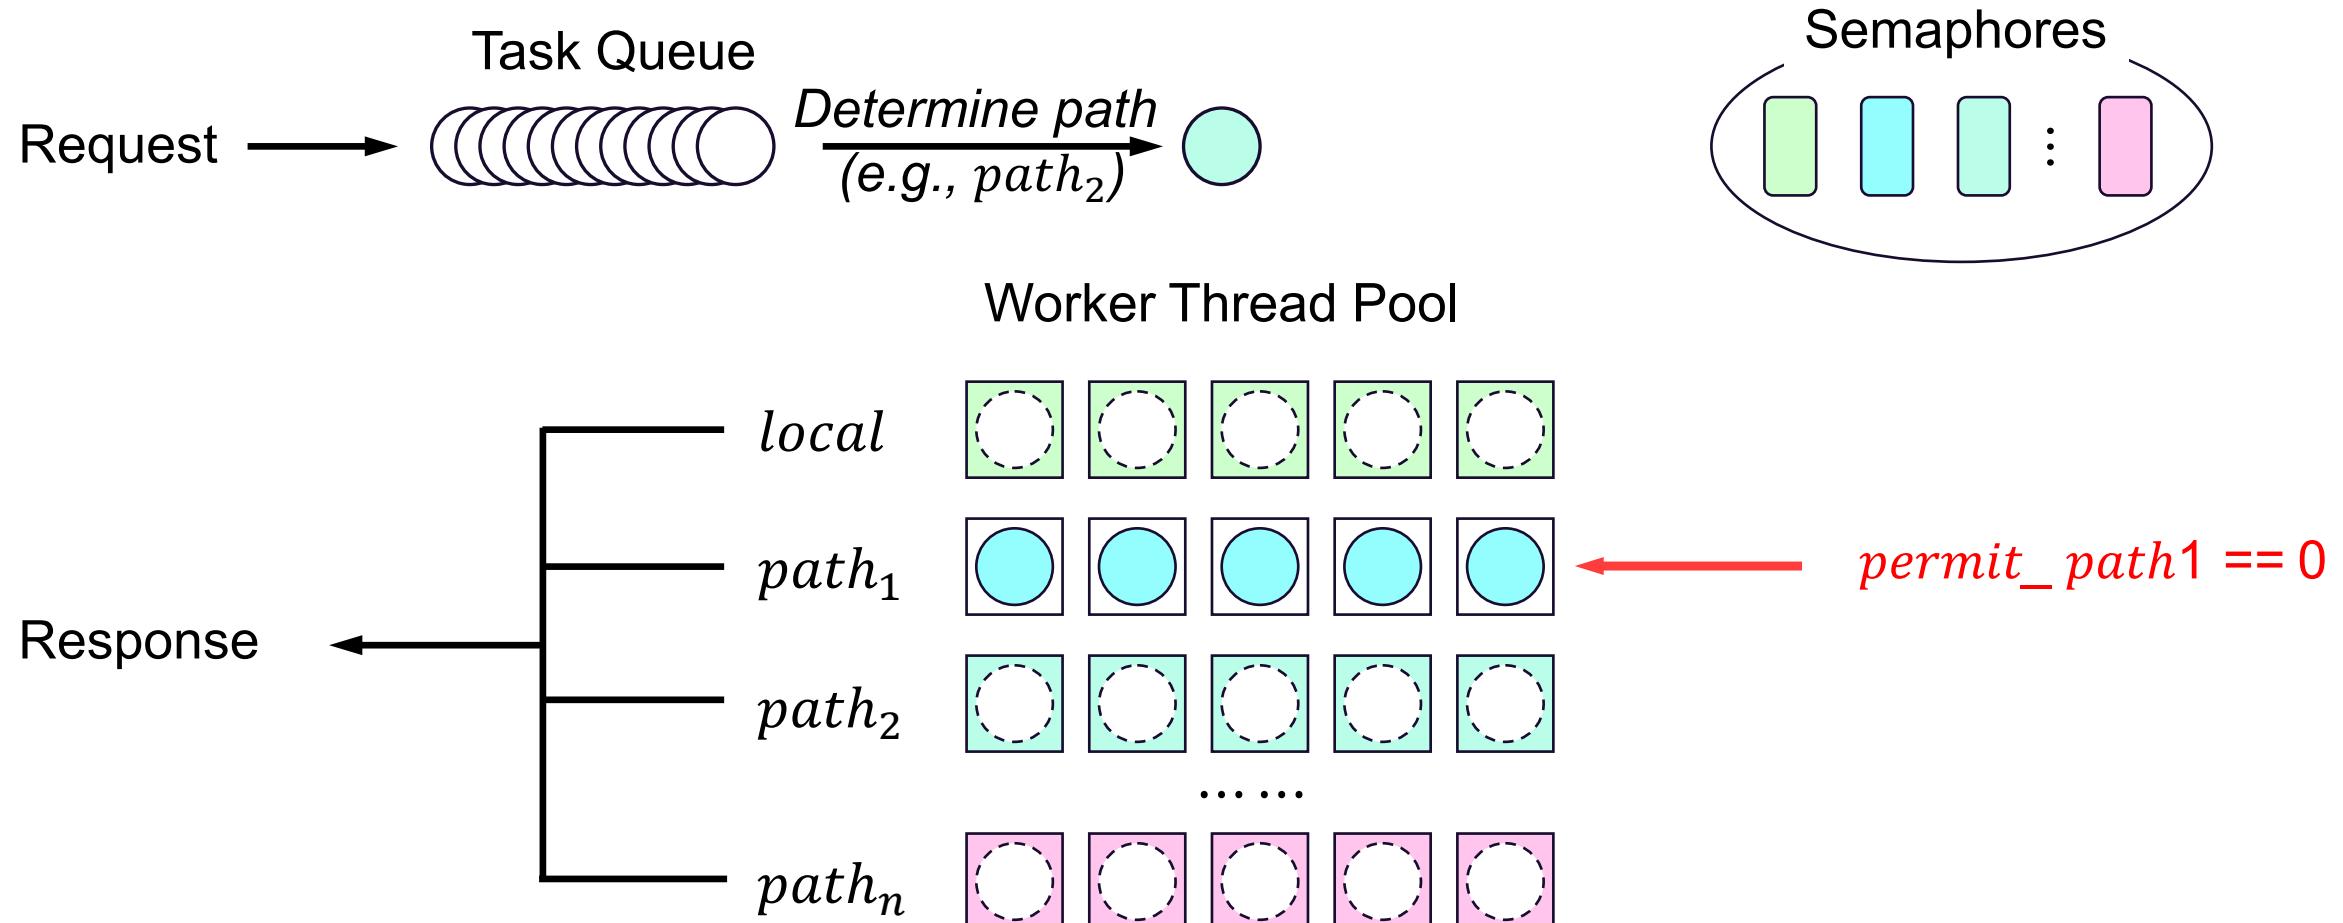
Eliminating Cross-Path Blocking



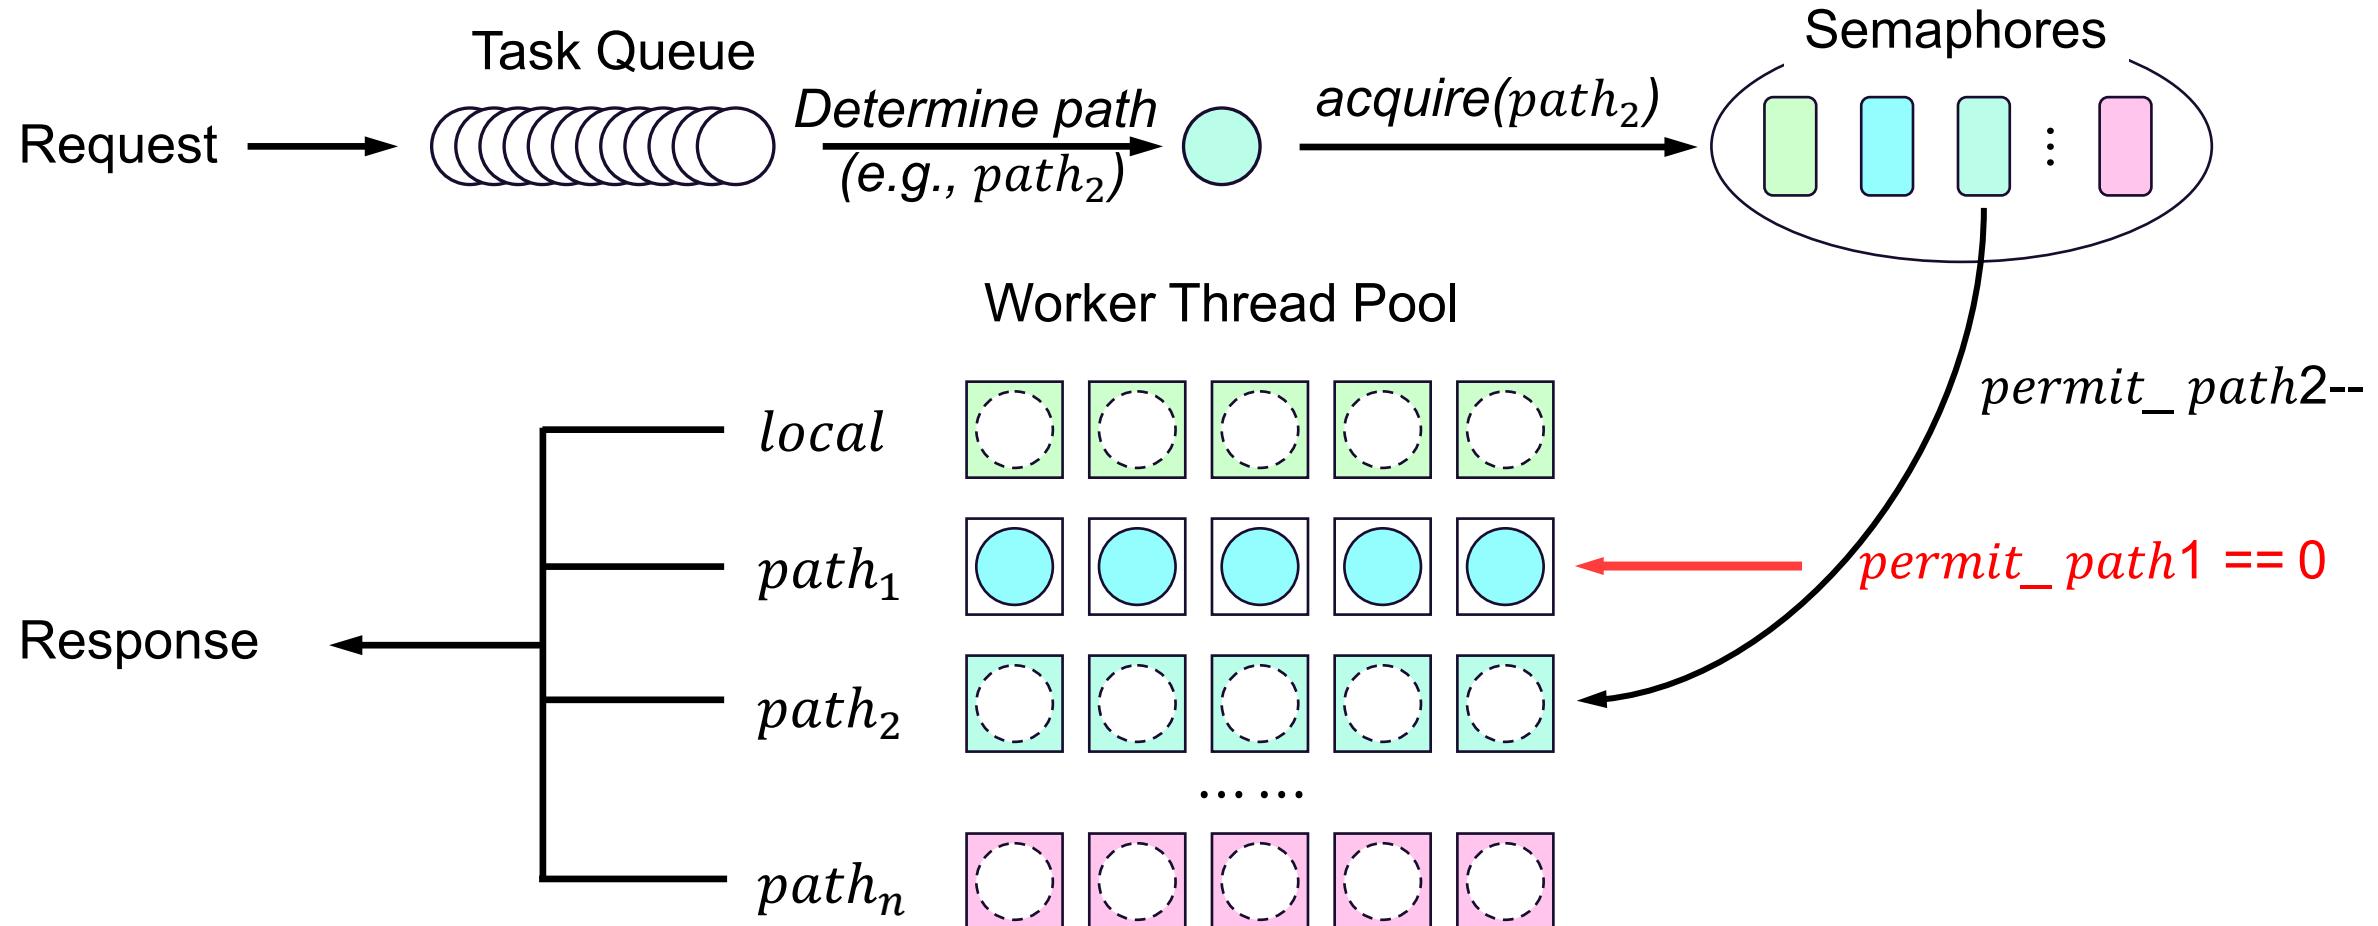
Eliminating Cross-Path Blocking



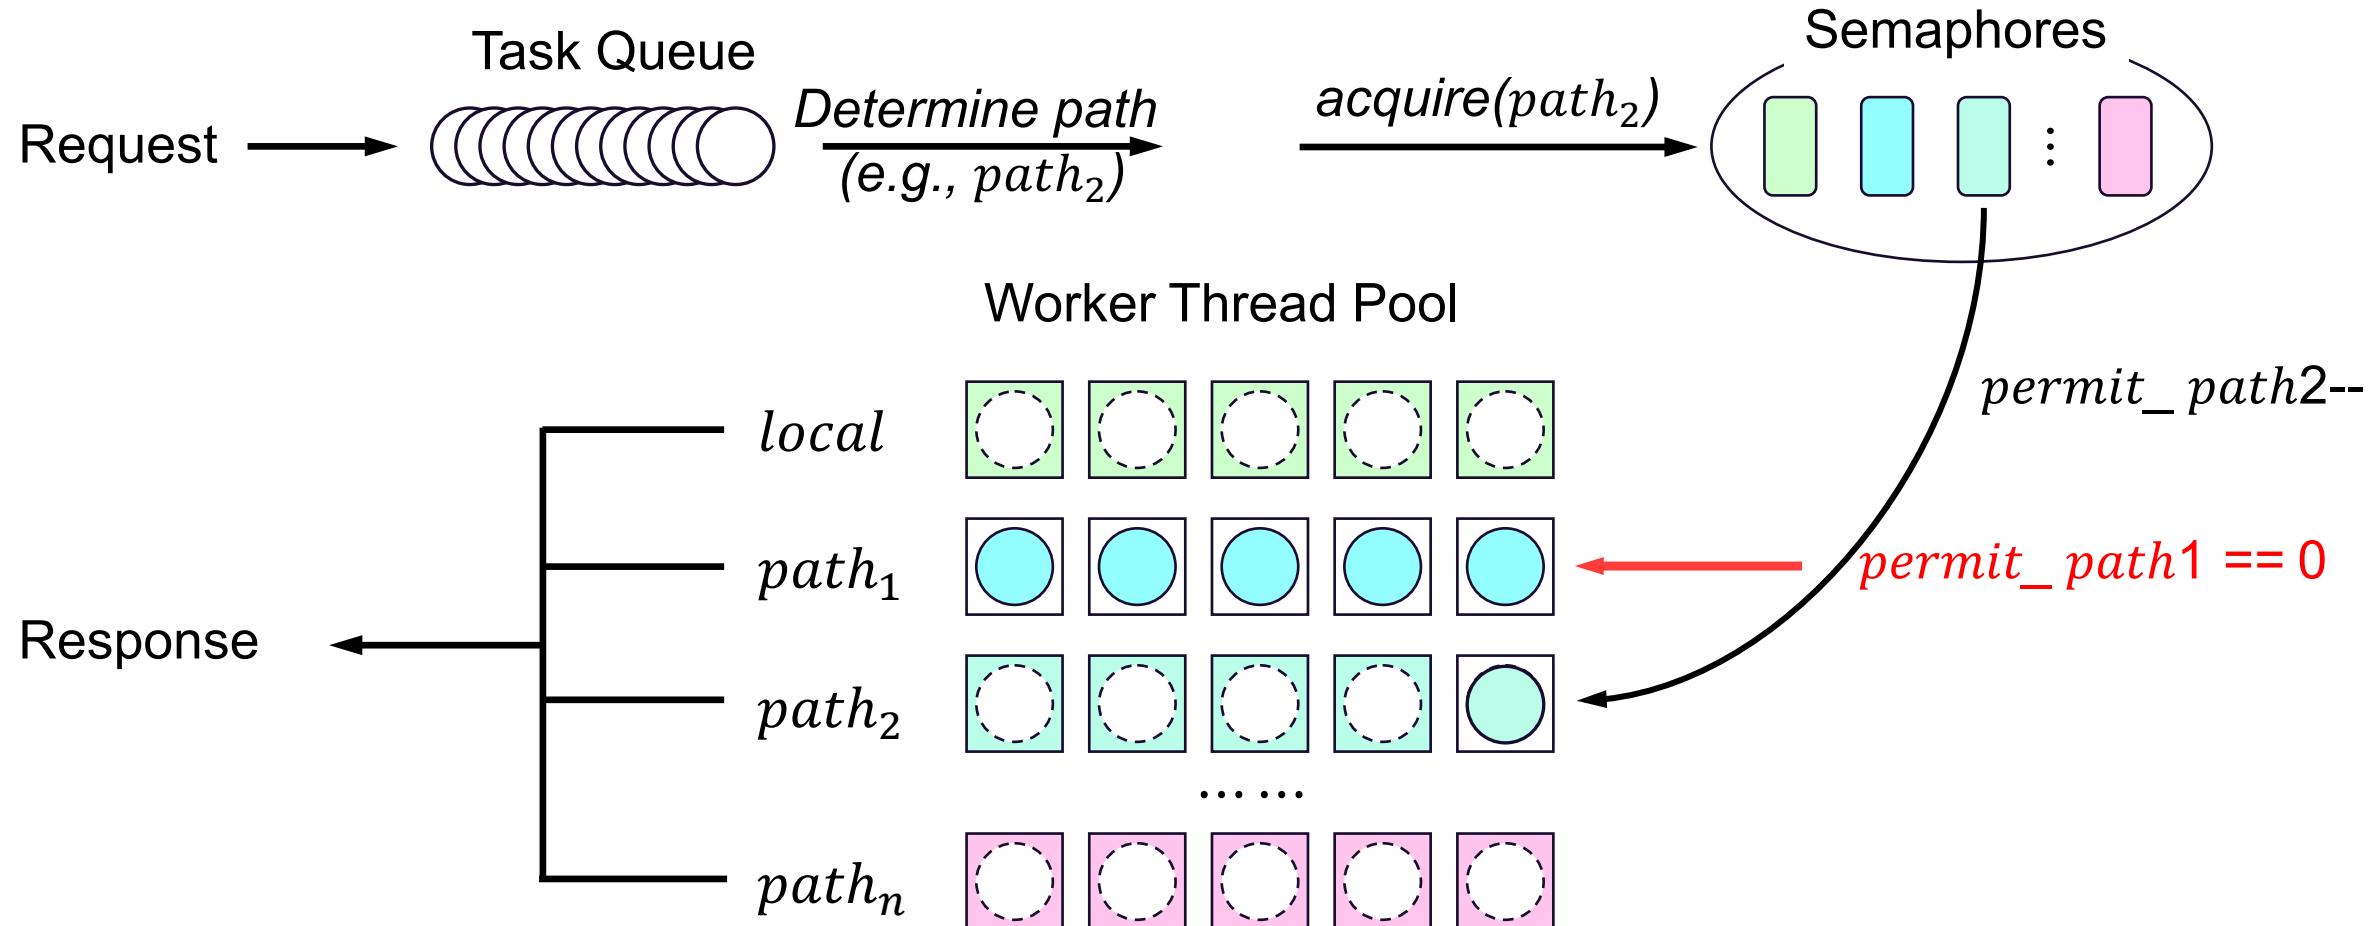
Eliminating Cross-Path Blocking



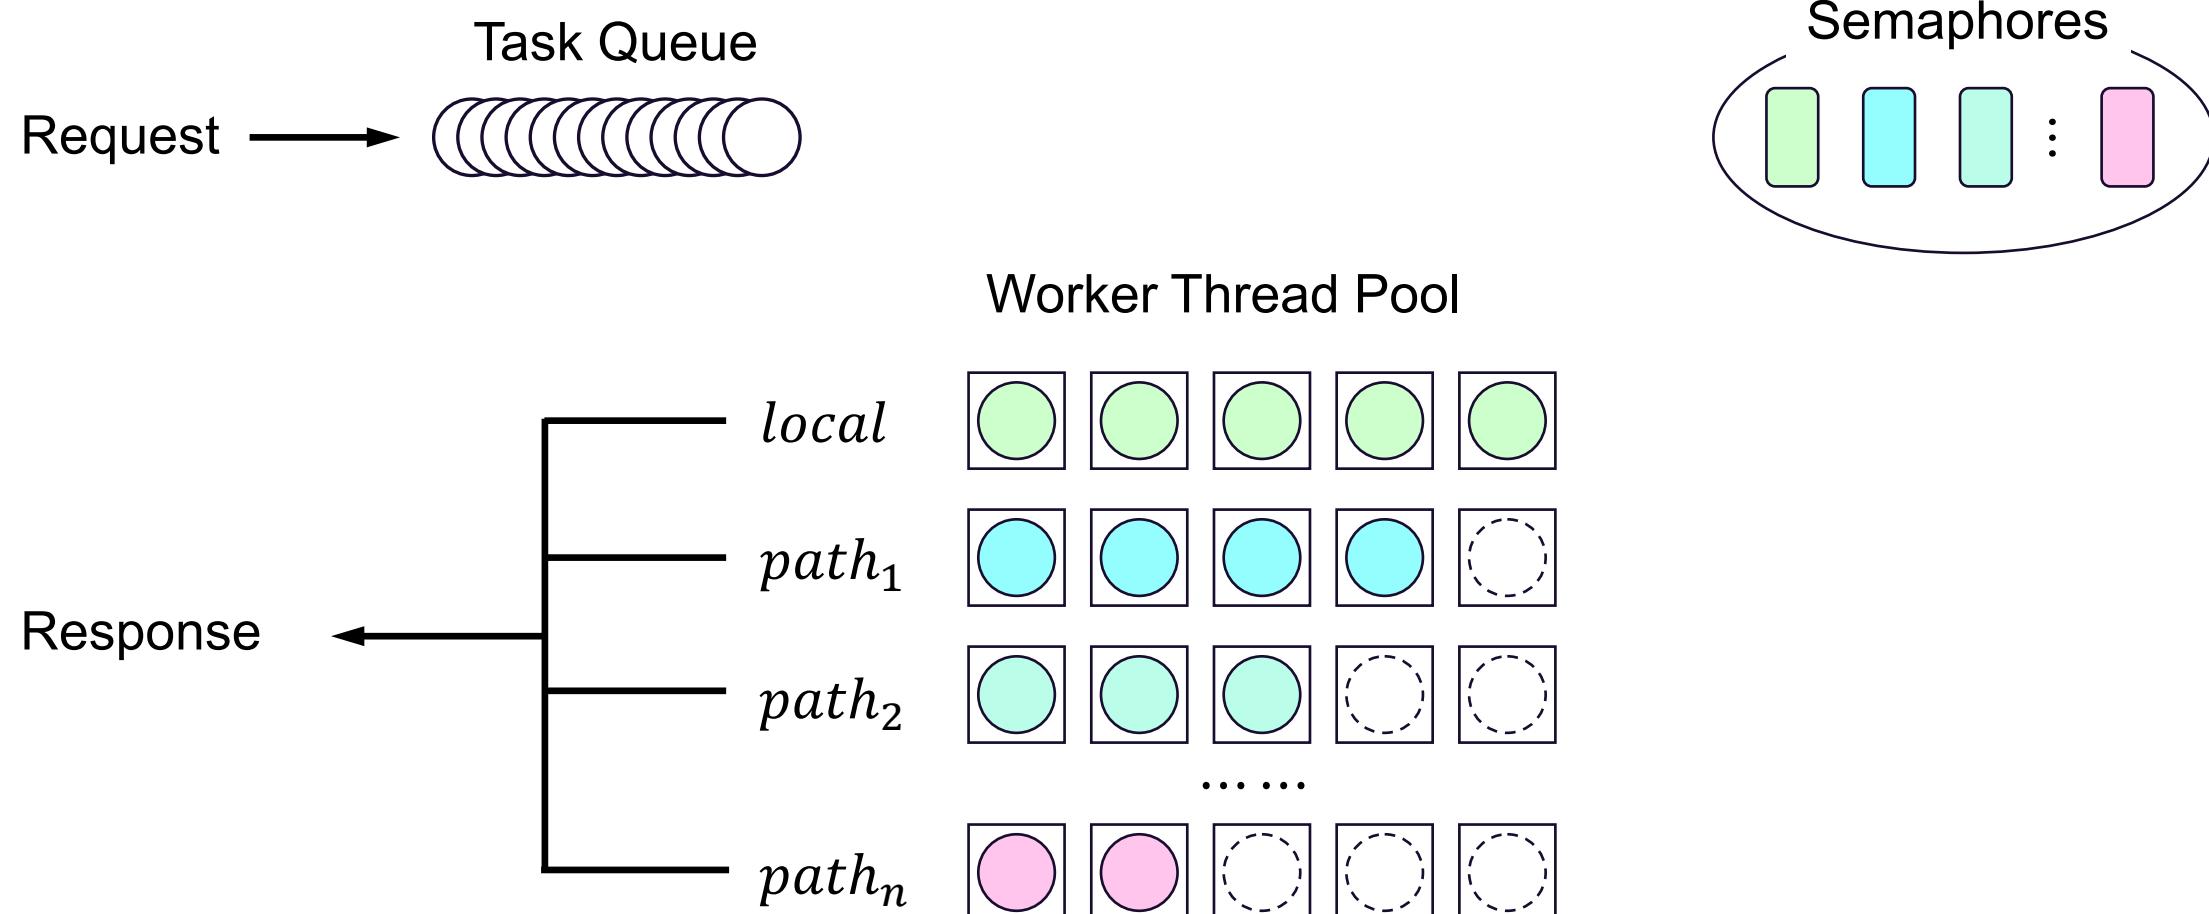
Eliminating Cross-Path Blocking



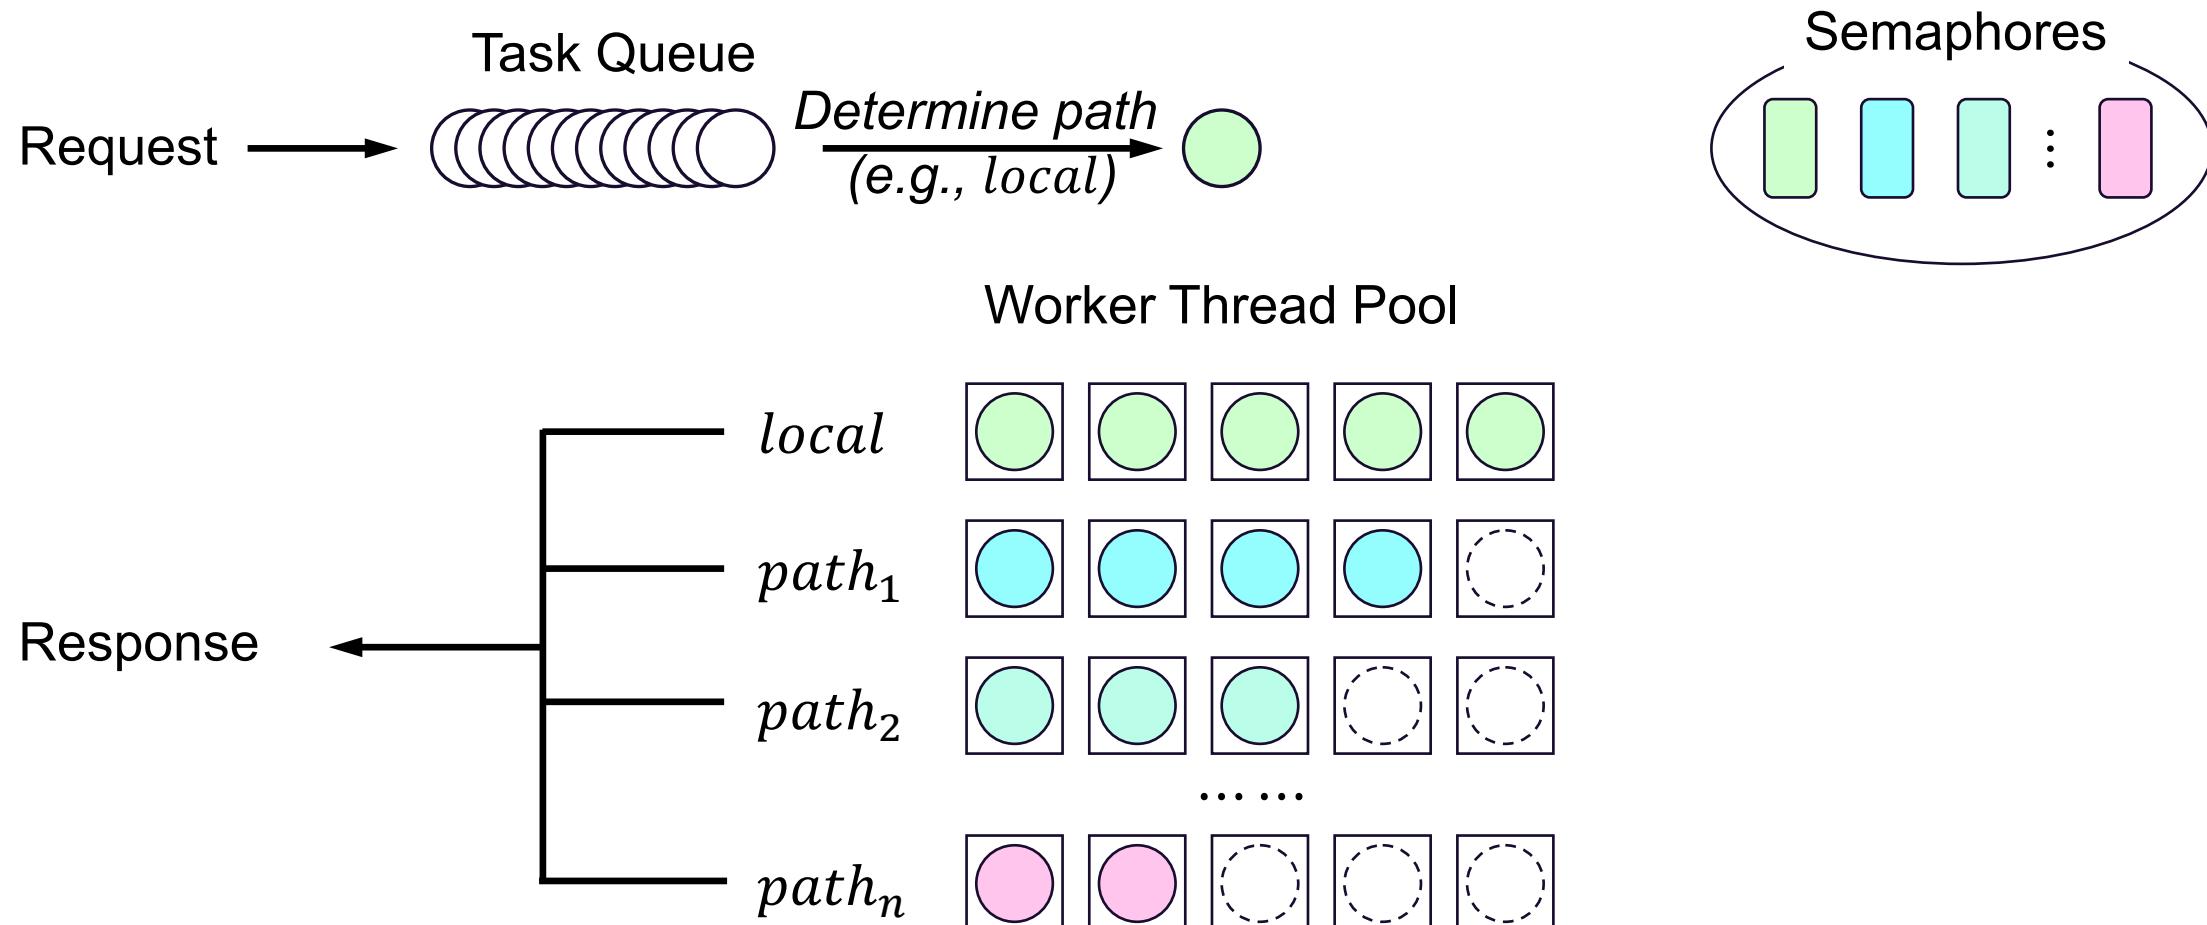
Eliminating Cross-Path Blocking



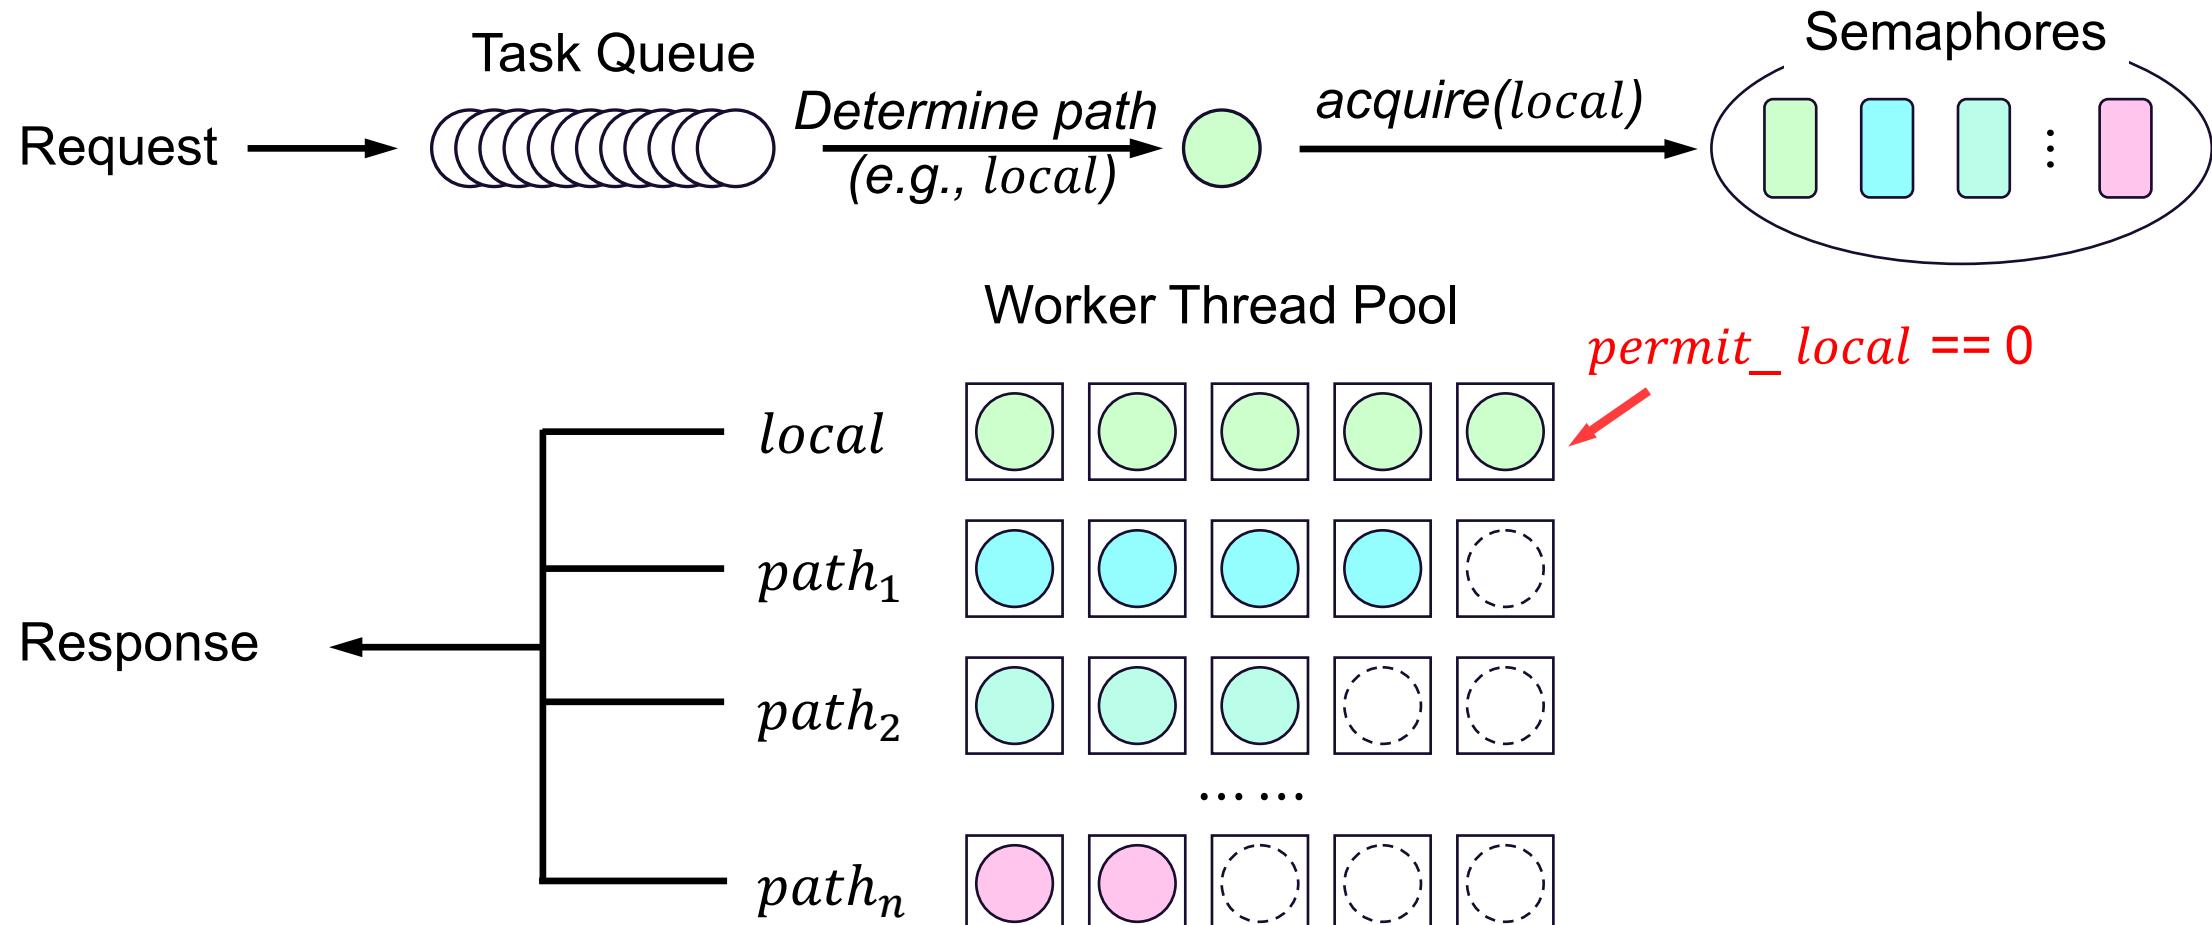
Prioritize Local Requests



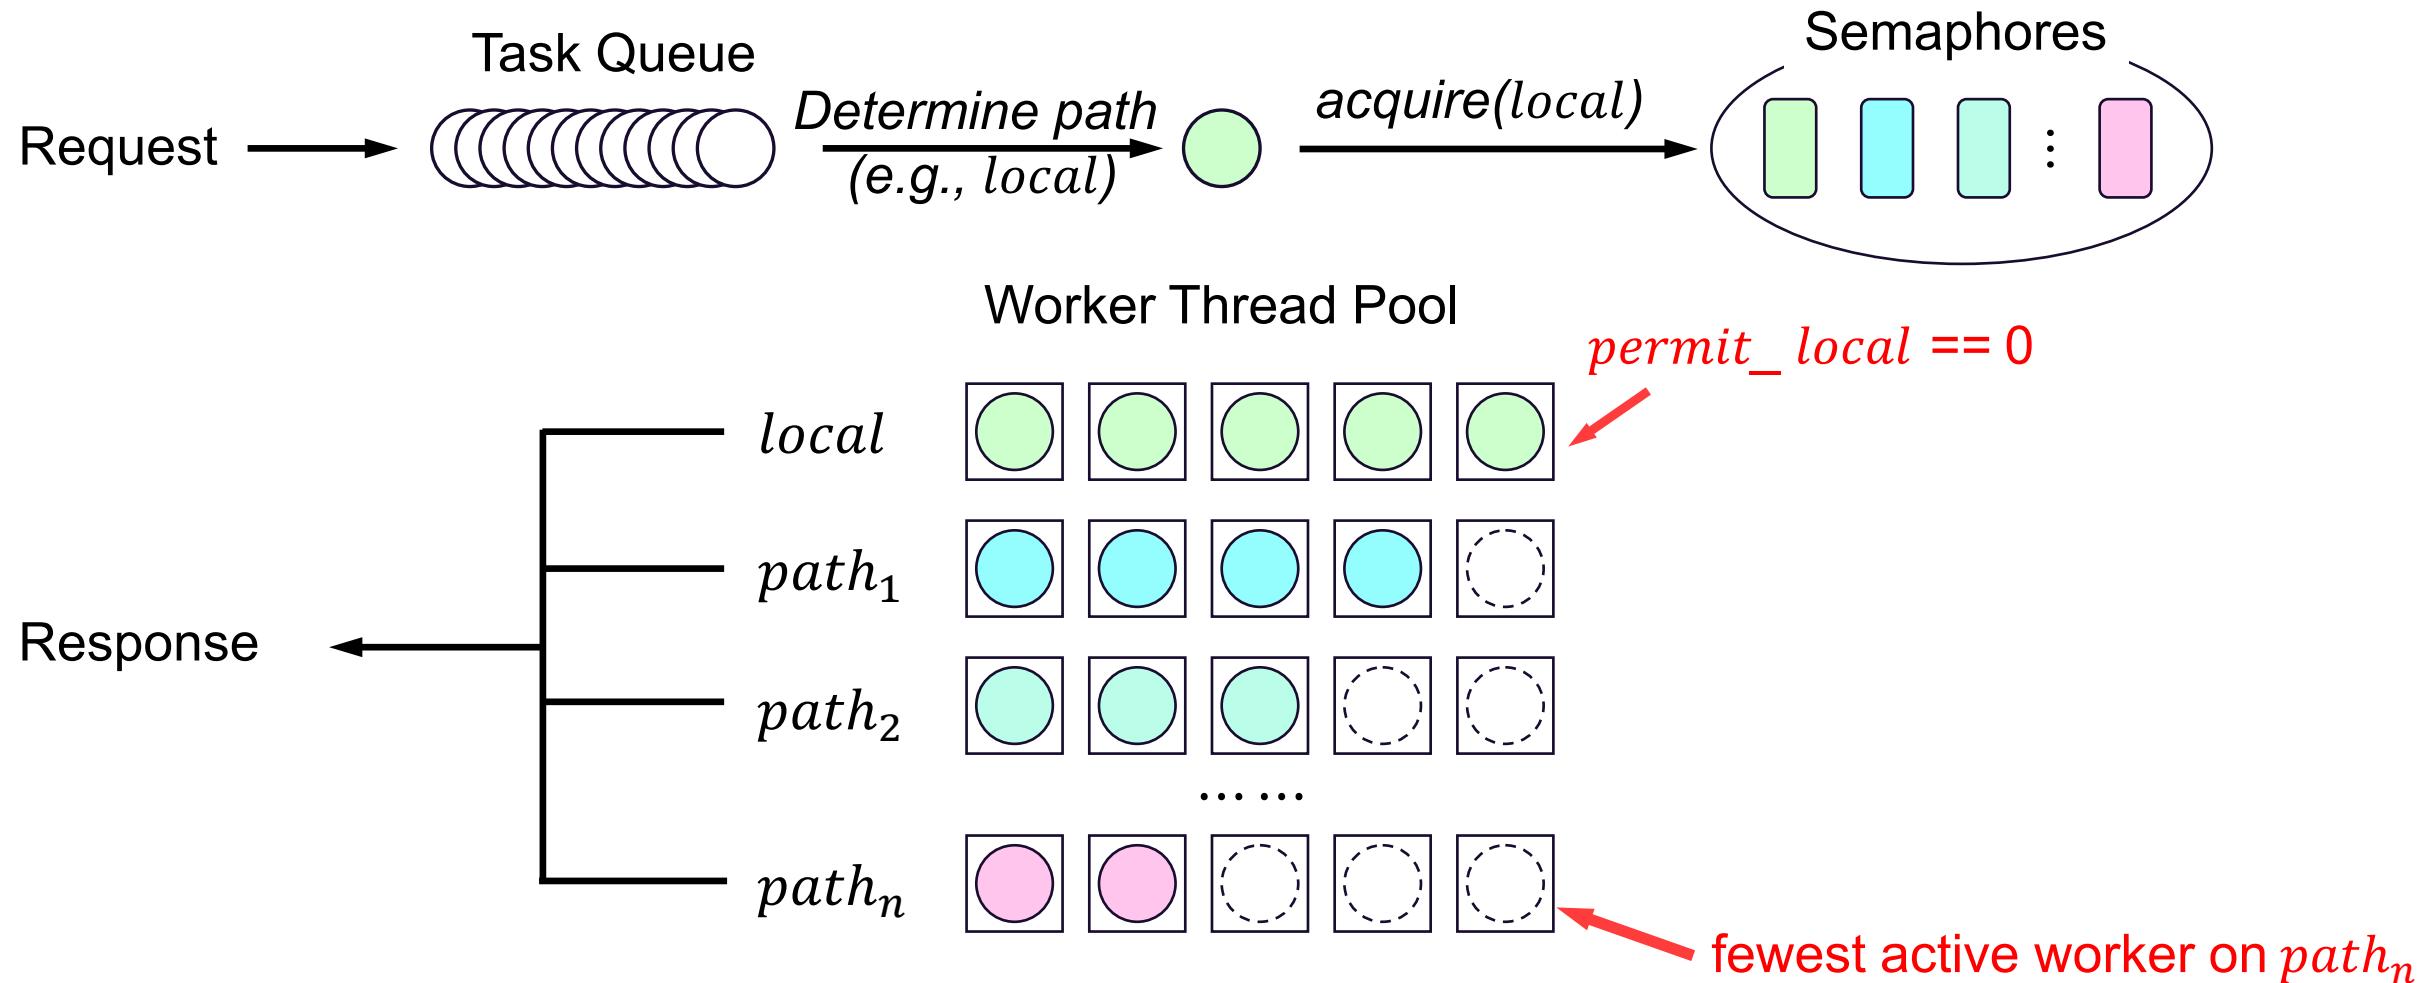
Prioritize Local Requests



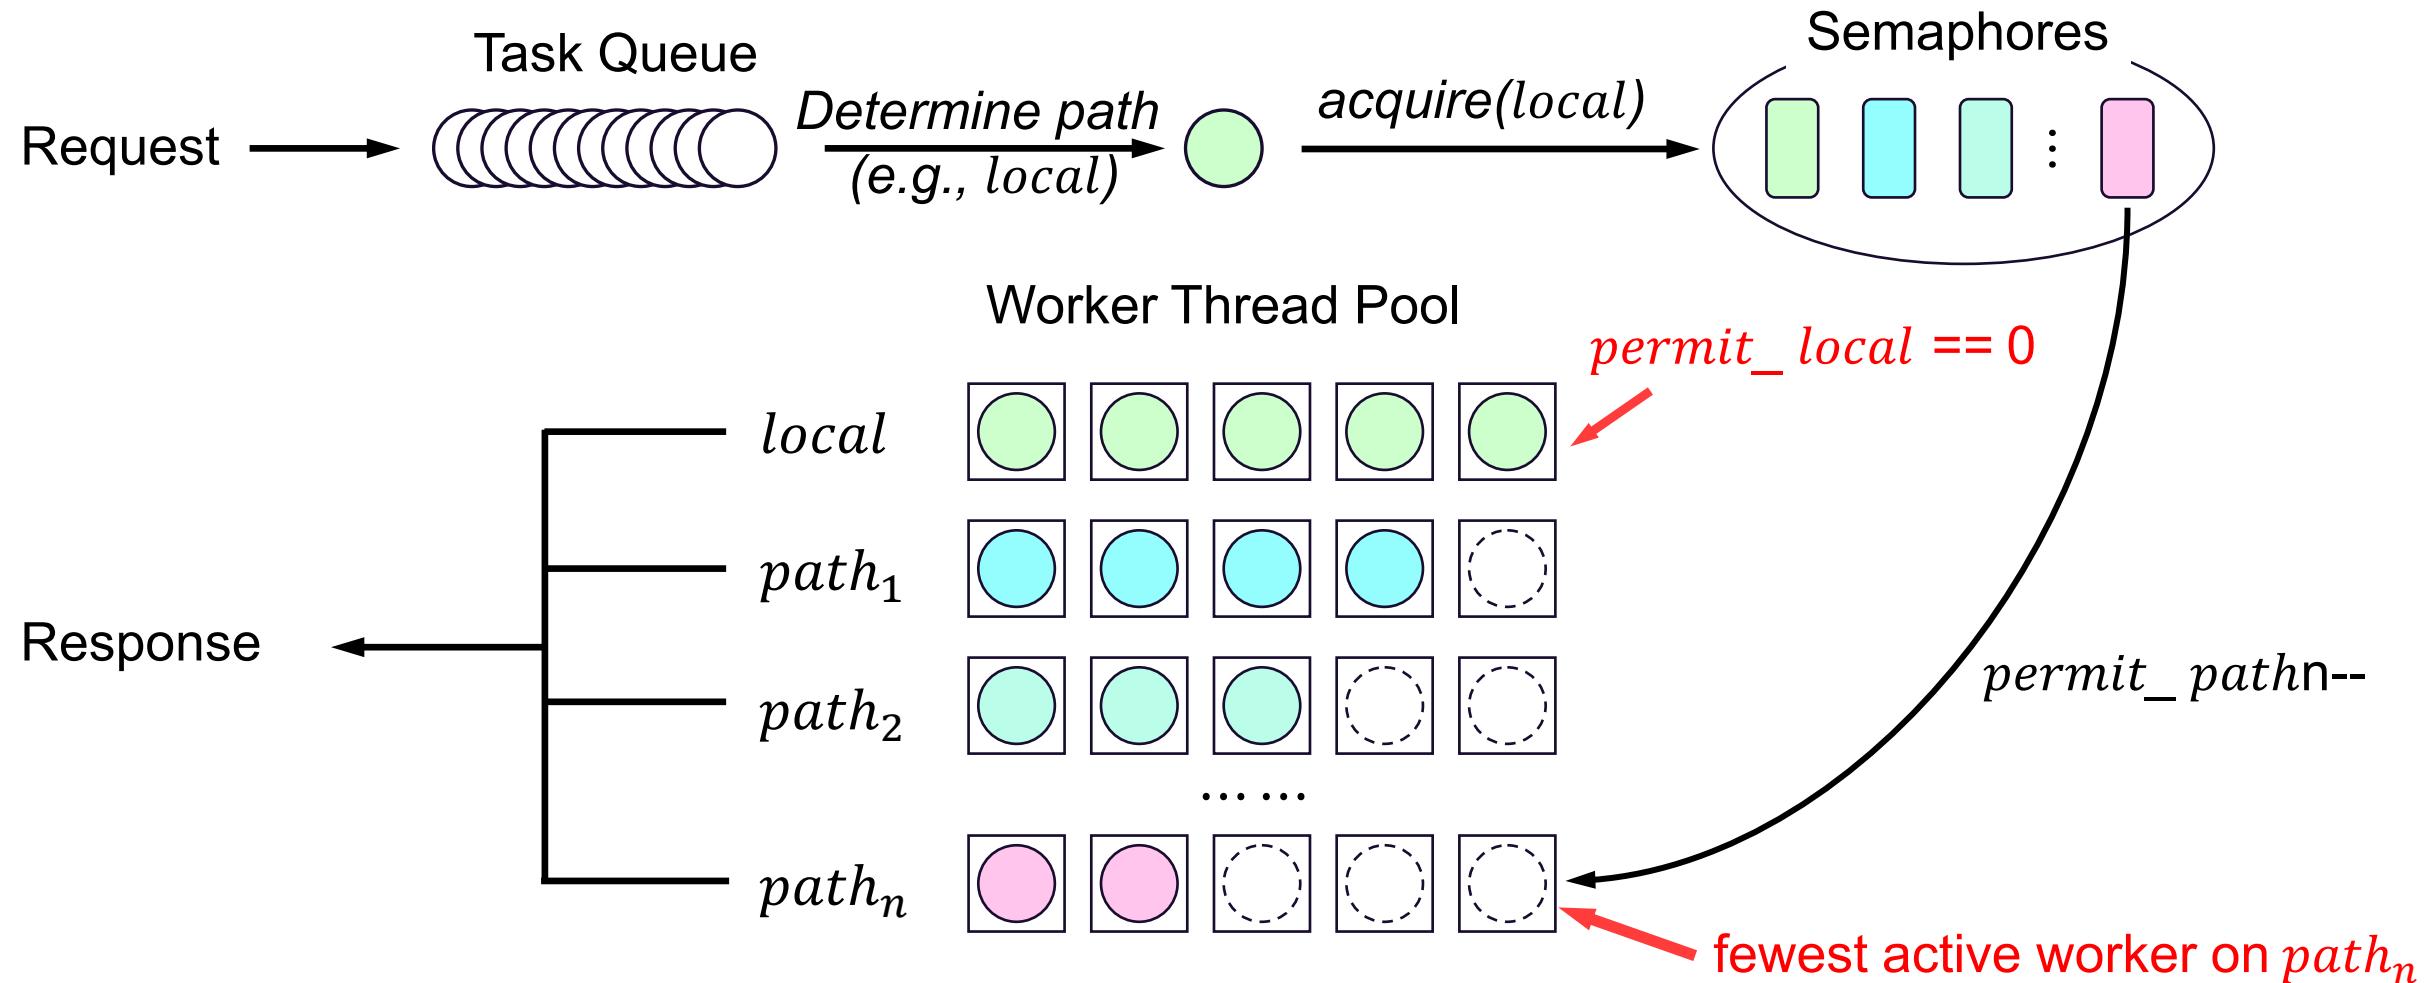
Prioritize Local Requests



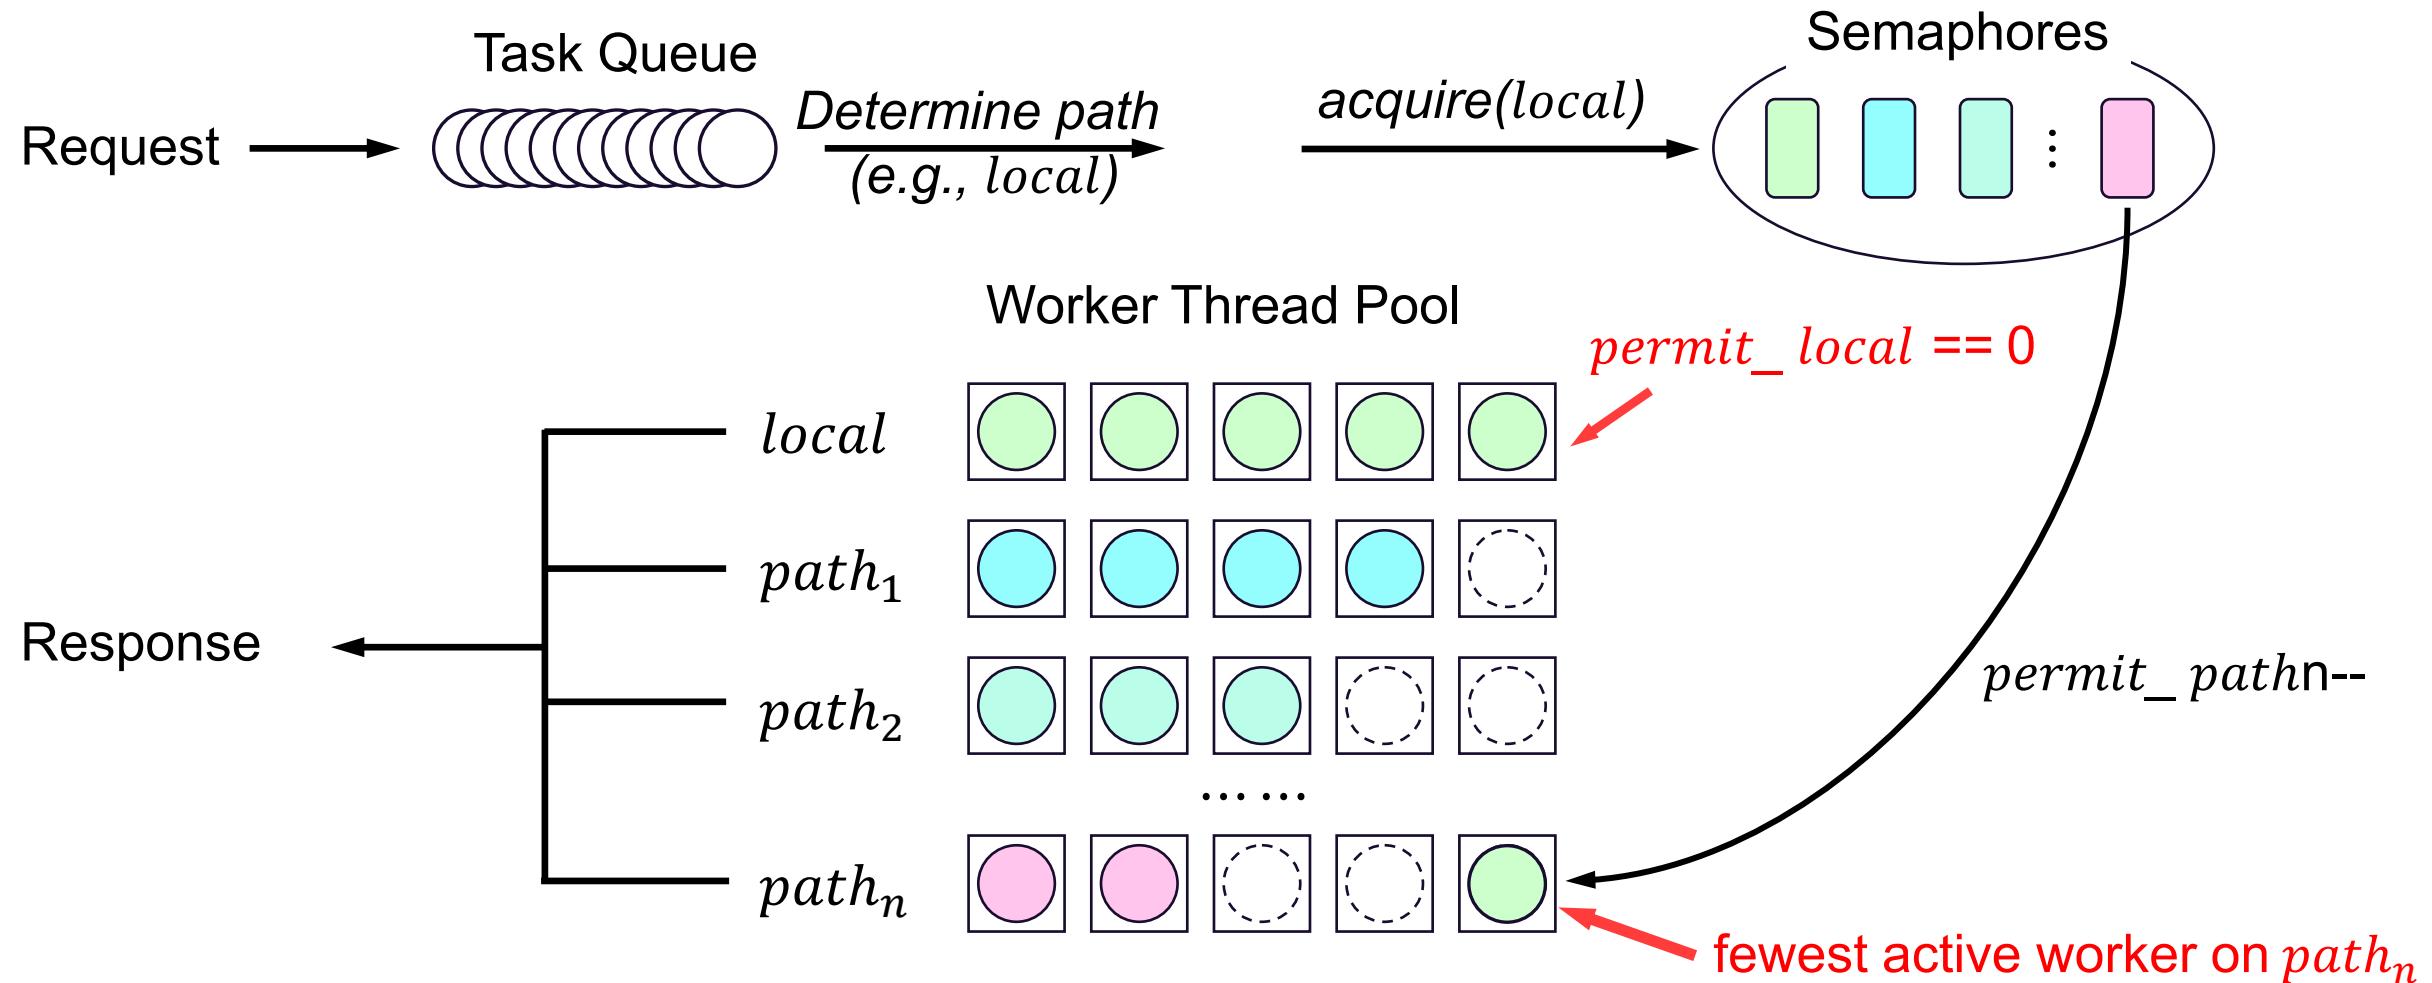
Prioritize Local Requests



Prioritize Local Requests

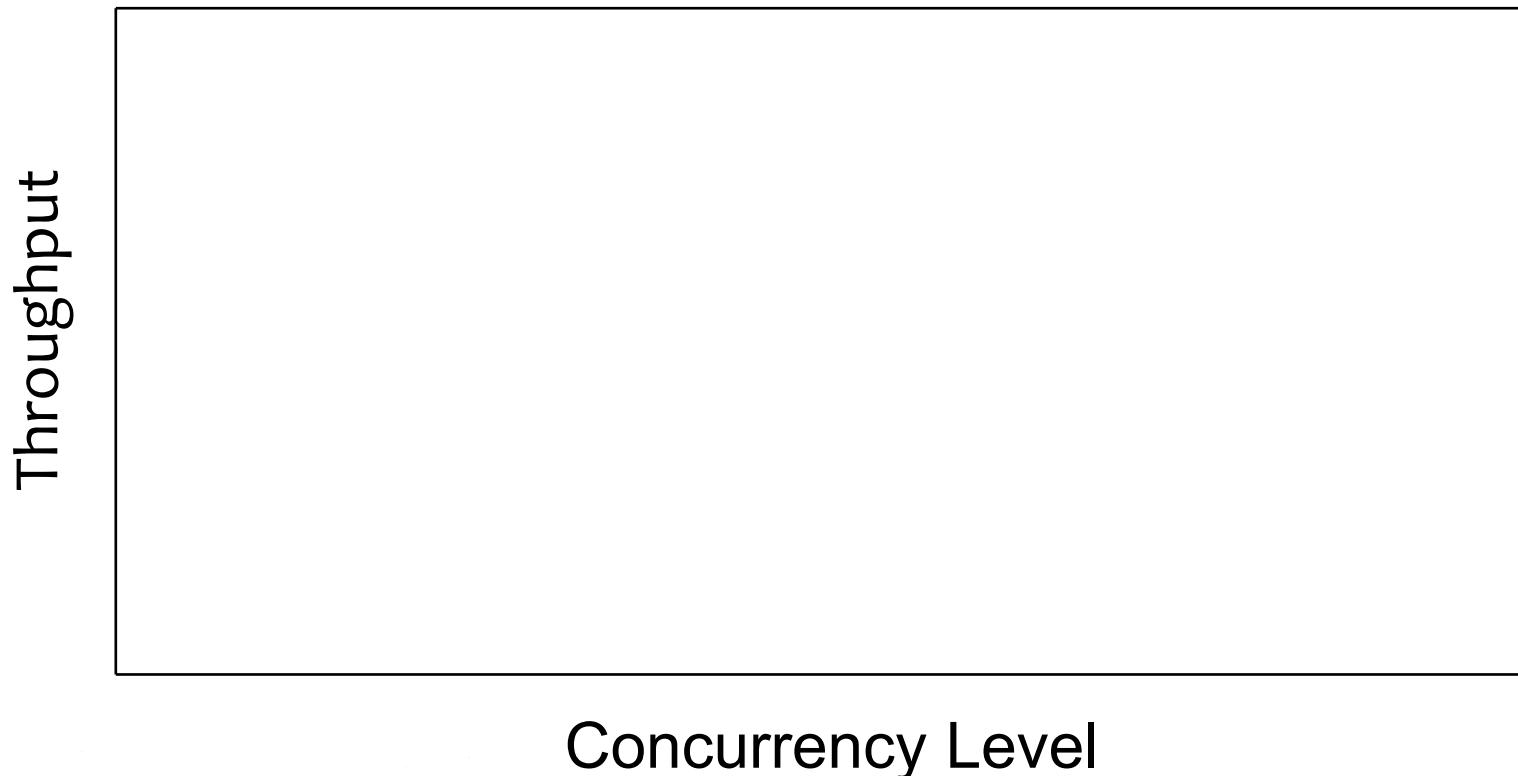


Prioritize Local Requests



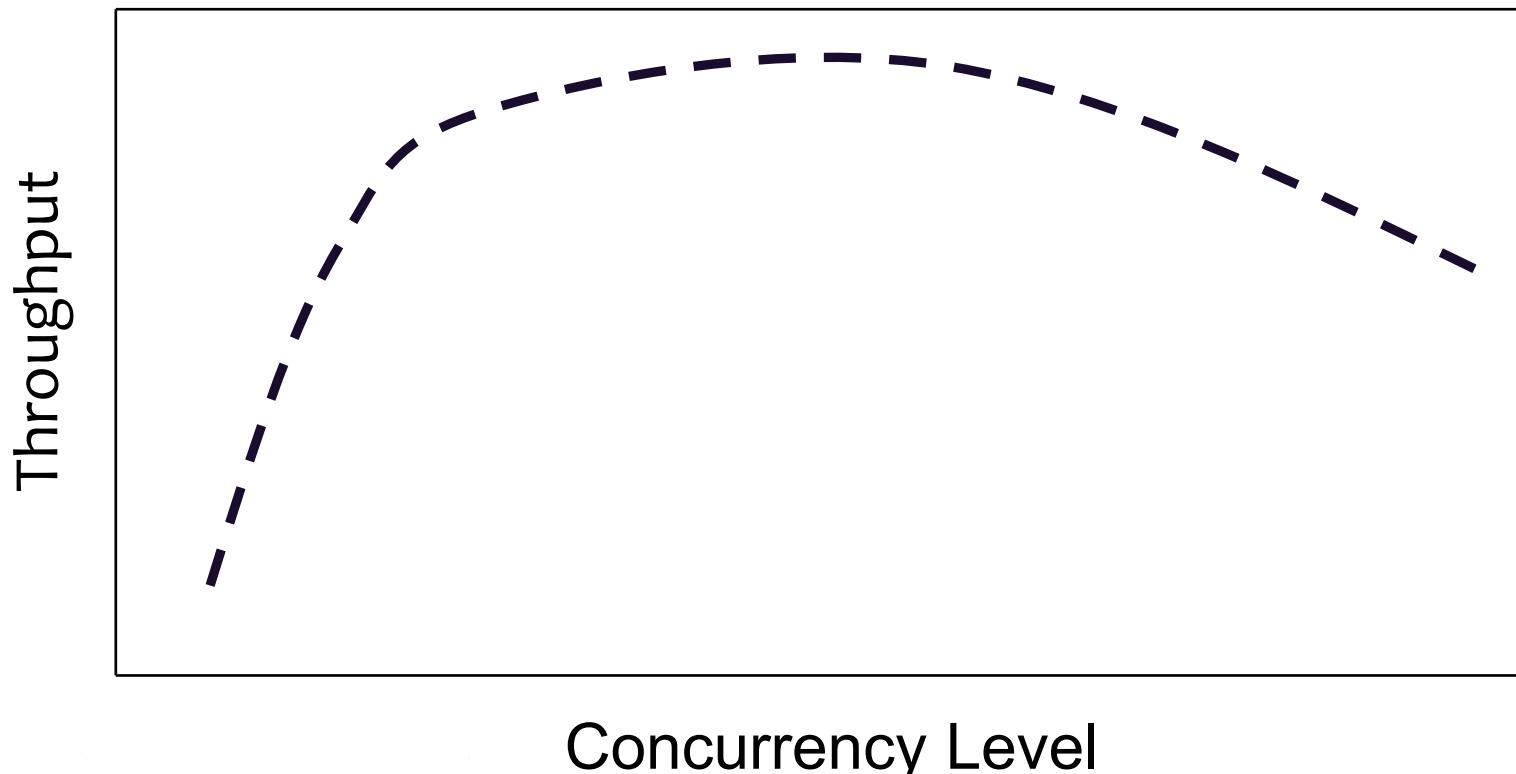
Determine Optimal Concurrency Level for Each Path

- ❑ Offline profiling + online updating
- ❑ Inspired by existing works
 - SCT model – [Liu et al. IPDPS' 20]
 - μ Tune – [Sriraman et al. OSDI' 18]



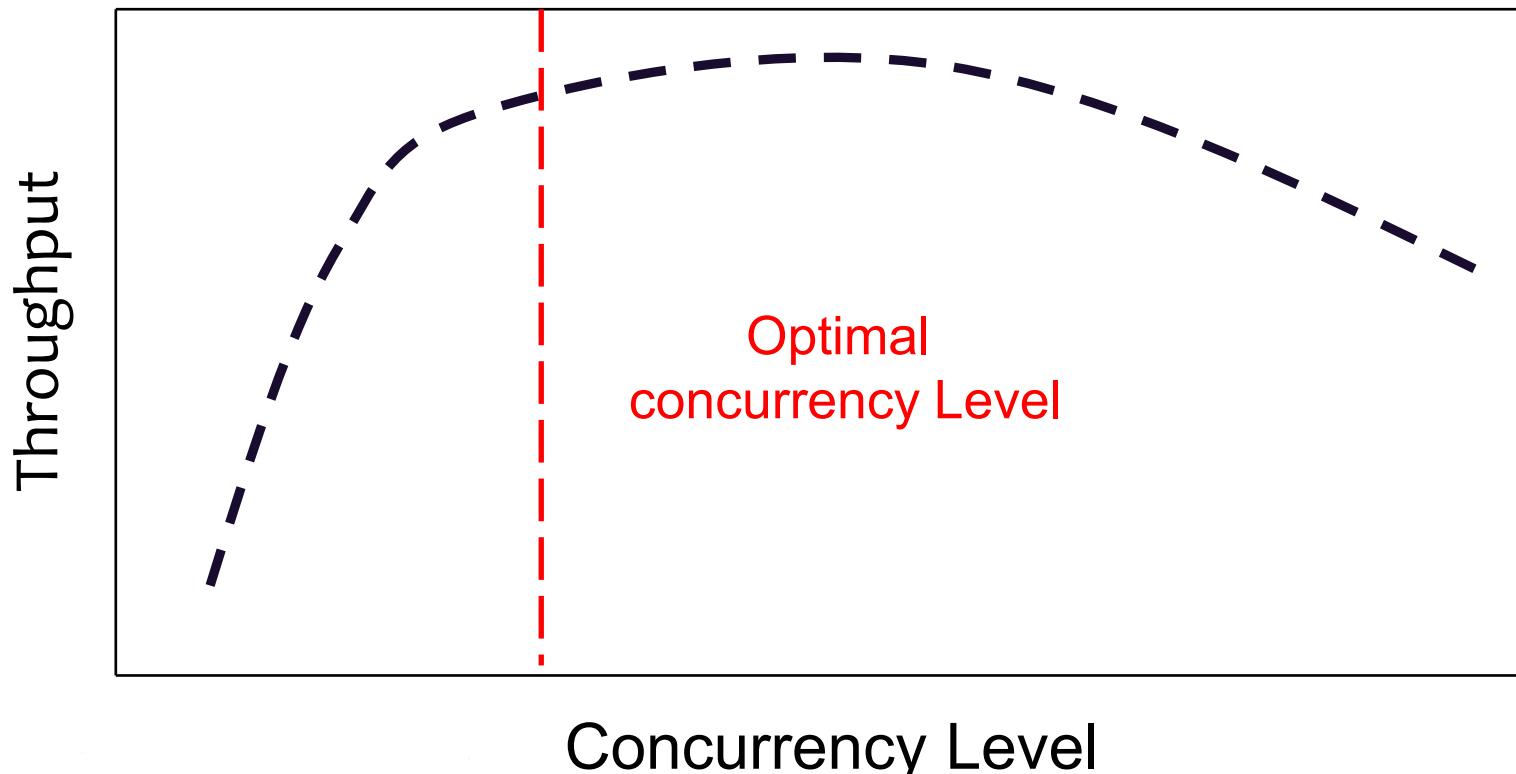
Determine Optimal Concurrency Level for Each Path

- ❑ Offline profiling + online updating
- ❑ Inspired by existing works
 - SCT model – [Liu et al. IPDPS' 20]
 - μ Tune – [Sriraman et al. OSDI' 18]



Determine Optimal Concurrency Level for Each Path

- ❑ Offline profiling + online updating
- ❑ Inspired by existing works
 - SCT model – [Liu et al. IPDPS' 20]
 - μ Tune – [Sriraman et al. OSDI' 18]

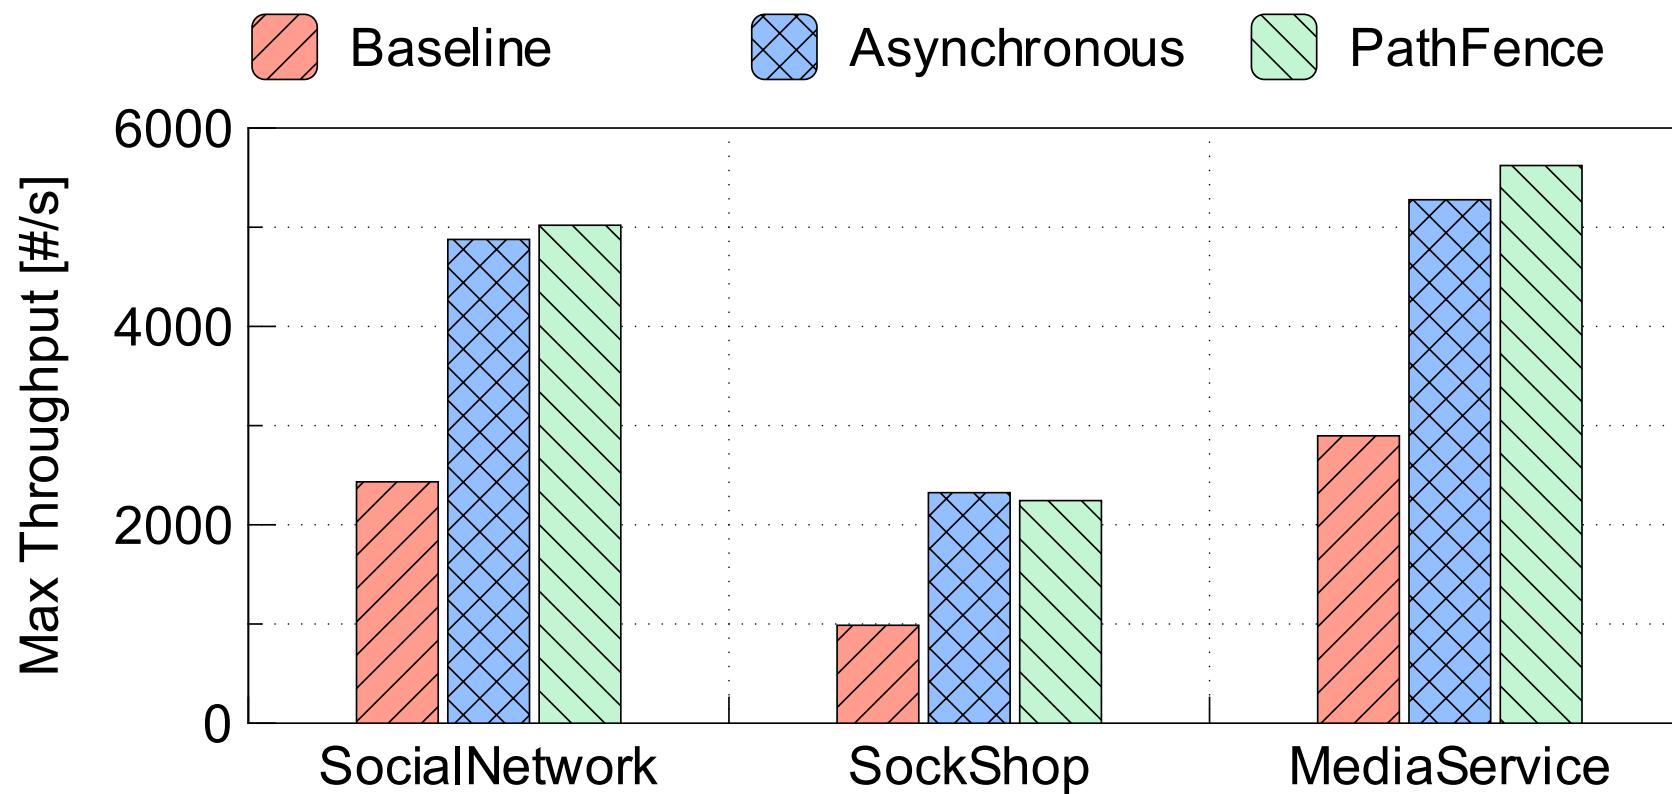


Implementation and Evaluation

- ❑ Implemented with **Apache Thrift** in C++ and **SpringBoot** in Java
- ❑ Deployed on a **Kubernetes cluster** of 20 physical nodes
- ❑ Evaluated with **3 real-world workload traces** and **3 microservices benchmarks**
- ❑ Compare with state-of-the-art (**Asynchronous invocation**)

Results

- ❑ Increases the average system throughput by up to **2x**
- ❑ Reduces the average tail latency by up to **80%**
- ❑ Reduces the number of dropped/timed out requests by over **90%**



Results

- Increases the average system throughput by up to **2x**
- Reduces the average tail latency by up to **80%**
- Reduces the number of dropped/timed out requests by over **90%**

App	Workload Variations	Average Response Time [ms]			99th Percentile Response Time [ms]			Dropped Requests [#]		
		Base	Async.	Our	Base	Async.	Our	Base	Async.	Our
SS	Large	249	171	153	3013	378	394	1833	0	0
	Quick	217	165	158	2374	319	327	797	0	0
	Slow	208	161	143	1178	293	299	273	0	0
SN	Large	314	255	223	7045	1094	1037	5981	173	189
	Quick	289	253	218	3097	1075	1021	2973	134	181
	Slow	273	244	252	1329	598	633	533	0	0
MS	Large	445	397	283	3470	1233	1092	3987	239	204
	Quick	416	311	332	3511	1204	997	2429	131	87
	Slow	391	297	286	1573	871	453	1266	0	0

Conclusion

Conclusion

- ❑ Empirical study of response time instability
 - Cross-path dependencies

Conclusion

- **Empirical study of response time instability**

- Cross-path dependencies

- **Common factors** cause resource contentions

- Internal collocation interference
 - External bursty workload

Conclusion

- **Empirical study of response time instability**
 - Cross-path dependencies
- **Common factors** cause resource contentions
 - Internal collocation interference
 - External bursty workload
- **Path-awareness** soft resources (i.e., threads) scheduling
 - Avoiding soft resource contention
 - Reducing cross-path dependencies

Conclusion

- **Empirical study of response time instability**

- Cross-path dependencies

- **Common factors** cause resource contentions

- Internal collocation interference
 - External bursty workload

- **Path-awareness** soft resources (i.e., threads) scheduling

- Avoiding soft resource contention
 - Reducing cross-path dependencies

- **PathFence achieve high performance** as Asynchronous invocation

- Significantly reducing programming overhead

Thank you!

Check out the full paper for more details!
(xgu5@lsu.edu)