

F₃: An FPGA-accelerated FaaS Framework


Charalampos Mainas, Martin Lambeck, Bruno Scheufler,
Laurent Bindschaedler, **Atsushi Koshiba**, Pramod Bhatotia



MAX PLANCK INSTITUTE
FOR SOFTWARE SYSTEMS

Function-as-a-Service (FaaS)

Function-as-a-Service (FaaS) simplifies the cloud computing model as **serverless**




FaaS achieves easy deployment, high scalability, and low cost (pay as you go)


Performance limitations of FaaS platforms

Compute-intensive tasks are being deployed on FaaS platforms

Compute-intensive workloads



Machine learning



Data analytics



LLM

CPU-centric FaaS platforms



Amazon



Microsoft




Google

Running compute-heavy workloads **only on CPUs** are not cost-/power-efficient

Accelerators in the cloud


Various types of accelerators are being adopted in the cloud



GPUs



Tensor Processing Units
(TPUs)



Field Programmable Gate Arrays
(FPGAs)


Our target

Can we leverage **FPGAs** to accelerate serverless functions?

FPGAs are a good fit to accelerate various, ever-changing cloud workloads

- Reconfigurability
- Optimizable for specific computations
- Potential performance benefits

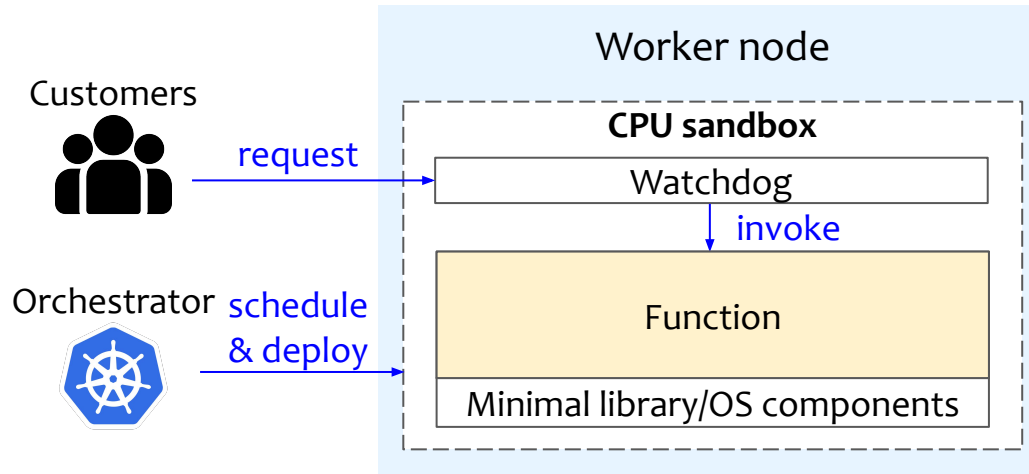
10-100x speedups



FPGA promises accelerating cloud workloads, **but...**

FPGA and serverless model mismatches

There are ***hard-to-reconcile gaps*** between FPGA and serverless computing model




FPGA adoption may lose serverless benefits:

- #1 Easy-to-develop
- #2 CPU isolation
- #3 Low invocation latency
- #4 CPU orchestration

FPGA and serverless model mismatches

There are **hard-to-reconcile gaps** between FPGA and serverless computing model



FPGA adoption may lose serverless benefits:

- #1 Development complexity
- #2 Lack of isolation
- #3 High invocation latency
- #4 No orchestration for FPGAs

How do we enable **FPGA acceleration for serverless functions** by reconciling mismatches between the serverless computing model and FPGA execution model?

F3: FPGA-accelerated FaaS Framework

The first end-to-end FaaS framework for FPGA acceleration and orchestration

Properties:

- Programmability
 - High-level APIs for FPGA control
- Multi-tenancy
 - Hypervisor-enforced isolation
- Low latency
 - Fast FPGA reconfiguration & CPU-FPGA communication
- Resource efficiency
 - FPGA-aware orchestration

Outline



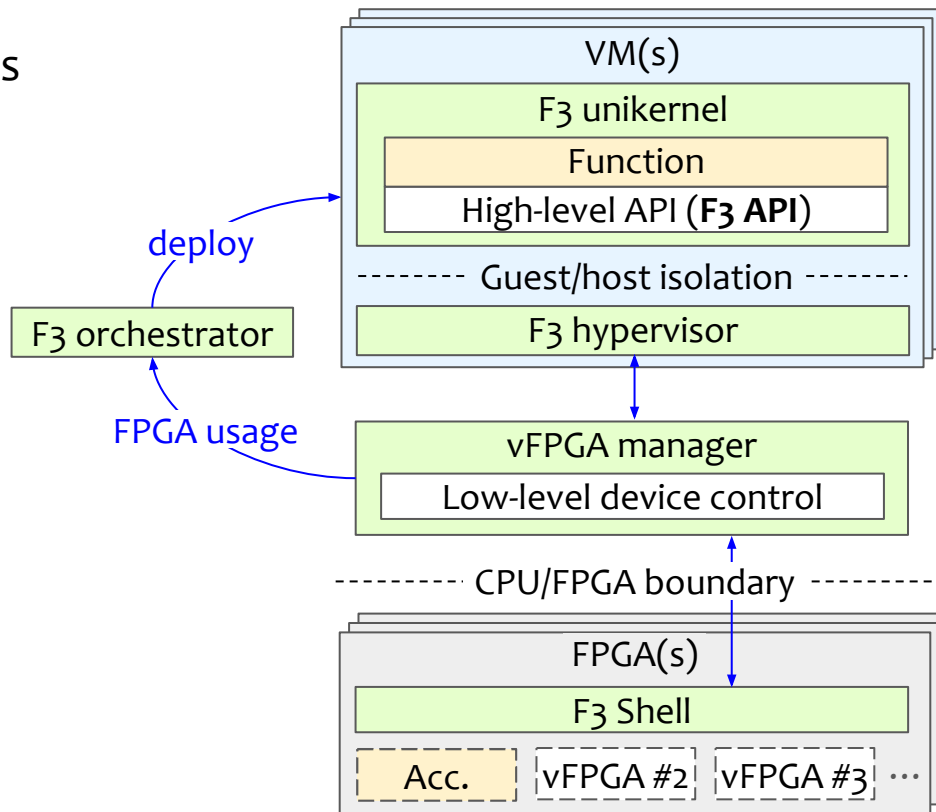
- ~~Motivation~~
- Overview
- Evaluation

F3 overview

An end-to-end HW/SW co-design bridges a gap between FPGA and serverless

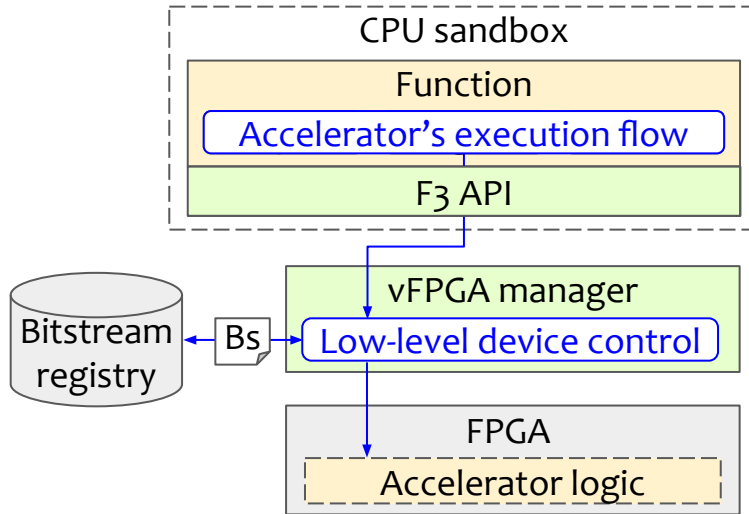
Key components:

- #1 F3 API
- #2 F3 unikernel & hypervisor
- #3 F3 Shell & vFPGA manager
- #4 F3 orchestrator




#1: F3 API

Accelerator-agnostic API that delegates low-level jobs to the backend



```
in  = f3_alloc_buffer(INPUT_SIZE);
tmp = f3_alloc_buffer(INPUT_SIZE);
out = f3_alloc_buffer(OUTPUT_SIZE);
id1  = f3_call_fpga('gzip', in, tmp);
id2  = f3_call_fpga('aes', tmp, out);
f3_wait_fpga(id1, id2);
```

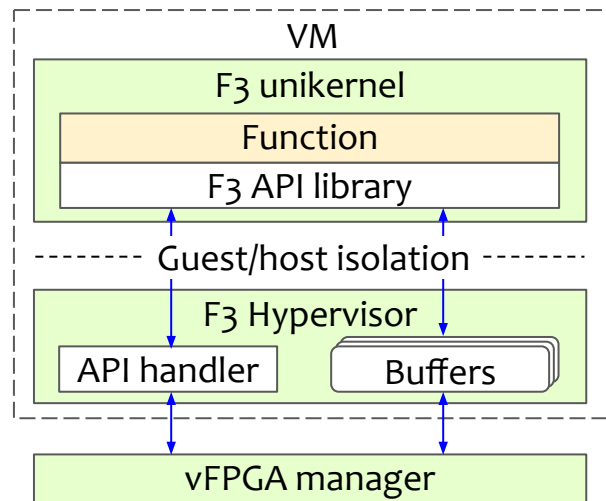


F3 API abstracts the complex FPGA execution flow from developers

#2: F3 unikernel and hypervisor

A *lightweight isolation mechanism* for latency-sensitive serverless functions

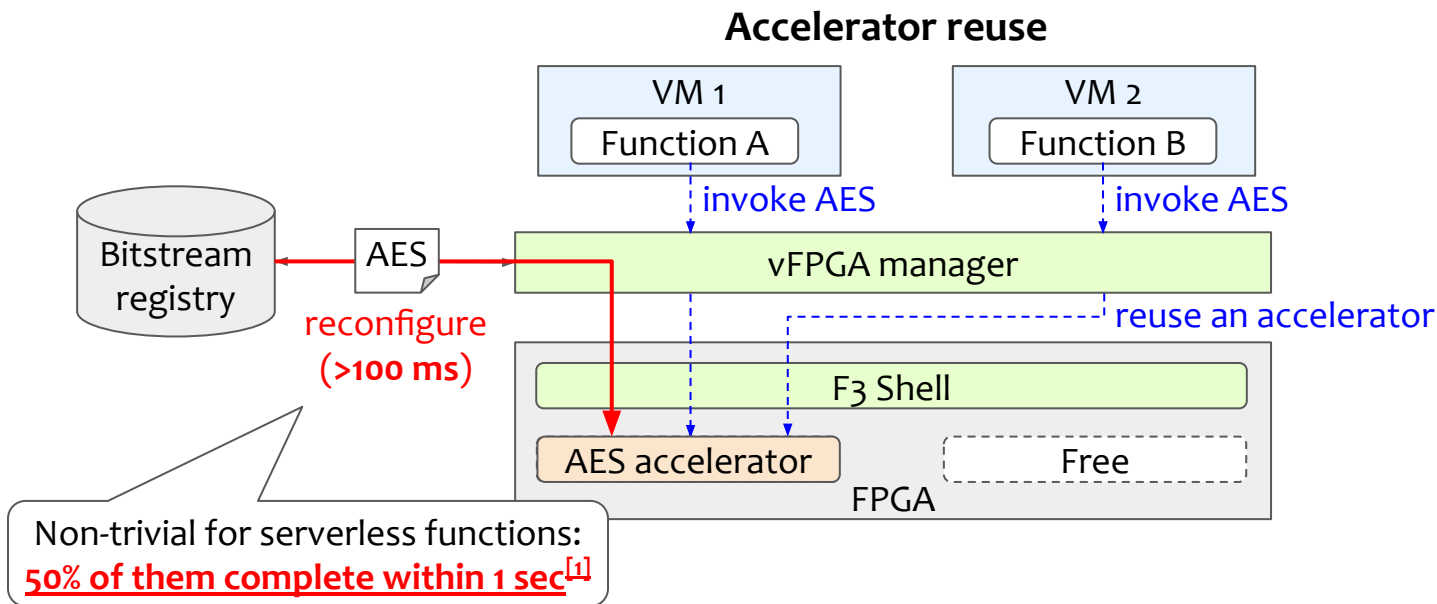
- F3 unikernel
 - ensures isolation for CPU contexts
- F3 hypervisor
 - bridges guest functions and FPGAs
- vFPGA manager
 - restricts access to FPGA resources



F3 guarantees multi-tenant isolation while retaining the lightweightness

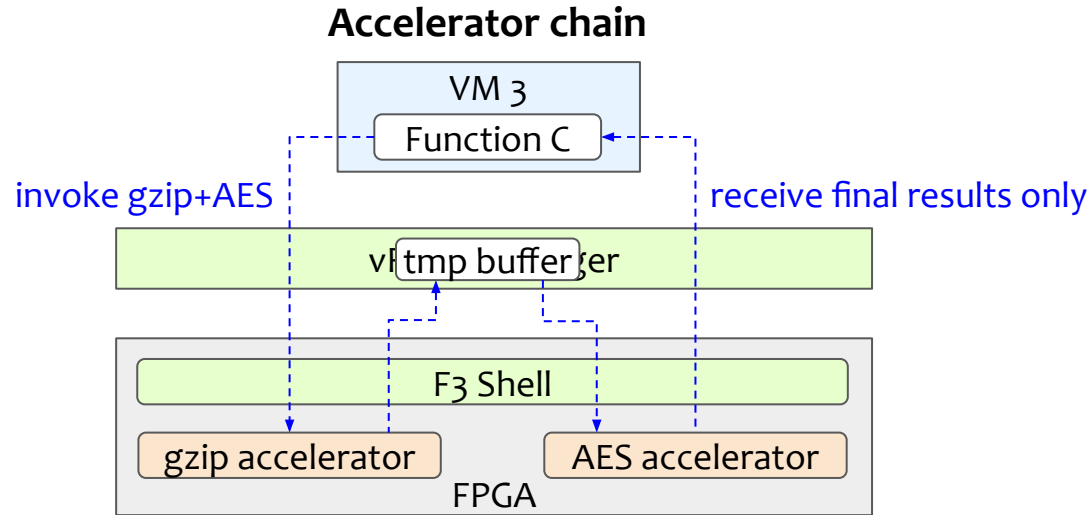
#3: F3 Shell and vFPGA manager

Low-level device management *mitigating FPGA-related overheads*



#3: F3 Shell and vFPGA manager


Low-level device management *mitigating FPGA-related overheads*



F3 leverages Shell features to reduce FPGA's start-up/communication latencies

#4: FPGA-aware orchestrator

Collecting FPGA-related metrics for *FPGA-centric orchestration*



F3 orchestrator achieves fair, resource-efficient function scheduling

Outline



- ~~Motivation~~
- ~~Overview~~
- Evaluation

Questions:

- What is the end-to-end performance gain of F3?
- How effective is FPGA function (accelerator) chaining?

**See our paper
for more results!**

Experimental setup:


- F3 prototype built upon OpenFaaS, Kubernetes, containerd, and Kata-runc
- Cluster : 1x leader (Xeon G5317@3.0GHz), 3x workers (Xeon G6238R@2.2GHz)
- FPGA : each worker node equipped with 1x AMD FPGA (Alveo U50)

Applications:

- 10 compute-intensive workloads suitable for FPGA acceleration
 - Encryption (AES-128-ECB)
 - Hashing (SHA3-512, MD5)
 - Compression (GZIP)
 - Convolutional neural network (HLS4ML^[2])
 - Data analytics (HyperLogLog)
 - Etc.


End-to-end performance

Lower is better ↓



F3 achieves **28.6x average speedups (up to 150.3x)** over a CPU-only baseline

Chaining GZIP and AES accelerators



F3 function chaining achieves **1.4x to 1.7x speedups** against OpenFaaS baselines

Hard-to-reconcile differences between serverless and FPGA execution models

- **Development complexity** due to low-level device management
- **Lack of isolation** for user code/data distributed across CPU & FPGA
- **High invocation latency** due to FPGA reconfiguration & communication overheads
- **No orchestration support** for FPGAs

F3: FPGA-accelerated FaaS Framework

- **High-level API** for easy FPGA execution
- **Multi-tenant isolation** for CPU & FPGA contexts
- **Low invocation latency** assisted by the FPGA Shell
- **FPGA-aware orchestration**

Paper



Code



<https://github.com/TUM-DSE/F3>