



HPDC 2025
July 20-23, 2025, Notre Dame, IN, USA



LiteForm: Lightweight and Automatic Format Composition for Sparse Matrix-Matrix Multiplication on GPUs

Zhen Peng (PNNL)
Polykarpos Thomadakis (PNNL)
Jacques Pienaar (Google)
Gokcen Kestor (Barcelona Supercomputing Center)

July 21, 2025

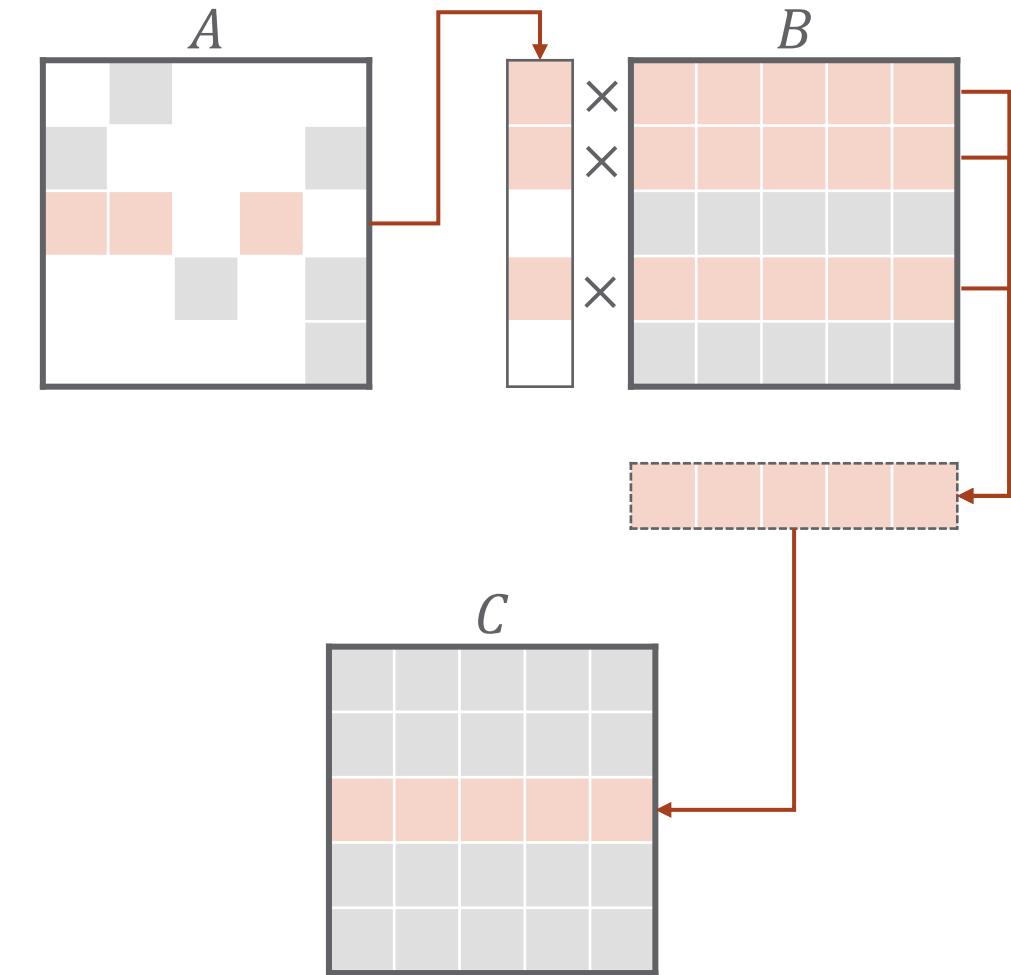


BATTELLE

PNNL is operated by Battelle for the U.S. Department of Energy

Motivation: Sparse Computations on GPUs

- GPUs
 - Excel at dense, regular tasks
 - Struggle with sparse ones like SpMM* ($C = A \cdot B$, A sparse, C and B dense)
- Challenges
 - Irregular memory access
 - Load imbalance
 - Warp divergence
- Key
 - Choosing right sparse format is crucial, but matrices vary in sparsity—single format often suboptimal
- Existing issues
 - Fixed formats lack flexibility
 - Composable formats require costly tuning



SpMM Computational Pattern

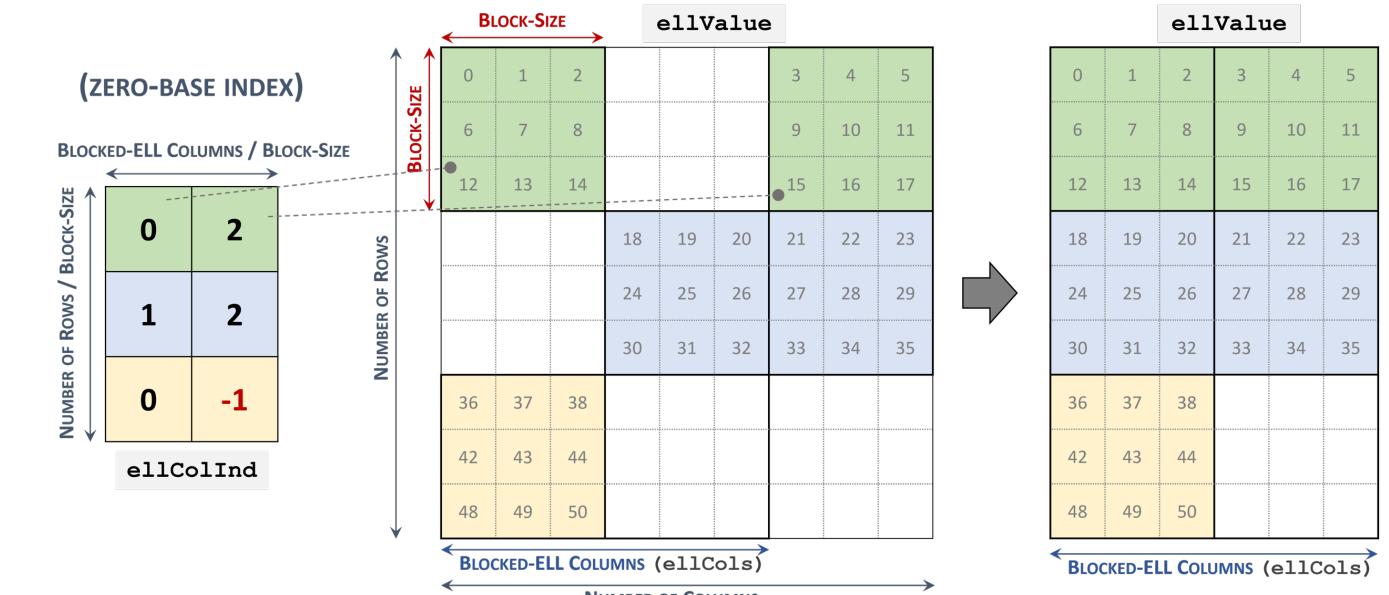
* Sparse matrix-matrix multiplication

Background: Sparse Formats

- Elementwise formats
 - COO, CSR, Ellpack (ELL)
- Blockwise formats
 - BCSR, Blocked-ELL, Sliced-ELL
- Benefits of blocks
 - Shared memory reuse
 - Aligned access
 - Loop unrolling
- Drawbacks
 - Padding ratio up to 99% → memory explosion
 - Less flexible for various sparsity

Sparse Matrix							
Rows	0	1	2	3	4	5	6
0	a					b	
1		c					
2	d	e			f		
3							
4	g	h	i	j	k	l	
5							
6	m	n					
7					o		

Sparse Matrix Ellpack Format



Blocked Ellpack Format

Prior Work and Limitations

- Fixed Formats
 - cuSPARSE, Triton, etc
 - Optimized but input-dependent
- Auto-Selection
 - Auto-SpMV, Seer, etc
 - Machine learning (ML) picks format but ignores intra-matrix sparsity variations
- Composable Formats
 - SparseTIR, STile
 - Adapt to patterns but high construction overhead (auto-tuning/microbenchmarks)

Type	Work	Auto Format Selection	Sparsity Aware	Format Construction Overhead
Fixed Formats	cuSPARSE ^[1] , Triton ^[2] , etc.	✗	✗	Low
Auto-Selection	Auto-SpMV ^[3] , Seer ^[4] , etc.	✓	✗	Low
Composable Formats	SparseTIR ^[5] , STile ^[6]	✗	✓	High
	LiteForm (ours)	✓	✓	Low

[1] NVIDIA, cuSPARSE

[2] P. Tillet, et al., Triton, MAPL 2019

[3] M. Ashoury, et al., Auto-SpMV, arXiv 2023

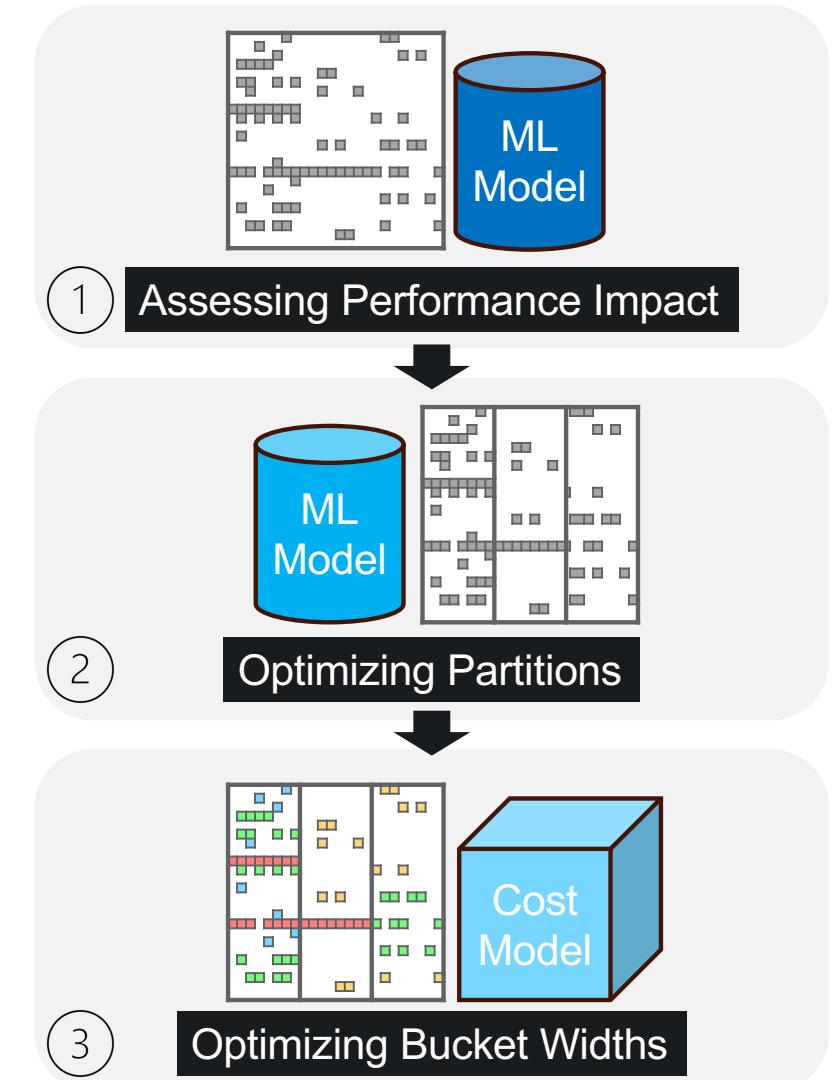
[4] R. Swann, et al., Seer, CGO 2024

[5] Z. Ye, et al., SparseTIR, ASPLOS 2023

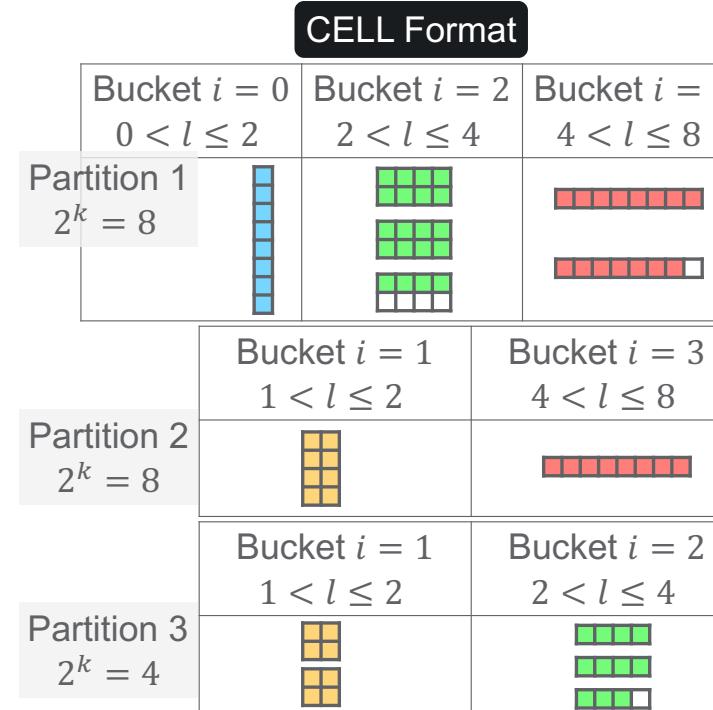
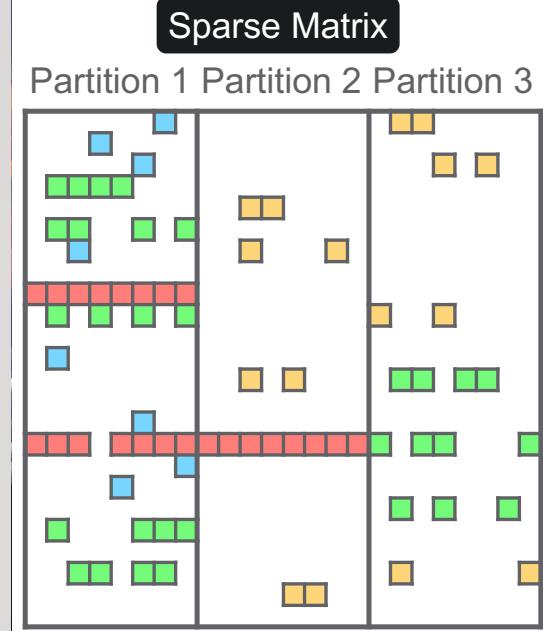
[6] J. Fang, et al., STile, PACMMOD 2024

Overview of LiteForm

- Lightweight framework for automatic format composition for SpMM
- LiteForm's Composable Format
 - CELL (Composable Ellpack) format
 - 3-level blockwise: partitions, buckets, blocks
- LiteForm's Workflow
 - ML predicts if CELL > fixed formats
 - ML sets partition count
 - Cost model + search optimizes bucket widths
- Contributions
 - CELL design
 - ML predictors and cost model
 - No runtime auto-tuning

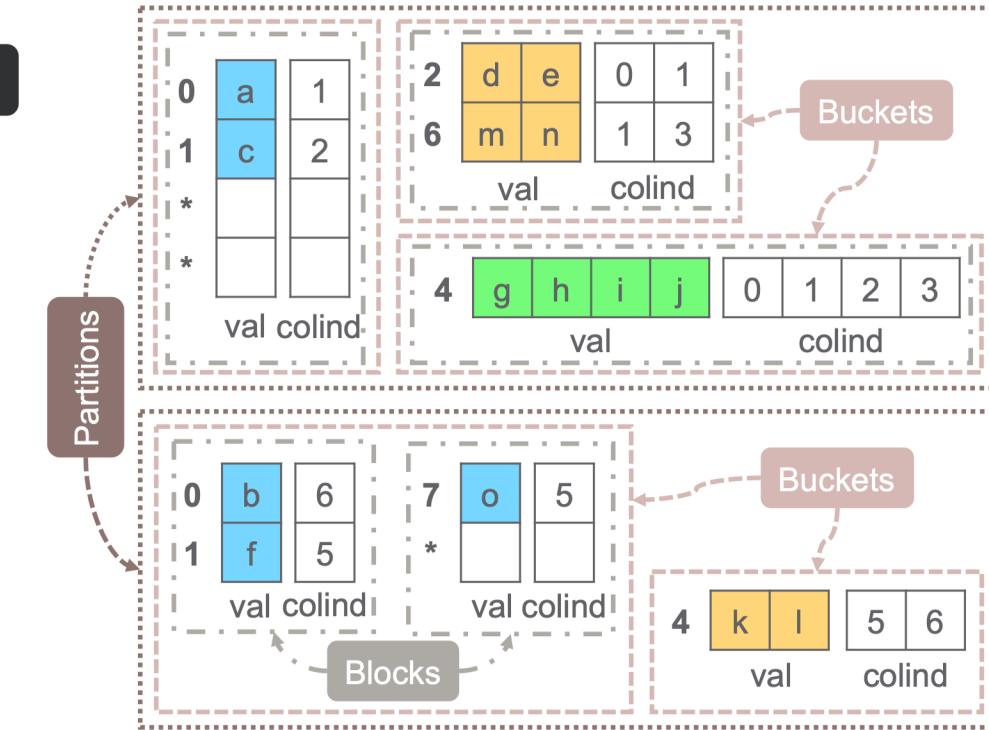


Composable Ellpack (CELL) Format Design



Sparse Matrix in CELL Format

		Columns							
		0	1	2	3	4	5	6	7
Rows	0	a					b		
	1			c					
Rows	2	d	e			f			
	3								
Rows	4	g	h	i	j		k	l	
	5								
Rows	6	m		n			o		
	7								



- 3 Levels:
 - Columns → Partitions (even divide, reduce padding)
 - Rows → Buckets (group by row length l that $2^{i-1} < l \leq 2^i$)
 - Elements → Blocks (2^{k-i} rows, fixed non-zeros 2^k , map to GPU thread blocks)
- Flexible buckets per-partition
- Balancing non-zeros in blocks

Automatic Format Composition

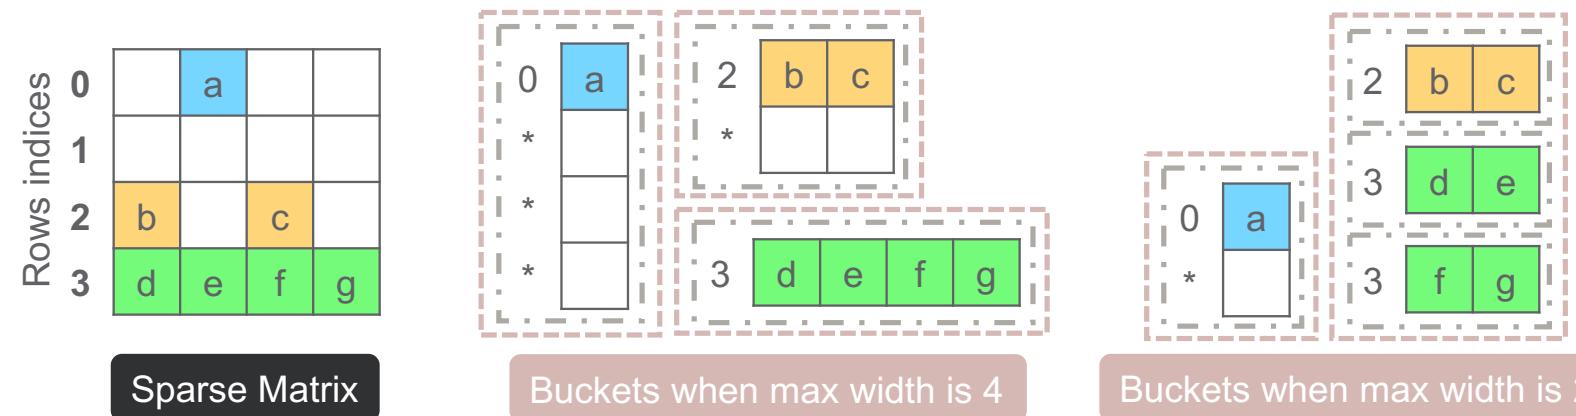
- Step 1: ML assesses CELL, predicts if $>1.1x$ speedup over CSR and BCSR
- Step 2: ML classifier predicts number of partitions
- Step 3: Use a cost Model and search for bucket widths (see in the next slide)
- Training: run SpMM on formats to collect best execution time and configuration.
 - Overhead is amortized over future uses
 - Use Random Forest after evaluation

Sparse matrix features to predict format
Number of rows
Number of columns
Number of non-zero elements
Average number of non-zeros per row
Minimum number of non-zeros per row
Maximum number of non-zeros per row
Standard deviation of non-zeros per row

Sparse matrix features to predict number of partitions
Number of rows
Number of columns
Number of non-zero elements
Average density of non-zeros per row
Minimum density of non-zeros per row
Maximum density of non-zeros per row
Standard deviation of non-zeros density per row
Product of other dimensions in the kernel

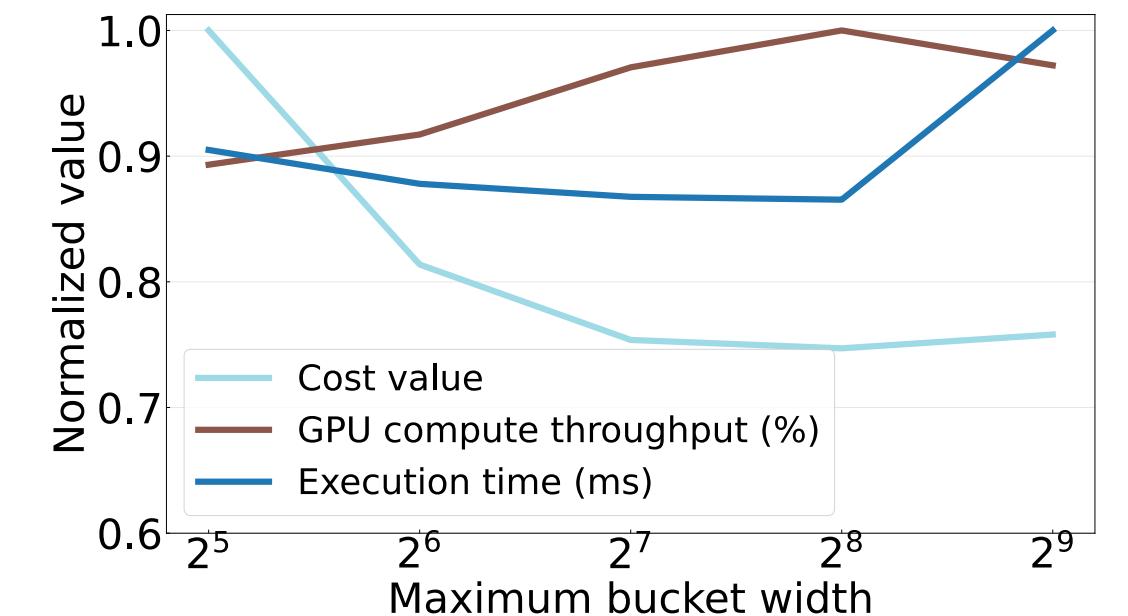
Optimizing Bucket Widths

- The max bucket width trade-off:
 - Larger widths → 👍 fewer row index accesses, coarser workloads, ⚠️ more zero padding
 - Smaller widths → 👍 less zero padding, ⚠️ more index accesses, more atomic writing



- Memory-centric cost model

$$\begin{aligned} \text{cost}(x) = & 2 * (\text{rows in bucket} * \text{width}) \text{ for } A \\ & + \text{unique columns} * J \text{ for } B \\ & + \text{Atomic} * (\text{rows in } C * J) \text{ for } C \end{aligned}$$
- Use a search algorithm to find optimal widths

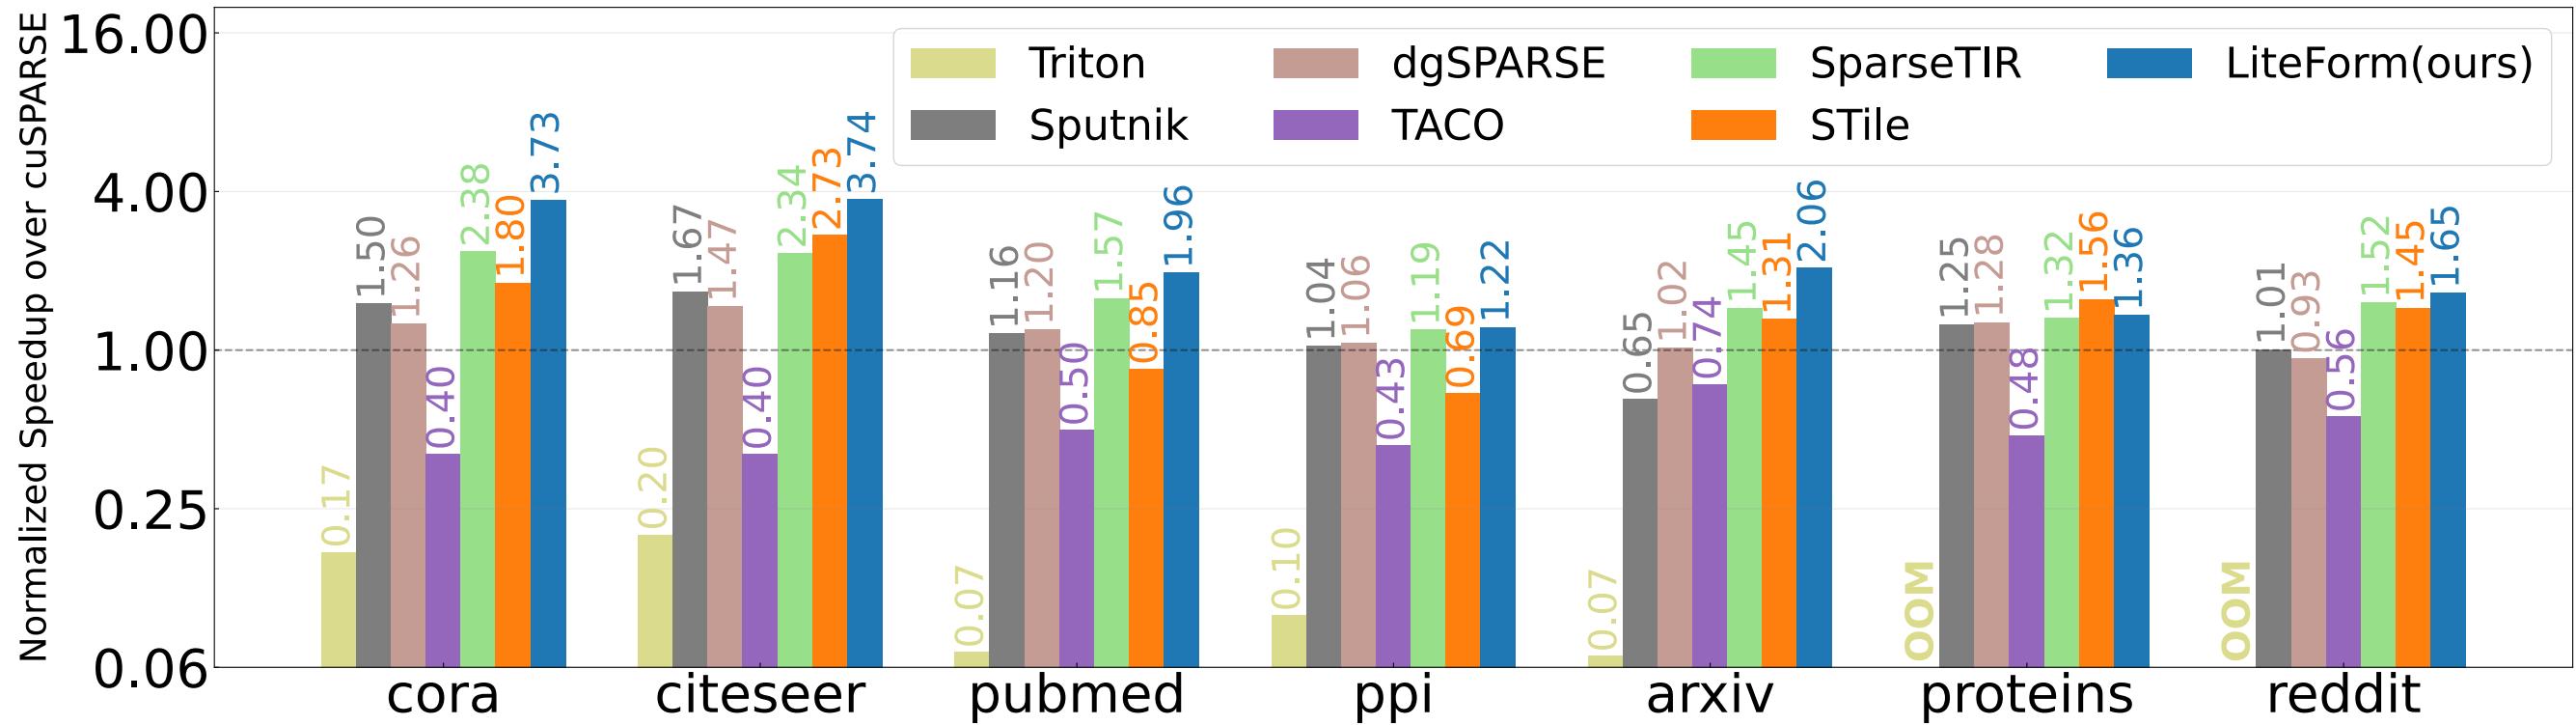


Implementation and Evaluation

- Built on SparseTIR and TVM; Use scikit-learn for ML model (Random Forest)
- Hardware: NVIDIA V100 GPU
- Baselines: cuSPARSE, Triton, Sputnik, dgSPARSE, TACO, SparseTIR, STile
- Datasets: 7 GNN graphs + 1,351 SuiteSparse matrices
- Metrics: Speedup, overhead

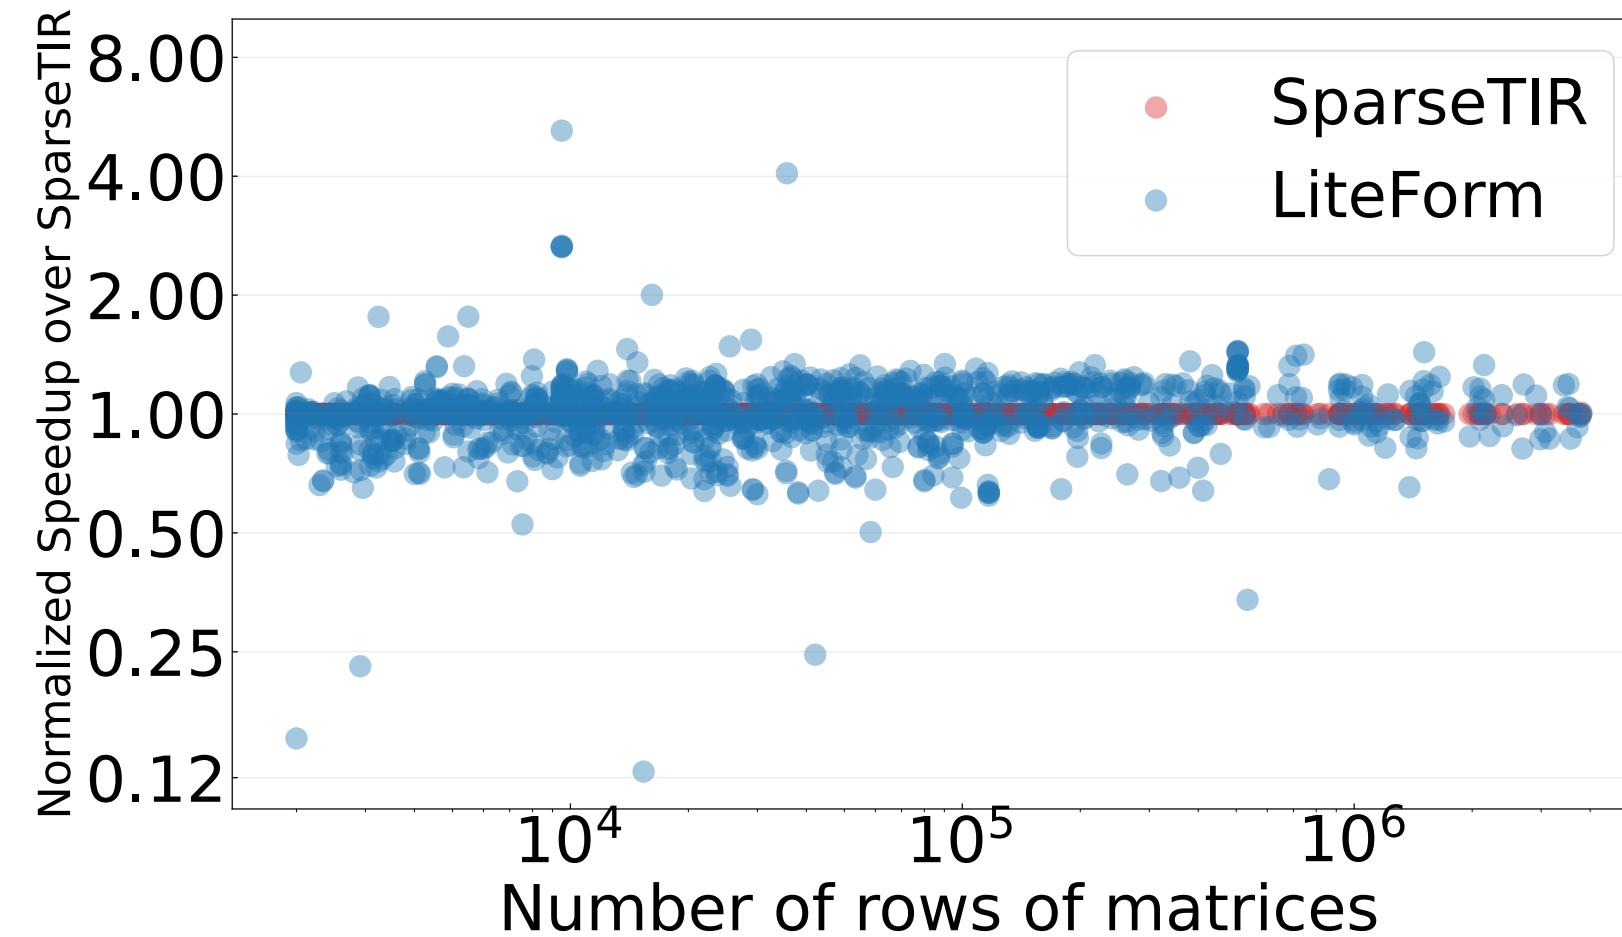
Graph	#nodes	#edges	Density
cora	2,708	10,556	1.44E-03
citeseer	3,327	9,228	8.34E-04
pubmed	19,717	88,651	2.28E-04
ppi	44,906	1,271,274	6.30E-04
arxiv	169,343	1,166,243	4.07E-05
proteins	132,534	39,561,252	2.25E-03
reddit	232,965	114,615,892	2.11E-03
SuiteSparse	2.0K–3.8M	3.1K–300.9M	8.7E-07–0.1

Performance Evaluation (GNN Graphs)



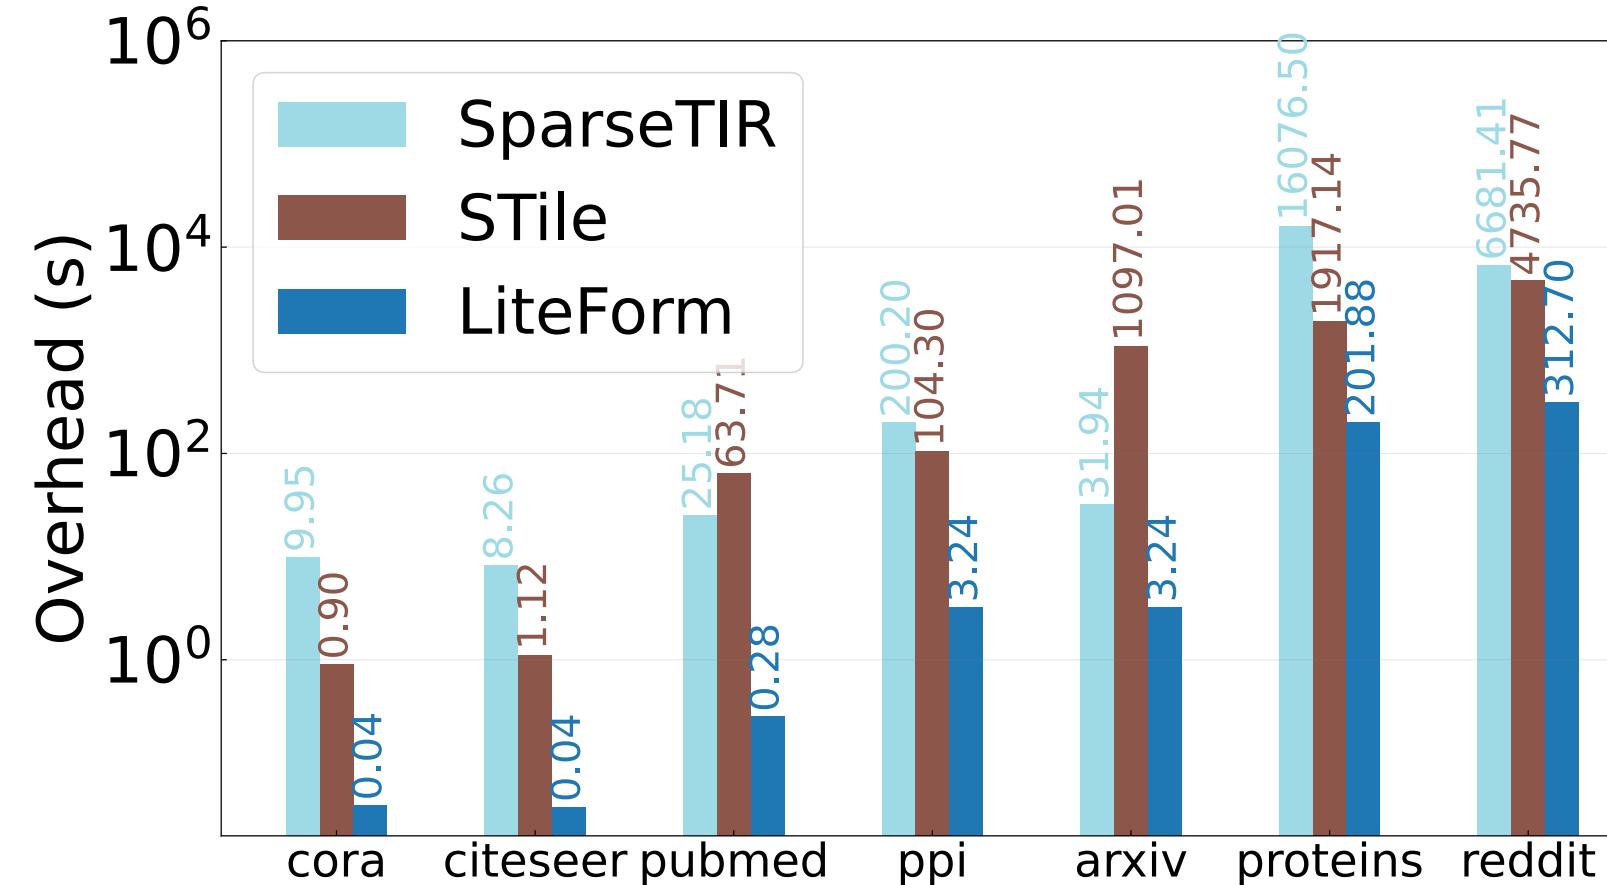
- LiteForm achieved 2.06X geometric mean speedup over cuSPARSE, 1.26X over SparseTIR, and 1.52X over STile

Performance Evaluation (SuiteSparse Graphs)



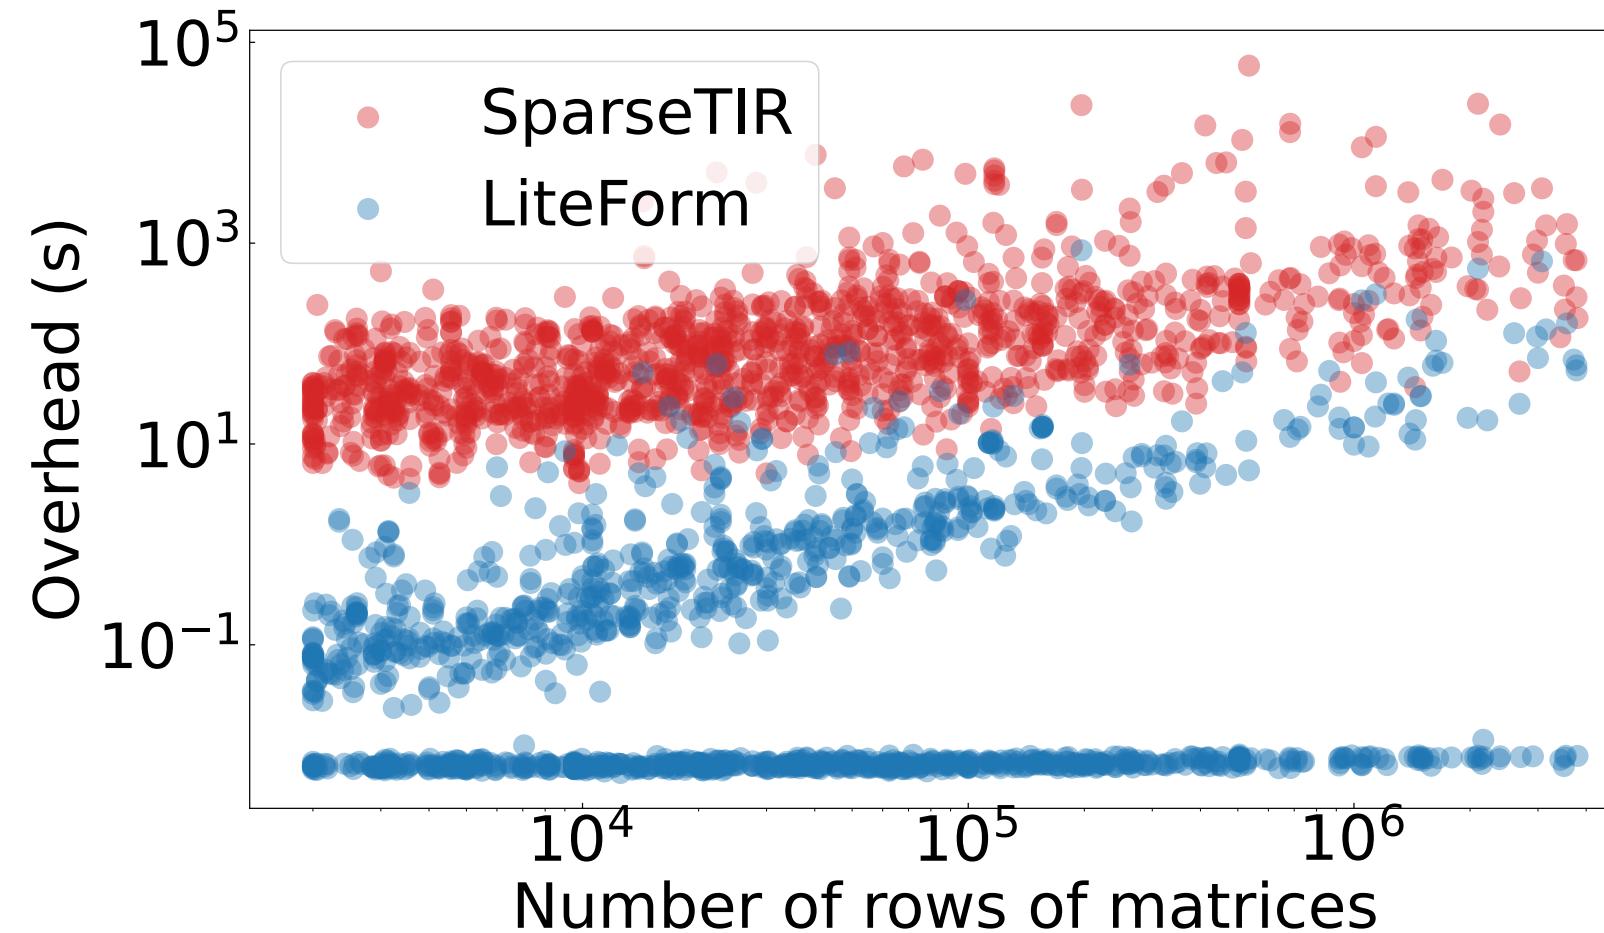
- LiteForm achieved 0.99X geometric mean speedup over SparseTIR
- SparseTIR used auto-tuning to determine the optimal configuration

Overhead Evaluation (GNN Graphs)



- SparseTIR: auto-tuning
- STile: microbenchmarking
- LiteForm: inference and searching (not including training)
- SparseTIR has 65.5X geometric mean overhead, STile has 42.3X

Overhead Evaluation (SuiteSparse Graphs)



- SparseTIR has 1150.2X geometric mean overhead

Limitations of LiteForm

- Needs collecting historical performance data, and requires model retraining for new architectures and kernels
- Historical data may not cover extreme cases, such as a large number of partitions, and extremely wide buckets
- Has not utilized Tensor Cores, or multiple GPUs

Conclusion

- LiteForm: a lightweight and automatic format composition framework for SpMM on GPUs
- Propose the Composable Ellpack (CELL) format with a 3-level blockwise design
- Utilize ML models and a cost model for automatic composition
- Eliminate the need for runtime auto-tuning

Thank you!

Zhen Peng
zhen.peng@pnnl.gov



- Backup slides

Cost Model

- For $C_{ij} = A_{ik} \cdot B_{kj}$, the cost of bucket is measured as memory accesses to matrix A , B , and C (dimensions I, J, K)
- $Cost(x) = cost^{(1)}(x) + cost^{(2)}(x) + cost^{(3)}(x)$
$$= 2 \cdot I^{(1)}W + |set(Ind[i, w])|J + Atomic \cdot I^{(2)}J$$
 - $I^{(1)}$: the number of rows in the bucket of matrix A
 - $I^{(2)}$: the number of corresponding rows in C . $I^{(1)}$ can $> I^{(2)}$ because of folded rows in A
 - W : the bucket width
 - $set(Ind[i, w])$: the set of unique column indices
 - $Atomic$: weight of atomic operation, can be set as $Atomic = \frac{I^{(1)}}{I^{(2)}}$
- Bucket width  (larger W) $\rightarrow cost^{(2)}$  ($|set(Ind[i, w])|$ larger), but $cost^{(3)}(x)$  ($Atomic \cdot I^{(2)}$ and $I^{(1)}$ smaller)
- Bucket width  (smaller W) $\rightarrow cost^{(2)}$  ($|set(Ind[i, w])|$ smaller), but $cost^{(3)}(x)$  ($Atomic \cdot I^{(2)}$ and $I^{(1)}$ larger)

Prediction Evaluation (Predict CELL)

Table 5: Overhead and accuracy of the tested ML models for predicting performance improvement of CELL format.

name	training(s)	inference(s)	accuracy	precision	recall	f1
Random Forest	0.2859	0.0079	88.92%	88.92%	88.92%	88.92%
KNeighbors	0.0024	0.0127	79.31%	79.31%	79.31%	79.31%
Linear SVM	0.0849	0.0098	67.00%	67.00%	67.00%	67.00%
RBF SVM	0.0856	0.0199	73.40%	73.40%	73.40%	73.40%
Gaussian Process	346.2509	0.0697	84.24%	84.24%	84.24%	84.24%
Decision Tree	0.0292	0.0004	85.96%	85.96%	85.96%	85.96%
Neural Net	2.8343	0.0016	66.50%	66.50%	66.50%	66.50%
AdaBoost	0.1828	0.0079	86.45%	86.45%	86.45%	86.45%
Naive Bayes	0.0018	0.0004	63.30%	63.30%	63.30%	63.30%
QDA	0.0022	0.0004	66.75%	66.75%	66.75%	66.75%

- Used 80% of 514 matrices as the training set and the other 20% as the test set
- Random Forest achieved the best accuracy

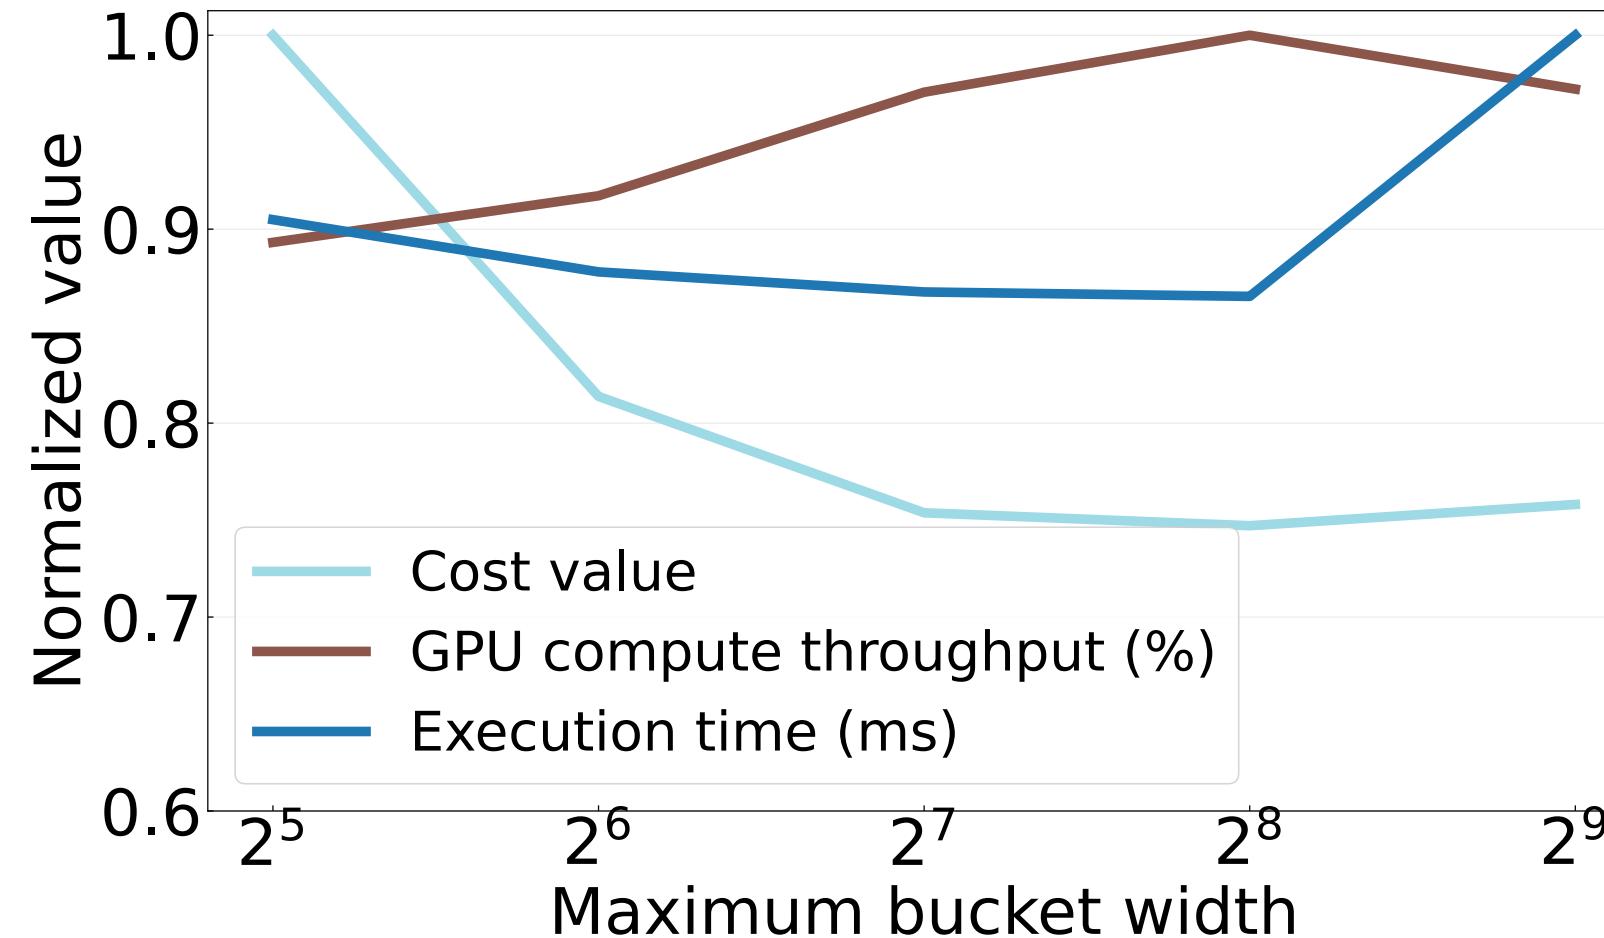
Prediction Evaluation (Predict Num. of Partitions)

Table 6: Overhead and accuracy of the tested ML models for predicting the optimal number of partitions in the CELL format. *cos_sim* stands for cosine similarity.

name	training(s)	inference(s)	accuracy	precision	recall	f1	cos_sim
Random Forest	0.4778	0.0127	87.30%	87.30%	87.30%	87.30%	0.77
KNeighbors	0.0046	0.0321	82.98%	82.98%	82.98%	82.98%	0.23
Linear SVM	0.2273	0.0244	82.45%	82.45%	82.45%	82.45%	0.25
RBF SVM	0.5688	0.0692	82.56%	82.56%	82.56%	82.56%	0.25
Gaussian Process	1481.1395	24.0115	82.56%	82.56%	82.56%	82.56%	0.25
Decision Tree	0.0200	0.0005	85.40%	85.40%	85.40%	85.40%	0.77
Neural Net	3.0432	0.0017	82.45%	82.45%	82.45%	82.45%	0.25
AdaBoost	0.1952	0.0106	82.13%	82.13%	82.13%	82.13%	0.25
Naive Bayes	0.0025	0.0008	56.41%	56.41%	56.41%	56.41%	0.29
QDA	0.0036	0.0011	0.21%	0.21%	0.21%	0.21%	0.25

- Random Forest achieved the best accuracy

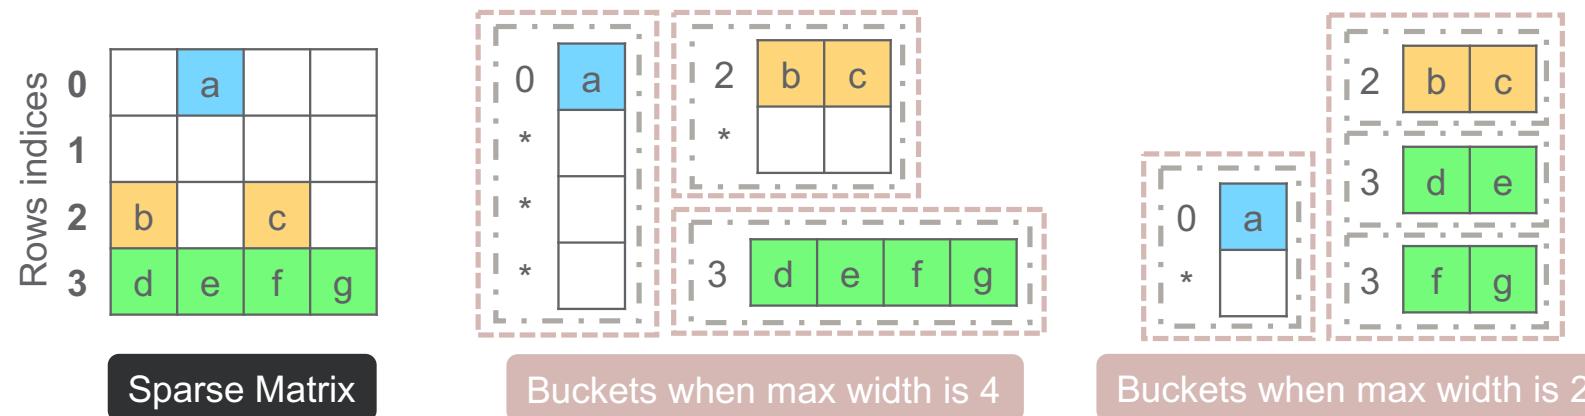
Cost Model Evaluation



- Tested the reddit data set
- The bucket width influence the cost value from the cost model
- When the cost value is the lowest, the GPU compute throughput reaches the highest and the execution time is the shortest

Optimizing Bucket Widths

- CELL format can represent a single long row as multiple rows
- The number of non-zeros in a block is set by the maximum bucket width



A larger maximum bucket width → fewer row index accesses, but more padding
A smaller maximum bucket width → more row index accesses, but fewer padding

- Use a cost model to estimate memory access overhead for given widths
 - Larger bucket width → More overhead to access B , but less overhead to access C
 - Smaller bucket width → Less overhead to access B , but larger overhead to access C
- Use a search algorithm to find optimal widths