

# Can Large Language Models Predict Parallel Code Performance?

Gregory Bolet<sup>1</sup>, Giorgis Georgakoudis<sup>2</sup>, Harshitha Menon<sup>2</sup>, Konstantinos Parasyris<sup>2</sup>, Niranjan Hasabnis<sup>3</sup>, Hayden Estes<sup>1</sup>, Kirk Cameron<sup>1</sup>, Gal Oren<sup>4</sup>


<sup>1</sup>Virginia Tech (VT), <sup>2</sup>Lawrence Livermore National Laboratory (LLNL), <sup>3</sup>Code Metal AI, <sup>4</sup>Technion & Stanford University

## Motivation

**Trend 1:** Large Language Models (LLMs) are becoming ubiquitous in Software Development.




**Trend 2:** Not many Performance Analysis sub-fields using LLMs for GPU execution profiling/analysis



**Trend 3:**

- New GPU hardware is becoming increasingly inaccessible (due to datacenter demand)
- Existing LLM-based GPU-code optimization works assume **hardware access** for profiling




Trend 1  
Trend 2  
Trend 3 **Idea:** Can LLMs predict GPU code performance *without* the need for profiling?

- Execution Time
- FLOP/s
- Cache Misses
- Bytes Read/Written
- FLOP/Byte
- Instructions/Cycle

**Problem:** LLMs are traditionally **BAD** at regression tasks

**Solution:** Focus on a simple *classification* task instead!




## Research Questions

How well can LLMs classify the Arithmetic Intensity (AI) of GPU codes?


**RQ1 (Baseline Roofline Classification)**

- Given the **GPU Roofline specs** and an **explicit AI value**, can an LLM correctly classify the value as BB/CB?



**RQ2 (Source Code Classification)**

- Given the **source code**, **necessary execution specs**, and **minimal instructions**, can an LLM correctly classify the program as BB/CB?



## Dataset Design Decisions


Built + Profiled:  
170 CUDA + 170 OpenMP  
HeCBench Codes

zjin-lcf / HeCBench

1) Concatenate each kernels' source files for prompting



Perf. Metrics  
SPFLOP (FP32),  
DPFLOP (FP64),  
INTOP,  
AI for each kernel



## RQ1: Roofline Understanding

### Experimental Setup

- 120 CB + 120 BB prompts
- Random Rooflines + AI values
- 2, 4, and 8-shot examples
- Fixed temp = 0.1, top\_p = 0.2
- Evaluation metric: **accuracy**

### Results:

Model Name	Reasoning	RQ1 Acc.	RQ1 CoT Acc.
o3-mini-high	✓	100	100
o1	✓	—	—
o3-mini	✓	100	100
gpt-4.5-preview	—	—	—
o1-mini-2024-09-12	✓	100	100
gemini-2.0-flash-001	91.25	92.50	
gpt-4o-2024-11-20	91.25	96.25	
gpt-4o-mini	90.00	100	
gpt-4o-mini-2024-07-18	90.00	100	

### ROI Prompting Template (w/ CoT)

CoT example 1 (shown below):  
Question: Given a GPU having a global memory with a max bandwidth of 45.9 GB/s and a peak performance of 52.22 GFLOP/s, if a program executed with an Arithmetic Intensity of 0.6 FLOP/Byte and a performance of 19.4 GFLOP/s, does the roofline model consider the program as compute-bound or bandwidth-bound?

Thought: The max bandwidth is 45.9 GB/s, and peak performance is 52.22 GFLOP/s. The balance point is at  $32.22 / 45.9 = 0.68$  FLOP/Byte. Because  $0.6 < 0.68$ , it is before the balance point, putting the program in the bandwidth-bound region. The roofline model would consider the program as bandwidth-bound.

Answer: Bandwidth

CoT examples 2-8 [redacted]

Question: Given a GPU having a global memory with a max bandwidth of 99.9 GB/s and a peak performance of 73.45 GFLOP/s, if a program executed with an Arithmetic Intensity of 1.55 FLOP/Byte and a performance of 32.8 GFLOP/s, does the roofline model consider the program as compute-bound or bandwidth-bound?

- All models have a reasonably-good understanding of AI
- Reasoning models** have good prediction accuracy w/ and w/o CoT
- 2 examples was often sufficient

## RQ2: Source Code Classification

### Experimental Setup

- 170 CB + 170 BB CUDA/OpenMP codes
- 2-shot examples
- Fixed temp = 0.1, top\_p = 0.2
- Evaluation metric: **accuracy**

### Pseudo-code examples

Model Name	Reasoning	Input/Output Cost (1M tokens)	RQ2 Acc.
o3-mini-high	✓	\$1.1 / \$4.4	64.12
o1	✓	\$15 / \$60	64.12
o3-mini	✓	\$1.1 / \$4.4	62.06
gpt-4.5-preview	—	\$75 / \$150	59.71
o1-mini-2024-09-12	✓	\$1.1 / \$4.4	59.64
gemini-2.0-flash-001	—	\$0.1 / \$0.4	55.59
gpt-4o-2024-11-20	—	\$2.5 / \$10	52.06
gpt-4o-mini	—	\$0.15 / \$0.6	50.59
gpt-4o-mini-2024-07-18	—	\$0.15 / \$0.6	50.29

### GPU specs

### Exec specs

- Non-reasoning models are akin to a coinflip
- Similar CUDA/OpenMP prediction accuracy
- Room for improvement** with o3-mini-high achieving highest accuracy of 64%


## Conclusions

- SoTA LLMs **do** understand the **Roofline Model** for GPU performance analysis
- SoTA LLMs **can predict parallel code performance** – when limited to classifying Arithmetic Intensity (AI) of CUDA/OpenMP programs
- Reasoning-equipped LLMs (e.g.: o3-mini-high) offer significantly better classification accuracy when compared to non-reasoning LLMs


## Next Steps

### Major Shortcomings:

- Binary classification
- Single-prompting approach



We currently have some success in applying Question Decomposition to estimate FLOPs



Target Name	Empirical FLOP Count	LLM-Estimated FLOP Count	% Diff
resize-cuda	16779307	16777216	0.012 %
zerocopy-cuda	1050389	1048576	0.17 %
iso2dfd-cuda	54419825	53196468	2.24 %
nll-cuda	6006	6273	4.44 %
backprop-cuda	3080240	3080192	0.001 %

## Acknowledgements

This work was funded in part by NSF awards:  
#1838271 VarSys: Managing Variability in High-Performance Computing Systems  
#1939076 iLORE: Computer Systems Performance Integrated Lineage Repository

